

Operating Systems!

OS Interaction (Part 1)

How to Get Stuff Done Using the OS, Processes, Threads, Etc.

Prof. Travis Peters

Montana State University

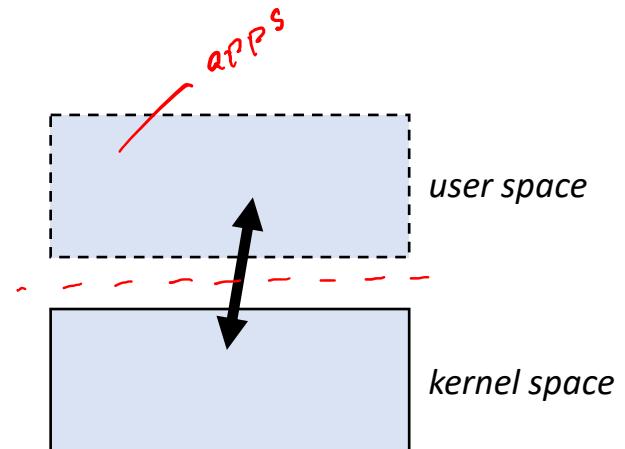
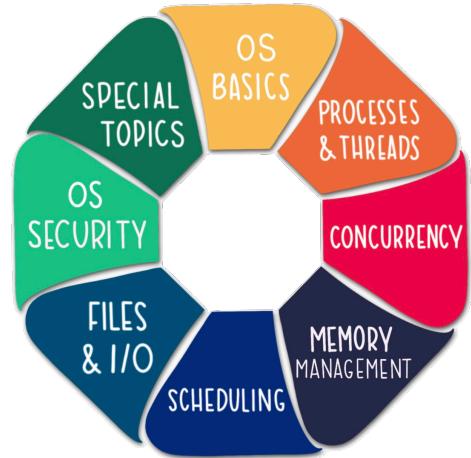
CS 460 - Operating Systems

Fall 2020

<https://www.cs.montana.edu/cs460>

Today

- Announcements
 - Heads up.... PA1 has been released!
Sunday [09/20/2020] @ 11:59 PM (MST)
NO D2L! Code pushed to GitHub == submitted!
- Learning Objectives
 - Understand the big ideas behind the “OS API”
 - **user vs. kernel, modes, syscalls, libraries**
 - Understand the big ideas behind process control
 - [theory] **control info, creation, termination, states, etc.**
 - [reality] **fork, exec, getpid, waitpid, etc.**



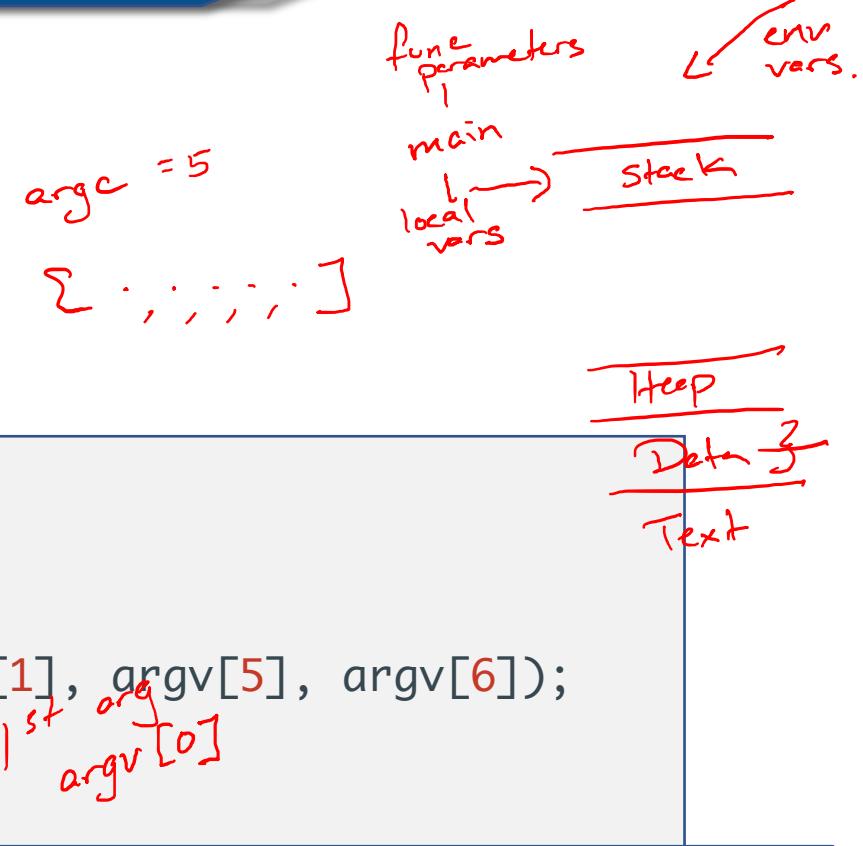
Warm-up (C Programming)

```
#include <stdio.h>
int main(int argc, char **argv[])
{
    printf("%d, %s, %s, %s, %s\n", argc, argv[0], argv[1], argv[5], argv[6]);
    return 0;
}
```

Run prog. What does this program do?
\$ gcc mystery1.c -o mystery1 && ./mystery1 what does this do ?

args — command line args

[5, what, does, ?, error]
* 6, ./mystery1, what, ?, (null), ..., "NULL"
#define NULL 0x00



Re-Orienting

An OS is the thing between the machine and the “user” that makes it easier for the user to get some work done.

Some Key Ideas:

- **multiprogramming:** the same processor can switch back and forth between different “processes” --- each process *thinks* it has the system to itself!
- **interrupt:** mechanism to stop a process and execute an “interrupt handler”
- **switching:** the actual mechanics of taking one process off of the processor and putting another one on.
- **modes:** an OS is typically divided into two (or more) modes of operation
 - **user** (“unprivileged”)
 - **kernel** (“privileged” / “supervisor” / “monitor”)

Modes (Kernel Mode vs. User Mode)

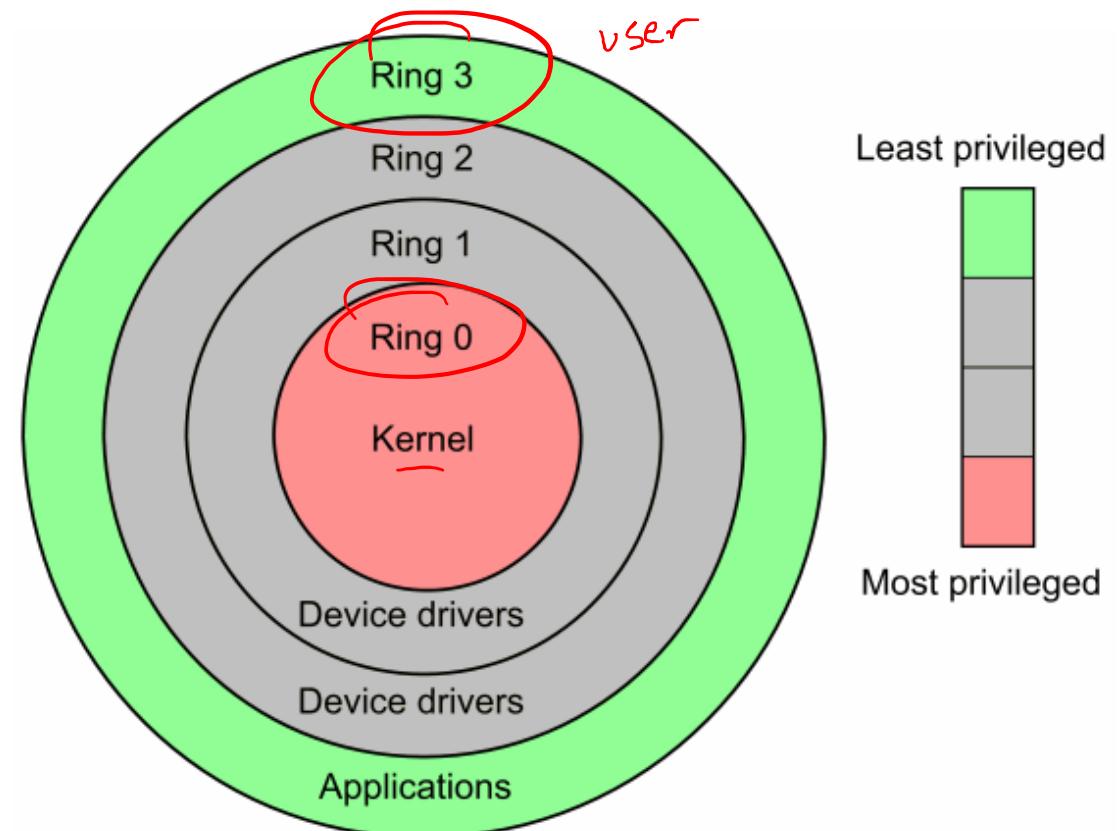
mode
User
-

An OS is typically divided into two (or more) modes of operation

- user ("unprivileged")
 - restricted instruction set
 - restricted access to memory
- kernel ("privileged" / "supervisor" / "monitor")
 - unrestricted instruction set
 - full access to memory and I/O

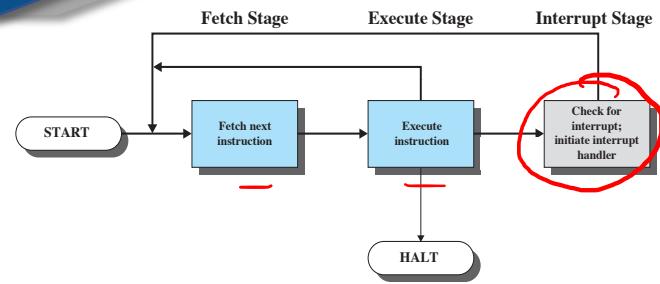


CPU must check whether instruction is permissible before executing it



Credit: <https://blog.codinghorror.com/understanding-user-and-kernel-mode/>

(Revisiting) Interrupts & Traps



Interrupt

- A hardware event that sets flag in the CPU; e.g.,
 - Timer event
 - I/O complete
- } types

Trap

- A software interrupt that is run under predefined conditions; e.g.,
 - Exceptions
 - System Calls
- } types

A Simple Flow

1. Set flag in CPU *(A bit; also some info about the type of interrupt)*
2. CPU cycle detects flag *(Recall fetch-decode-execute cycle w/ "check")*
3. CPU pauses typical execution sequence
 - a) stash current processor state *(By pushing some things onto the stack)*
 - b) switch to kernel mode *(Flip the "mode" bit)*
 - c) start executing at predefined location in memory *(A "handler" - look in the Interrupt Vector; which handler to run depends on the interrupt type)*

While there are subtle distinctions between definitions for interrupts, traps, and exceptions, they really are all just types of *interrupts*... They all *interrupt* the processor and cause it to take some other action.



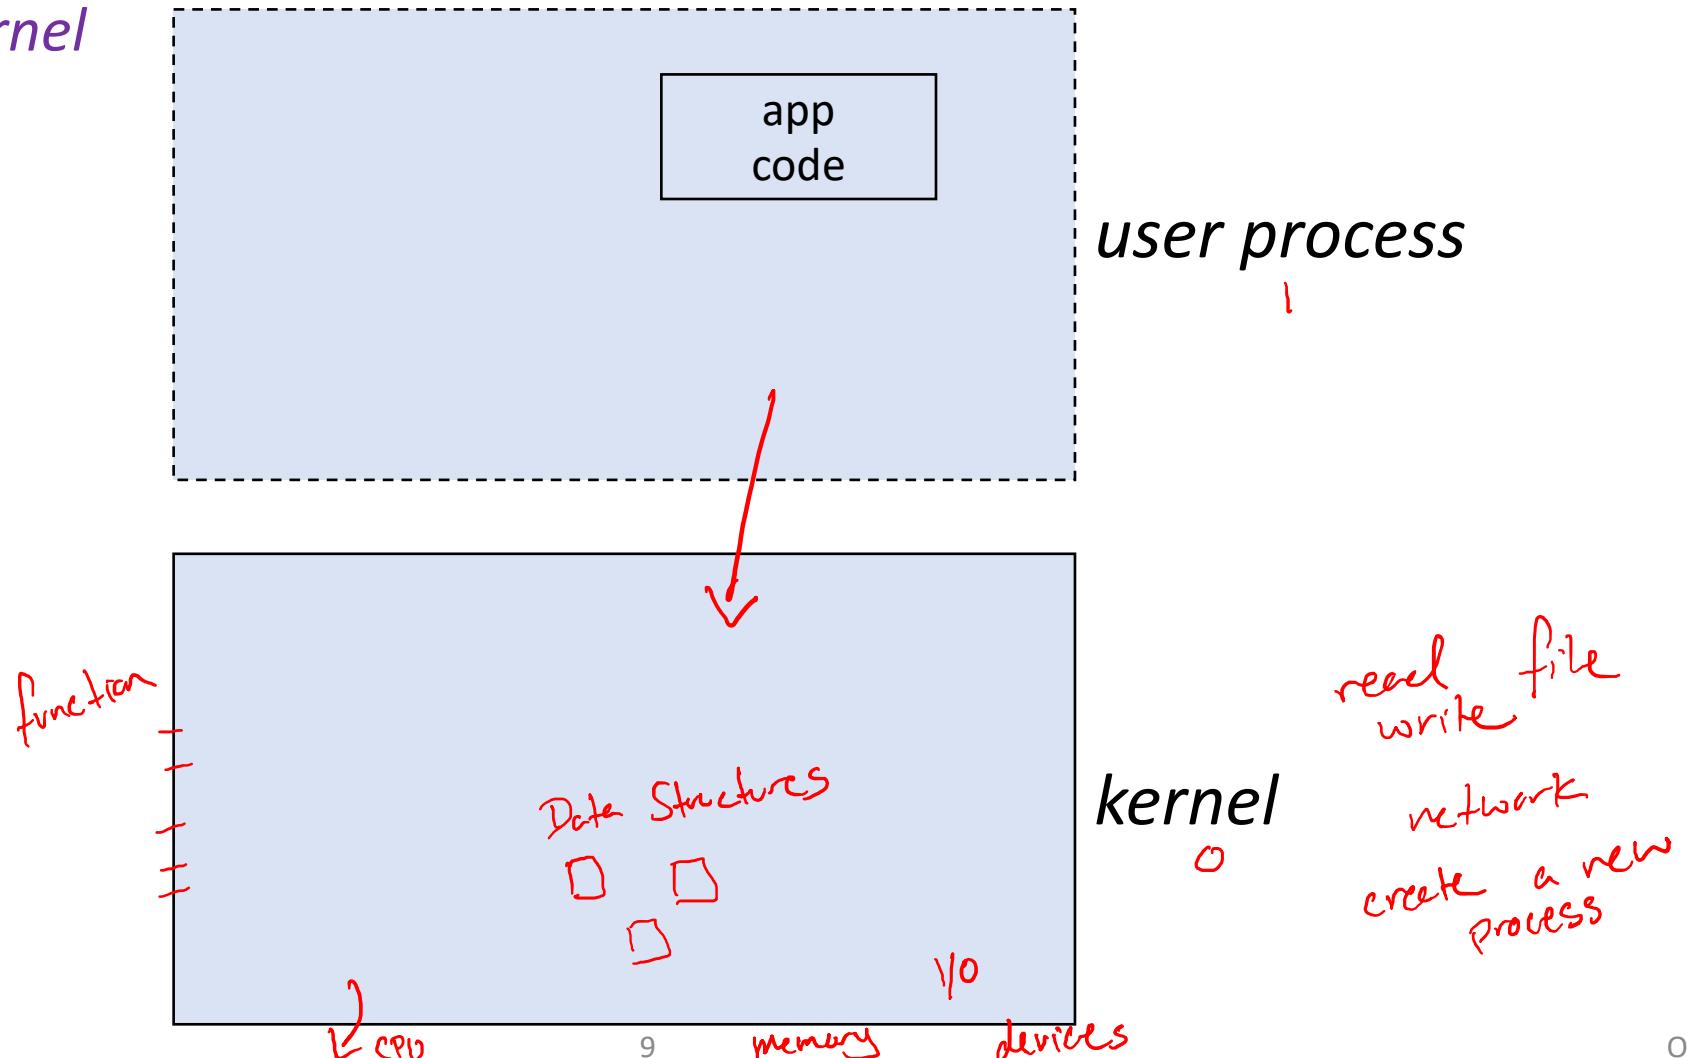
(Go read this SO post: <https://stackoverflow.com/a/3149217>)

Interacting With Your OS

System Calls ("Syscalls")

An application program explicitly requests an OS service by making a system call ("syscall").

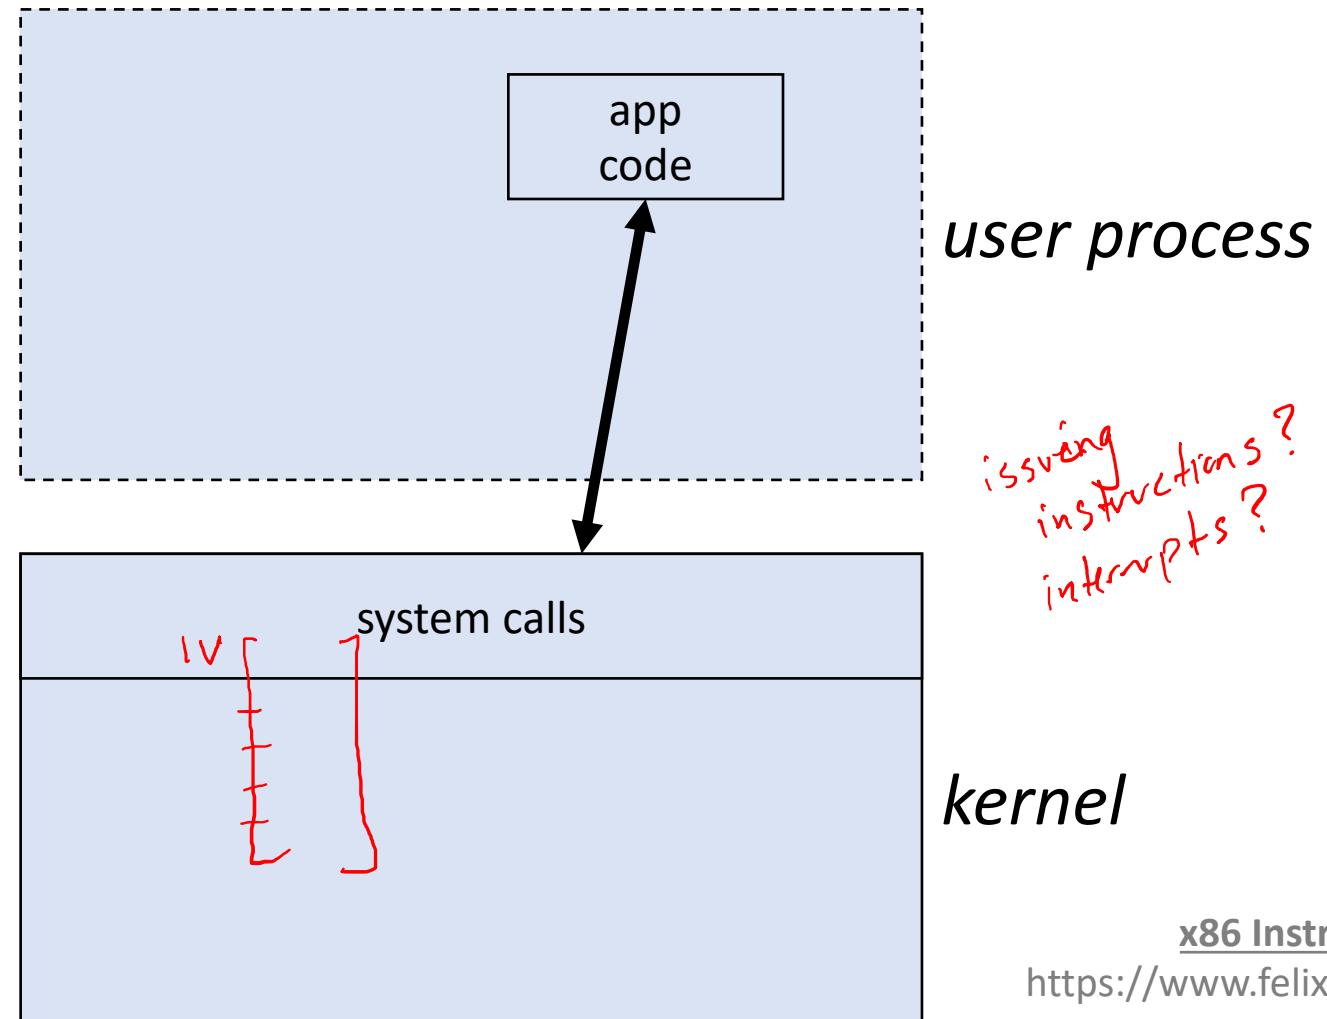
HOW do I use the kernel
and its services?



System Calls ("Syscalls")

An application program explicitly requests an OS service by making a system call ("syscall").

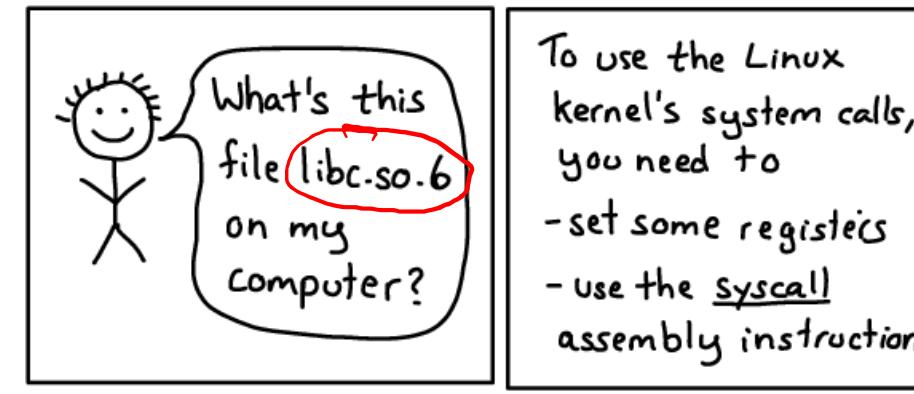
A syscall is invoked with a software interrupt after putting values into registers or on the stack



x86 Instructions: ~~x86~~
<https://www.felixcloutier.com/x86/>

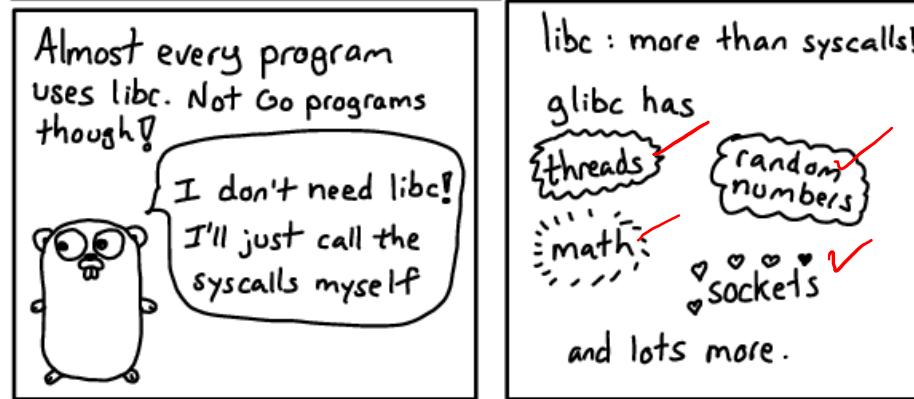
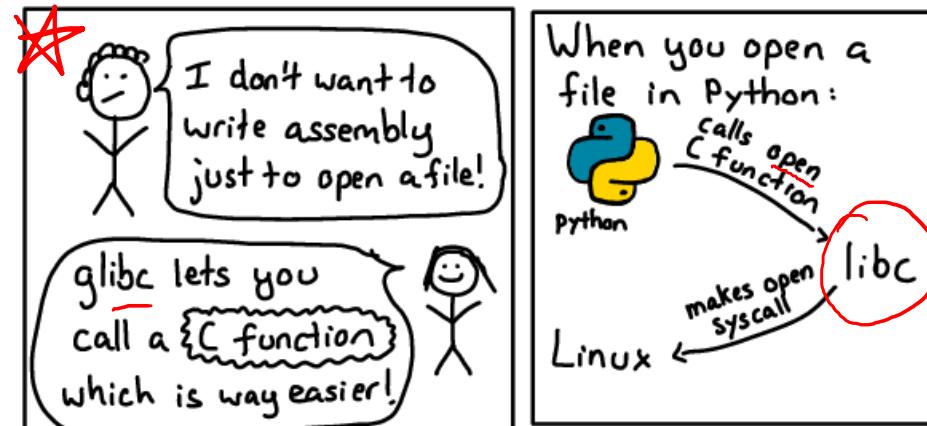
♥ libc ♥

JULIA EVANS
@bork



To use the Linux kernel's system calls, you need to

- set some registers
- use the syscall assembly instruction



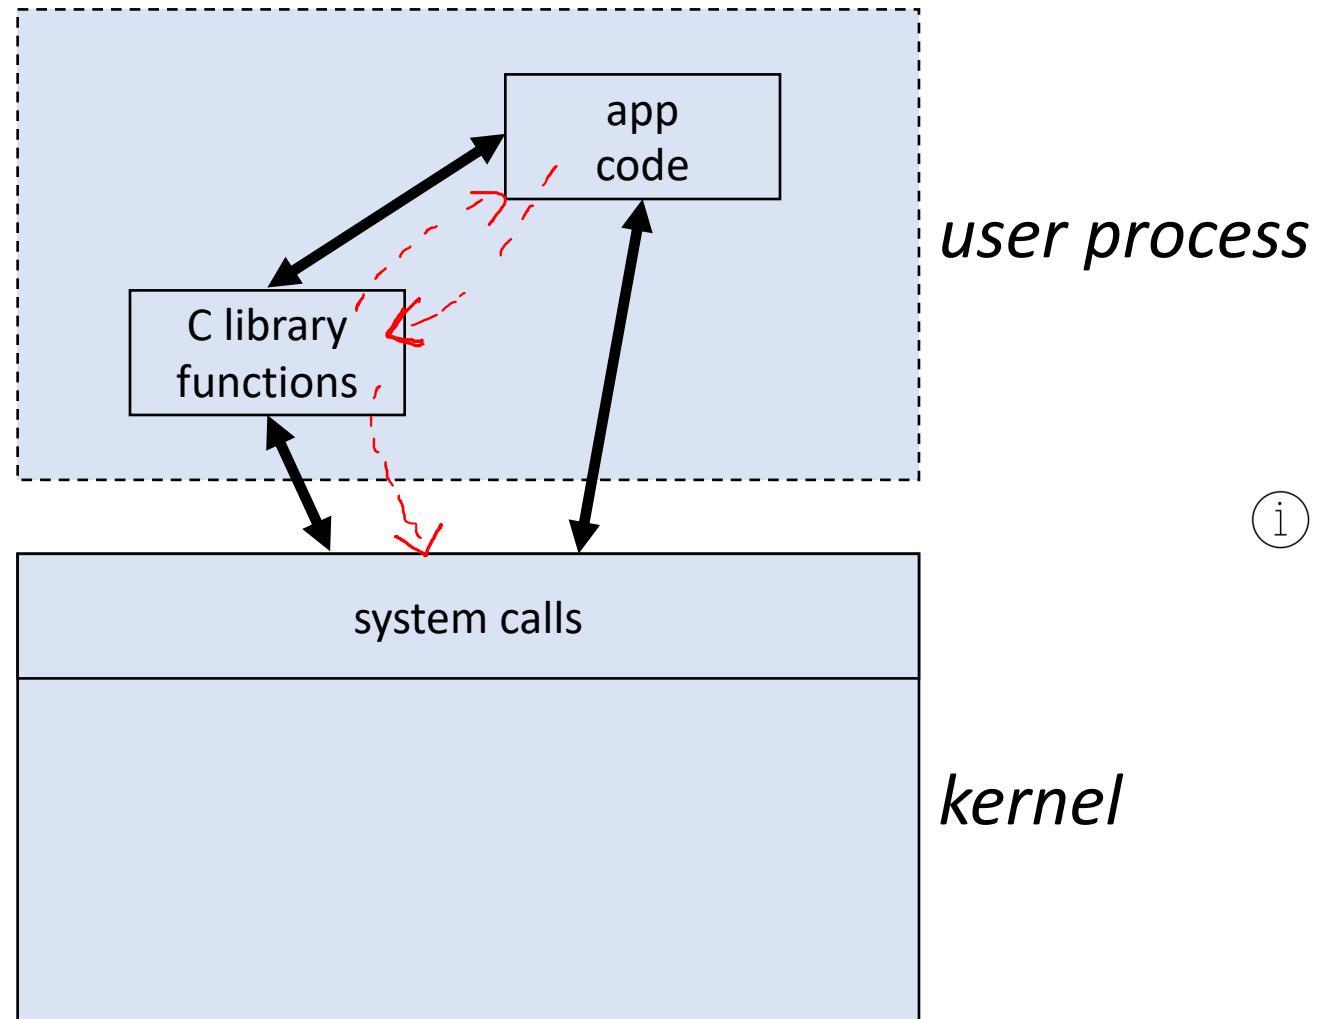
System Calls ("Syscalls")

An application program explicitly requests an OS service by making a system call ("syscall").

A syscall is invoked with a software interrupt

after

putting values into registers or on the stack



try week03/test.c
to play with ideas
of lib vs. OS

Assembler Fragment

A *syscall* is often realized in assembly,
which adheres to the processor's conventions
(this fragment is actually from Yalnix!)

```
// basic macro: move the 4 args and the 'code' into registers
// and throw a software trap (picked up in support/exception.c).
//
// ideally, we should have a different version for each argument count.

#define YSYSCALL(code,A,B,C,D) \
    int rc; \
    __asm__ __volatile__ ( "int $04\n\t" \
                          : "=eax" (rc) \
                          : "eax"(A), "b"(B), "c"(C), "d"(D), "D" (code) ); \
    return rc
```