

Operating Systems!

OS Interaction (Part 2)

How to Get Stuff Done Using the OS, Processes, Threads, Etc.

Prof. Travis Peters

Montana State University

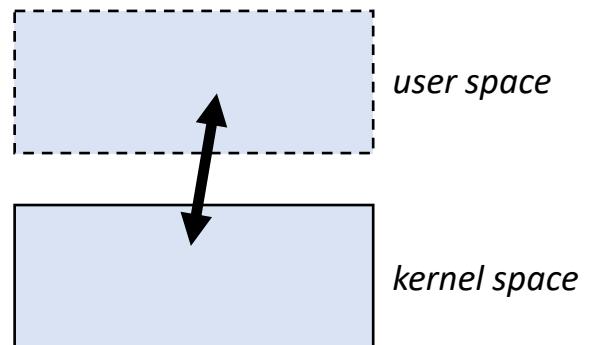
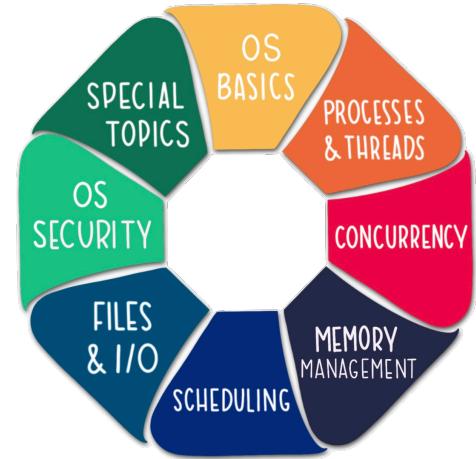
CS 460 - Operating Systems

Fall 2020

<https://www.cs.montana.edu/cs460>

Today

- Announcements
 - Heads up.... PA1 has been released!
Sunday [09/20/2020] @ 11:59 PM (MST)
NO D2L! Code pushed to GitHub == submitted!
- Learning Objectives
 - Understand the big ideas behind the “OS API”
 - user vs. kernel, modes, syscalls, libraries
 - Understand the big ideas behind process control
 - [theory] **control info, creation, termination, states, etc.**
 - [reality] **fork, exec, getpid, waitpid, etc.**



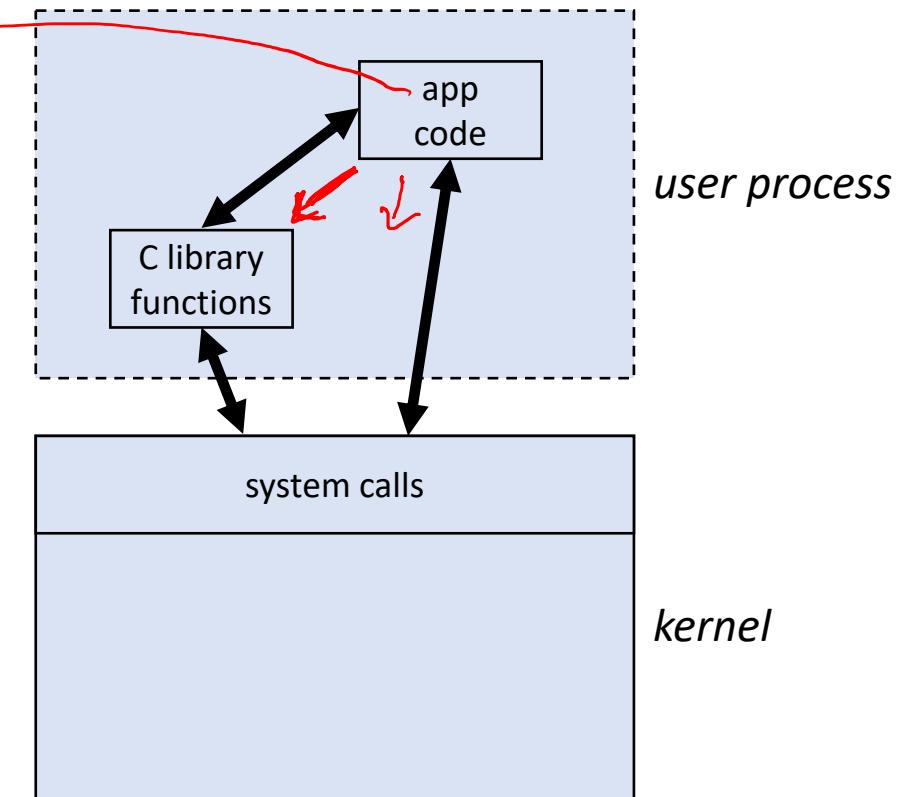
System Calls ("Syscalls")

An example!

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
}
```

```
# Run myexec. What happens?
$ gcc exec1.c -o exec1
$ ./exec1
```



Syscalls

(cont.)
An example!

Registers & Invoking Syscalls

EAX

System Call Number

EBX

1st Argument

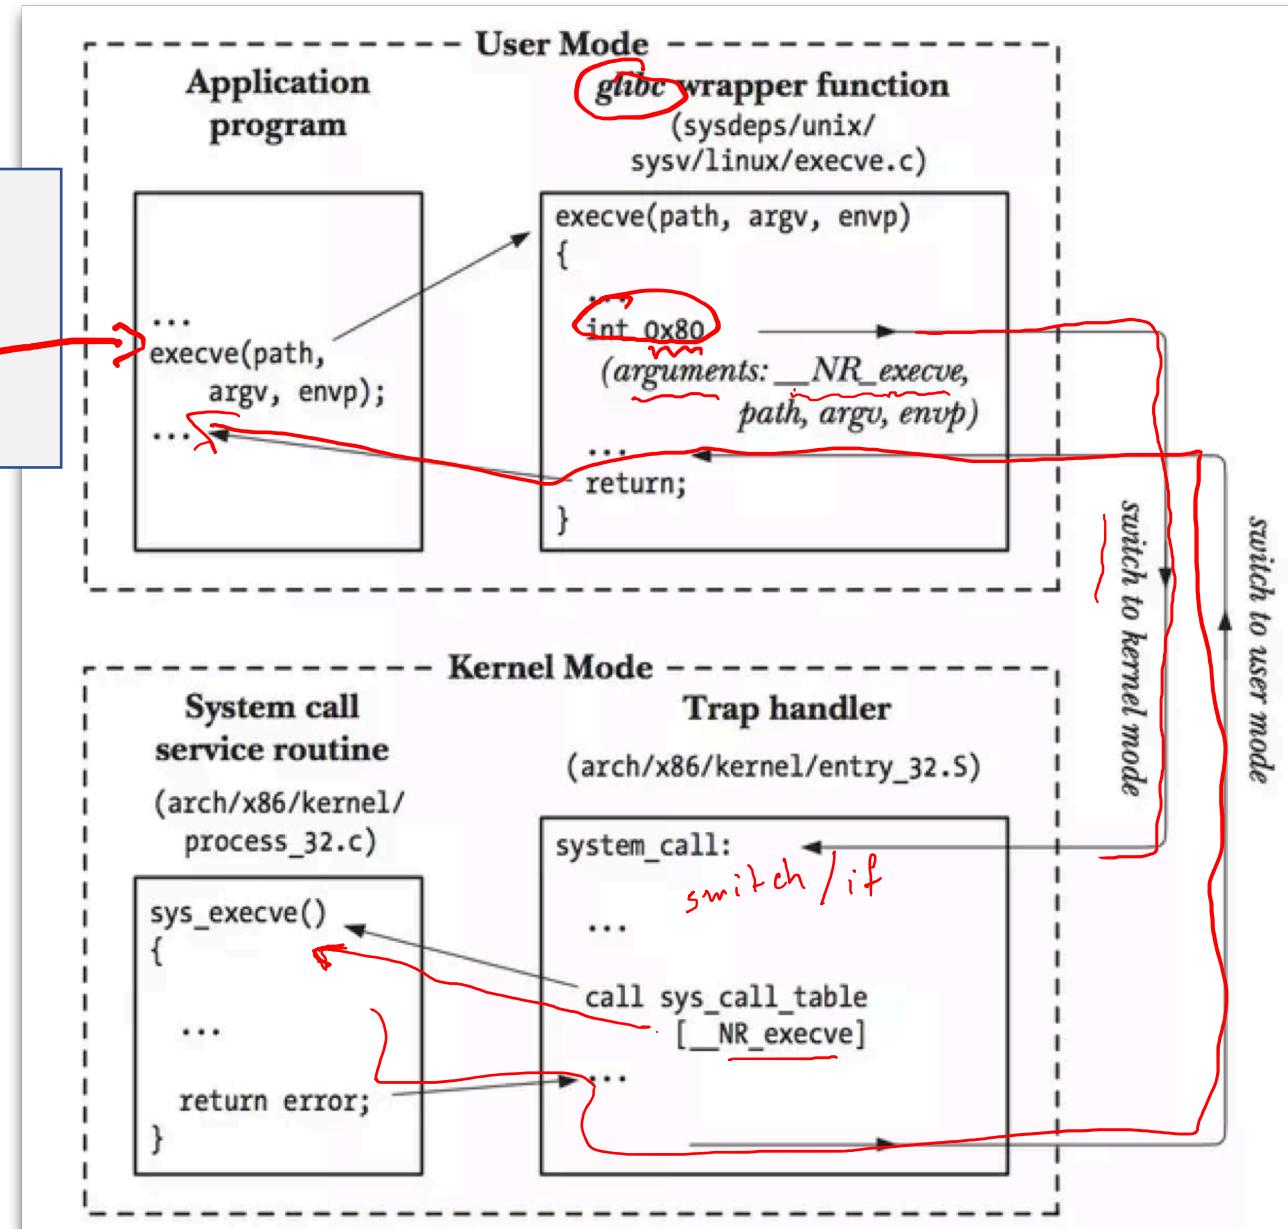
ECX

2nd Argument

EDX

3rd Argument

```
int main()
{
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
}
```



Syscalls

(cont.)
An example!

Registers & Invoking Syscalls

EAX

EBX

ECX

EDX

```
int main()
{
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
}
```

System Call Number

0x0000000b (11) – *the value for the execve syscall

1st Argument

address of “/bin/sh”

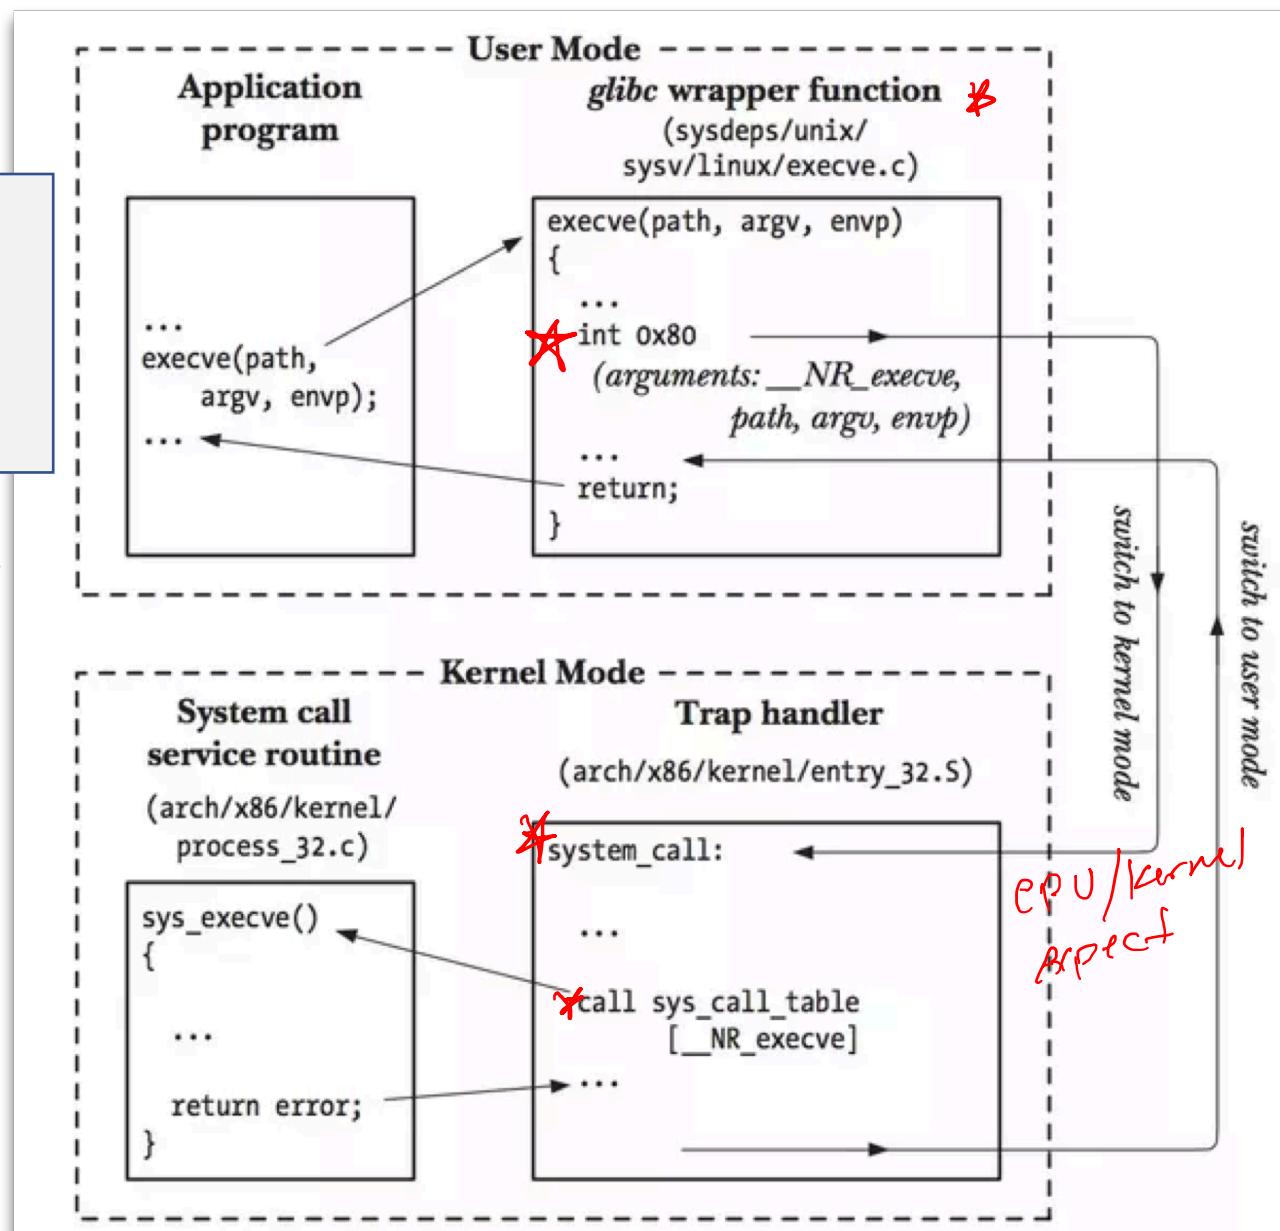
2nd Argument

0 – a.k.a. NULL - i.e., don't pass on environment variables

3rd Argument

INT 0x80 – trap to kernel and invoke
the syscall identified in the EAX register

*Example based on program compiled for 32-bit systems



Processes & Process Control

The *WHAT* and *WHY* (and a *little* of the *HOW*)

A (User) Process In Memory

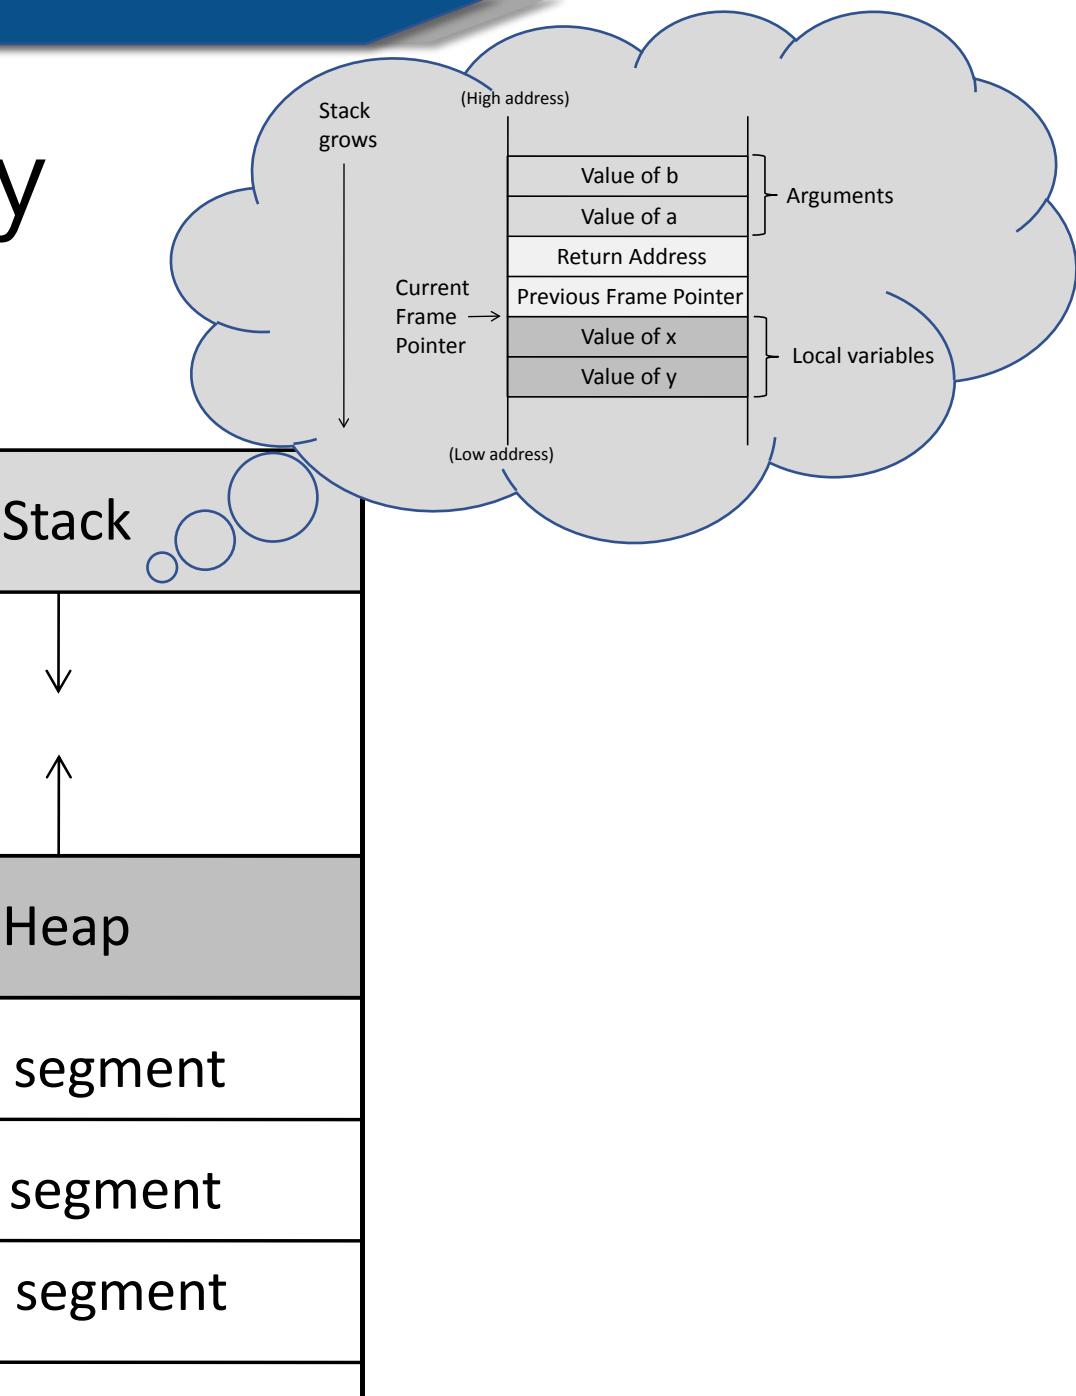
(Revisited)

The standard way to
set up a process in memory

env. vars.

(High address)

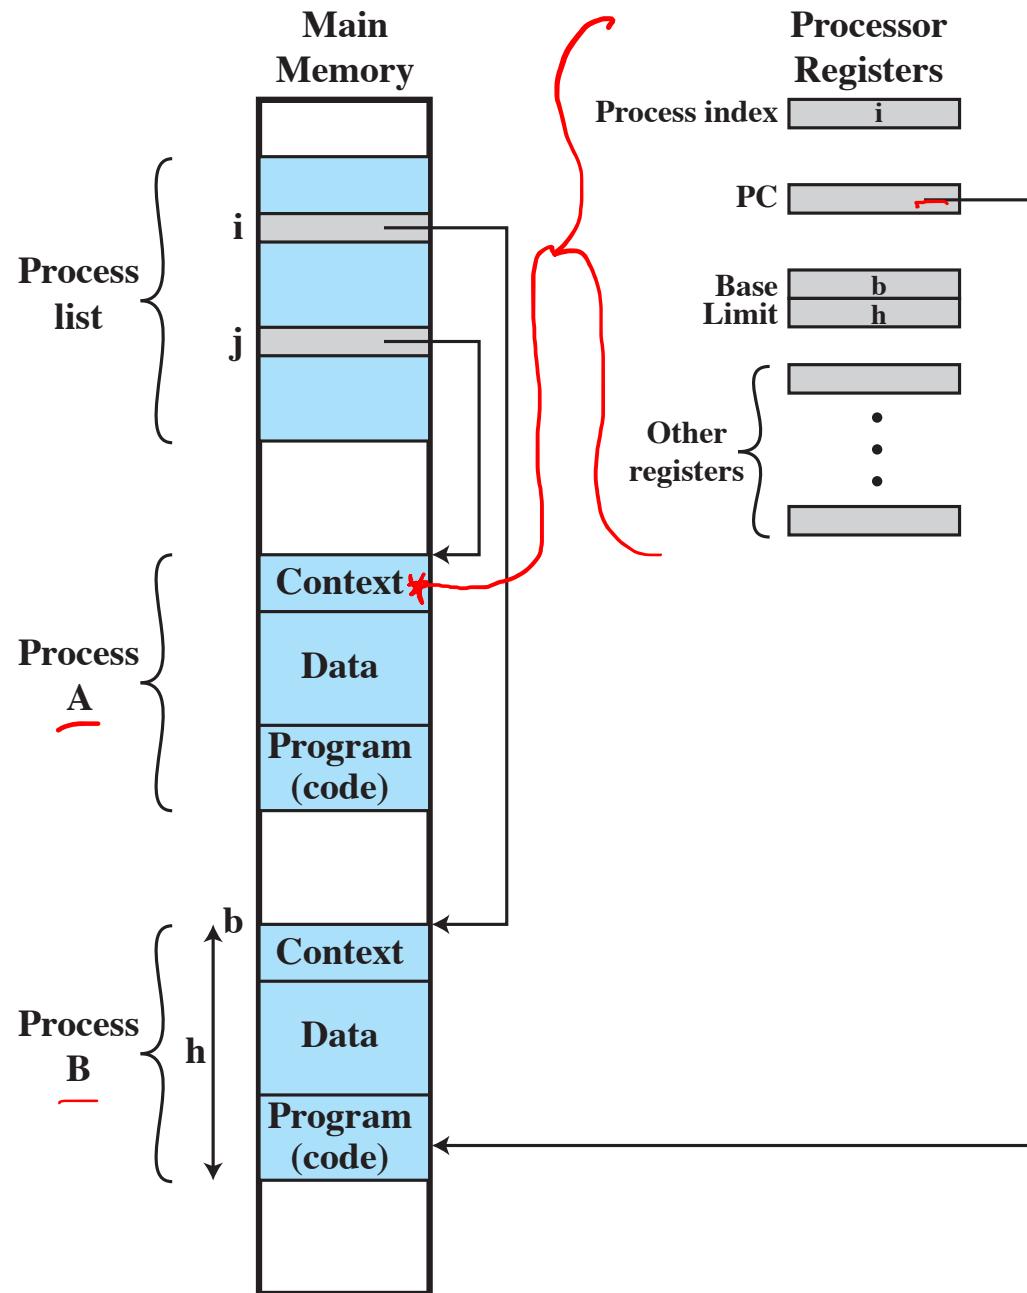
(Low address)



The Idea of a “Process”

(Revisited)

- An *instance of a program running on a computer*
- Consists of
 - (1) an executable program (**code**),
 - (2) associated **data**,
 - (3) **execution context**
(i.e., info the OS needs to manage the process)



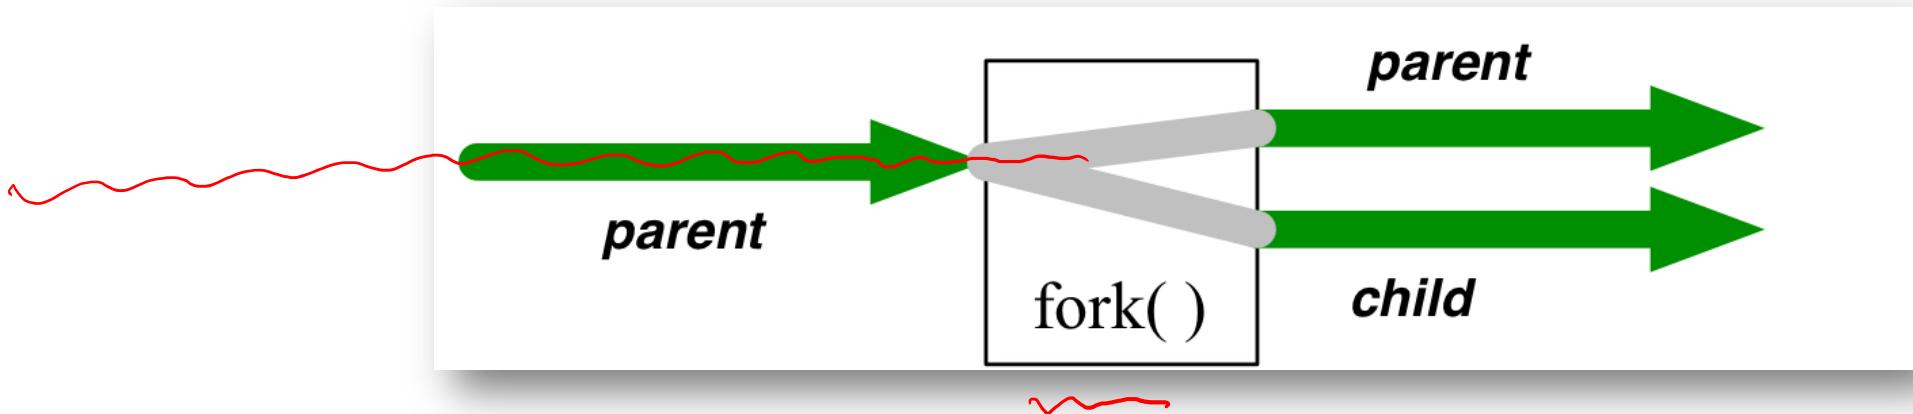
Processes & Process Control

The *HOW*: *create, control, identify, synchronize, terminate, etc.*

Creating a Process

(Who's Your Daddy? Or Mommy? Parent? Guardian?)

- In the Unix family (and others), we use **fork()** to create a new process
- **fork()** returns twice...
 - It creates a **child process** with an exact copy of the **parent process**.
 - same PC, same address space, ...
 - since PC is the same, both pick up executing at the same place (right after **fork()**)



Creating a Process

(Who's Your Daddy? Or Mommy? Parent? Guardian?)

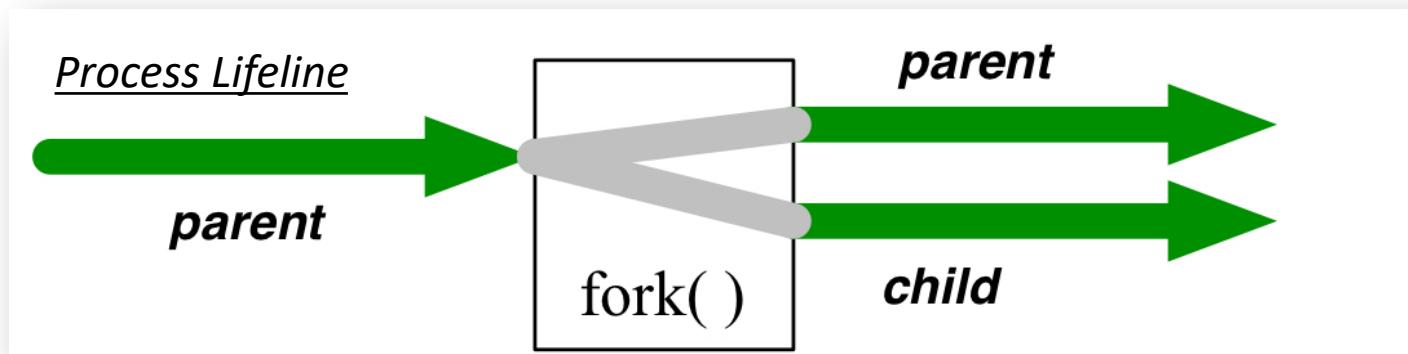
- In the Unix family (and others), we use **fork()** to create a new process
- **fork()** returns twice...
 - It creates a **child process** with an exact copy of the **parent process**.
 - same PC, same address space, ...
 - since PC is the same, both pick up executing at the same place (right after **fork()**)

Q: How do you tell the difference?

- A. You just know...
- B. You control this by setting the input arguments to fork()
- C. Check the return value of fork()
- D. Use another syscall to query the OS

Creating a Process (cont.)

- A. You just know...
- B. You control this by setting the input arguments to fork()
- C. Check the return value of fork()**
- D. Use another syscall to query the OS

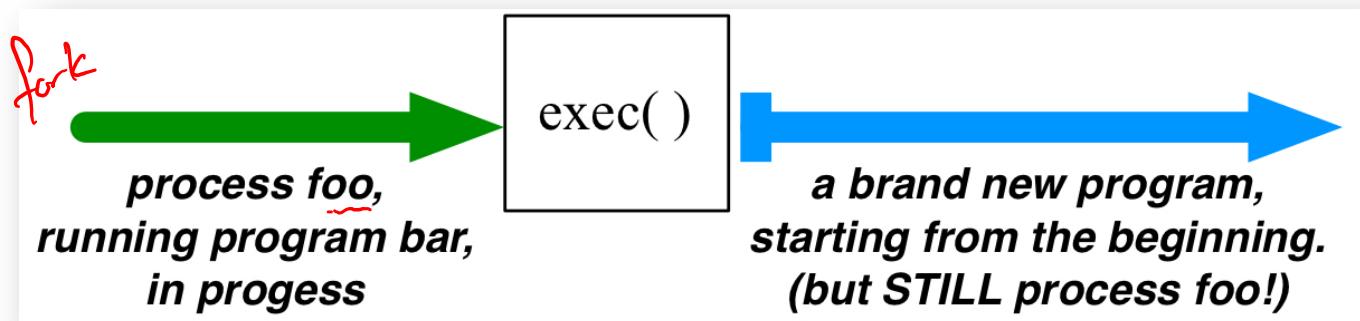


```
pid = fork();           pid > 0
if (0 == pid) {
    // I'm the child
    printf("Hi, I'm the child. Me no know how talk talk.\n");
    exit(0);
}
printf("I'm the parent. My child has pid %d\n", pid);
```

Changing Your Program

(But I don't want to be like my parents...)

- But what if I don't want my child to be EXACTLY like me...?
"But I want the child to run a different program!! What do I do?"
- The **exec** family of function calls
 - ...tell the OS to wipe out the address space in the calling process
 - ...+ repopulate with a new program and start running again from scratch → think *main()*

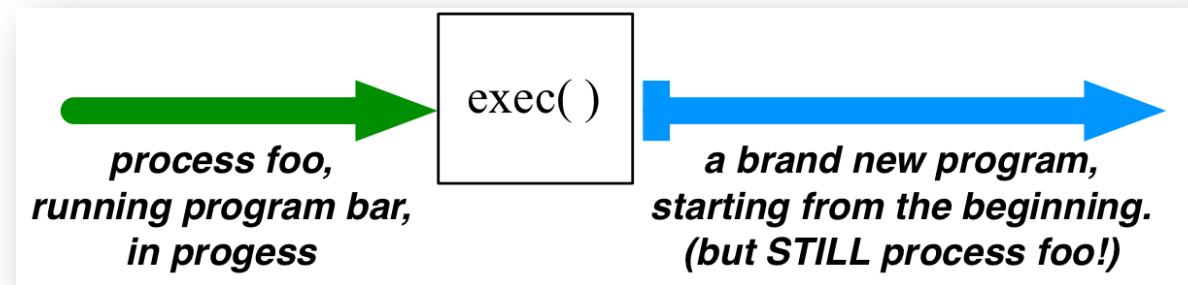


```
int main()
{
    char *name[2];
    name[0] = "/bin/sh";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
}
```

Changing Your Program

(But I don't want to be like my parents...)

- But what if I don't want my child to be EXACTLY like me...?
"But I want the child to run a different program!! What do I do?"



- There are a bunch of versions, varying on:
 - Do you specify the new program as a filename, or a full path?
 - Do you give the arguments as a list of NULL-terminated parameters, or as an argv[] array?
 - Do you want to bother specifying a new set of environment variables?
 - ...*lots of conventions & special semantics*

execP



Try `man exec`

...oops, that's the wrong answer.

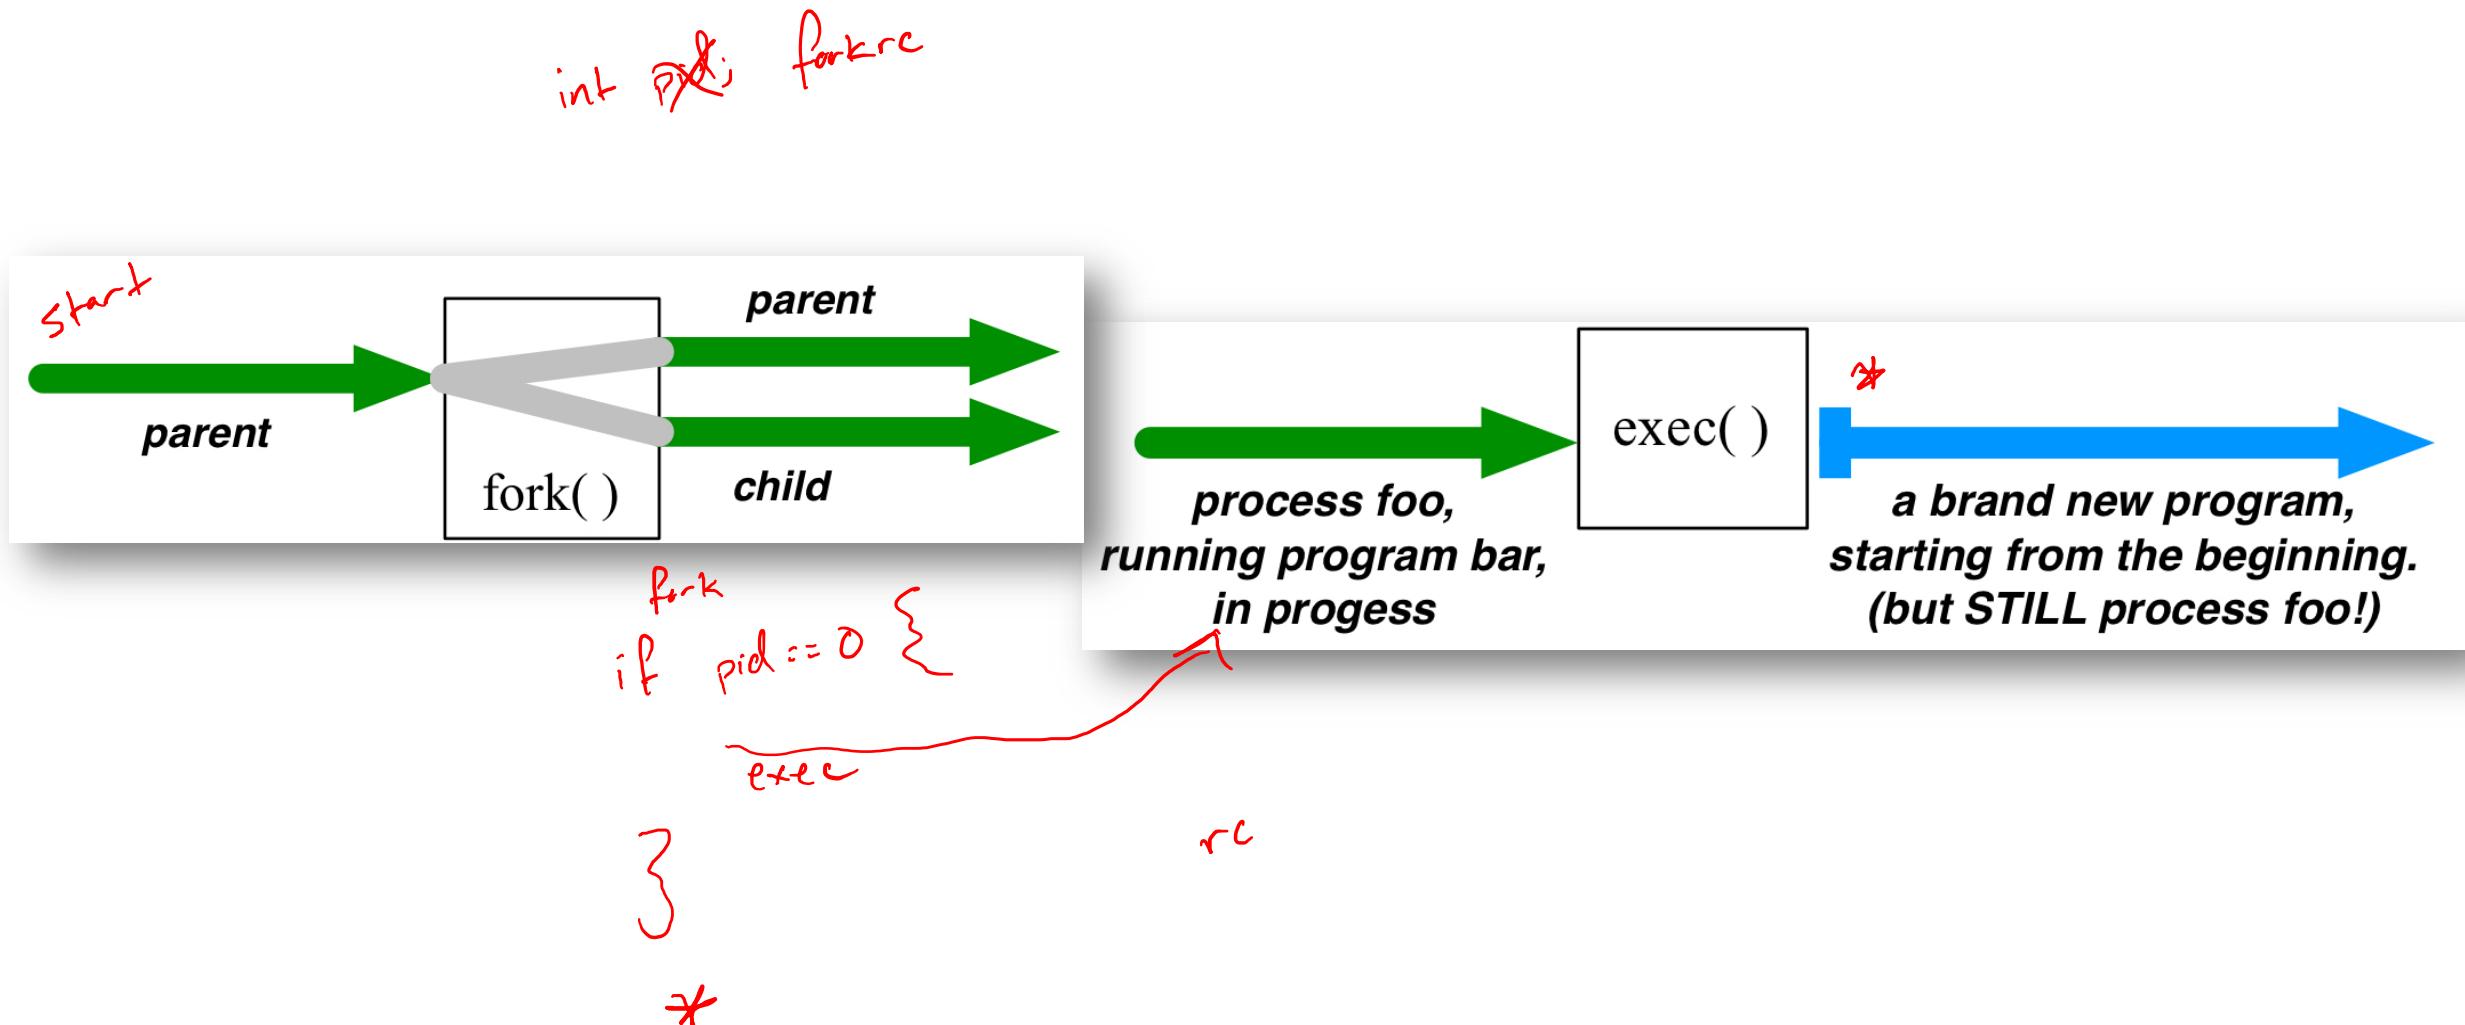
...try `man -k exec` → there's a bunch!

... try `man -S3 exec` (or `man 3 exec` or `man exec.3`)

... try `man execv*`

Changing Your Program

How we typically create a new process that runs a different program!



Process Identity

(And who exactly do you think you are?!)

A process can also explore its “family tree”.

- `getpid` gets your PID / this process's PID.
- `getppid` gets your parent's PID.
- The Init process is typically distinguished by `PID = 1`

[vagrant@vbox] [~/code/week04]\$ ps axif								
PPID	PID	PGID	SID	TTY	TPGID	STAT	UID	TIME COMMAND
0	2	0	0	?	-1	S	0	0:00 [kthreadd]
2	4	0	0	?	-1	I<	0	0:00 _ [kworker/0:0H]
0	1	1	1	?	-1	Ss	0	0:01 /sbin/init
1	418	418	418	?	-1	S<s	0	0:00 /lib/systemd/systemd-journald
1	425	425	425	?	-1	Ss	0	0:00 /sbin/lvmetad -f
1	437	437	437	?	-1	Ss	0	0:00 /lib/systemd/systemd-udevd
1	468	468	468	?	-1	Ss	0	0:00 /sbin/rpcbind -f -w
1	478	478	478	?	-1	Ss	100	0:00 /lib/systemd/systemd-networkd
1	514	514	514	?	-1	Ss	101	0:00 /lib/systemd/systemd-resolved
1	621	621	621	?	-1	Ssl	0	0:00 /usr/bin/lxcfs /var/lib/lxcfs/
1	624	624	624	?	-1	Ss	103	0:11 /usr/bin/dbus-daemon --system --address=
1	676	676	676	?	-1	Ssl	0	0:00 /usr/bin/python3 /usr/bin/networkd-dispa
1	677	677	677	?	-1	Ss	1	0:00 /usr/sbin/atd -f
1	678	678	678	?	-1	Ssl	102	0:00 /usr/sbin/rsyslogd -n
1	680	680	680	?	-1	Ssl	0	0:03 /usr/lib/accountsservice/accounts-daemon
1	684	684	684	?	-1	Ss	0	0:00 /usr/sbin/cron -f
1	691	691	691	?	-1	Ss	0	0:00 /lib/systemd/systemd-logind
1	763	763	763	?	-1	Ssl	0	0:00 /usr/lib/polkit-1/polkitd --no-debug
1	881	881	881	?	-1	Ss	0	0:00 /usr/sbin/sshd -D
881	1730	1730	1730	?	-1	Ss	0	0:00 _ sshd: vagrant [priv]
1730	1821	1730	1730	?	-1	S	1000	0:04 _ sshd: vagrant@pts/0
1821	1822	1822	1822	pts/0	12346	Ss	1000	0:01 _ -bash
1822	12346	12346	1822	pts/0	12346	R+	1000	0:00 _ ps axjf

Termination

(I'll be back.... No you won't.)

- Most often* a process terminates by...
 - `exit()`ing
 - “falling off” the end of `main`
 - receiving/catching a “terminate” **signal** (we’ll talk more about signals soon...)
 - signal == software interrupt
 - e.g., SIGTERM (15) - “please terminate...pretty please...”
 - being killed by the OS

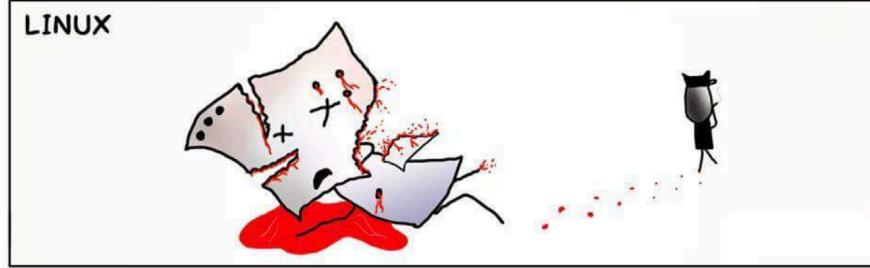
*Occasionally, some systems will enforce cascading termination: if a parent dies, its child must die too.

Name	Description	ISO C	SUS	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10	Default action
SIGABRT	abnormal termination (abort)	•	•	•	•	•	•	terminate+core
SIGALRM	timer expired (alarm)	•	•	•	•	•	•	terminate
SIGBUS	hardware fault	•	•	•	•	•	•	terminate+core
SIGCANCEL	threads library internal use							ignore
SIGCHLD	change in status of child	•	•	•	•	•	•	ignore
SIGCONT	continue stopped process	•	•	•	•	•	•	continue/ignore
SIGEMT	hardware fault							terminate+core
SIGFPE	arithmetic exception	•	•	•	•	•	•	terminate+core
SIGFREEZE	checkpoint freeze							ignore
SIGHUP	hangup	•	•	•	•	•	•	terminate
SIGILL	illegal instruction	•	•	•	•	•	•	terminate+core
SIGINFO	status request from keyboard							ignore
SIGINT	terminal interrupt character	•	•	•	•	•	•	terminate
SIGIO	asynchronous I/O							terminate/ignore
SIGIOT	hardware fault							terminate+core
SIGJVM1	Java virtual machine internal use							ignore
SIGJVM2	Java virtual machine internal use							ignore
SIGKILL	termination	•	•	•	•	•	•	terminate
SIGLOST	resource lost							terminate
SIGLWP	threads library internal use							terminate/ignore
SIGPIPE	write to pipe with no readers	•	•	•	•	•	•	terminate
SIGPOLL	pollable event (poll)							terminate
SIGPROF	profiling time alarm (setitimer)							terminate
SIGPWR	power fail/restart							terminate/ignore
SIGQUIT	terminal quit character	•	•	•	•	•	•	terminate+core
SIGSEGV	invalid memory reference	•	•	•	•	•	•	terminate+core
SIGSTKFLT	coprocessor stack fault							terminate
SIGSTOP	stop	•	•	•	•	•	•	stop process
SIGSYS	invalid system call	XSI	•	•	•	•	•	terminate+core
SIGTERM	termination	•	•	•	•	•	•	terminate
SIGTHAW	checkpoint thaw							ignore
SIGTHR	threads library internal use							terminate
SIGTRAP	hardware fault	XSI	•	•	•	•	•	terminate+core
SIGTSTP	terminal stop character	•	•	•	•	•	•	stop process
SIGTTIN	background read from control tty	•	•	•	•	•	•	stop process
SIGTTOU	background write to control tty	•	•	•	•	•	•	stop process
SIGURG	urgent condition (sockets)	•	•	•	•	•	•	ignore
SIGUSR1	user-defined signal	•	•	•	•	•	•	terminate
SIGUSR2	user-defined signal	•	•	•	•	•	•	terminate
SIGVTALRM	virtual time alarm (setitimer)	XSI	•	•	•	•	•	terminate
SIGWAITING	threads library internal use							ignore
SIGWINCH	terminal window size change	XSI	•	•	•	•	•	ignore
SIGXCPU	CPU limit exceeded (setrlimit)	XSI	•	•	•	•	•	terminate or terminate+core
SIGXFSZ	file size limit exceeded (setrlimit)	XSI	•	•	•	•	•	terminate or terminate+core
SIGXRES	resource control exceeded	XSI	•	•	•	•	•	ignore

Try:

kill -l

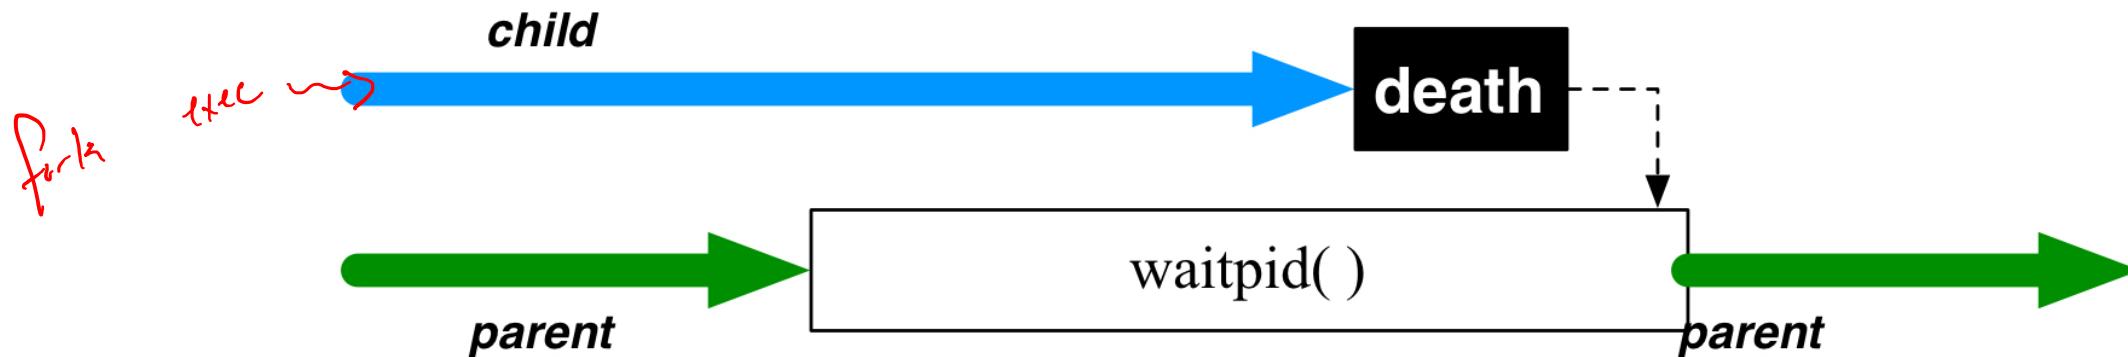
HANDLING NON-RESPONDING & FROZEN APPLICATIONS



Waiting

(Hey! Wait for meeeee)

- A process might want to check if another process has finished.
- The "classic" way of doing this is with **wait()**:
"please block me until the child process I named terminates, and then tell me its exit code."



- In Linux, **waitpid()** is maybe the simplest way to do this
 - *the parent can specify exactly which child (PID) it is waiting for.*

Demos

- Go and try out week04/ demo code
 - fork
 - exec
 - waitpid
 - kill
 - getpid
 - + more!

