

Operating Systems!

Concurrency: Race Conditions

Prof. Travis Peters

Montana State University

CS 460 - Operating Systems

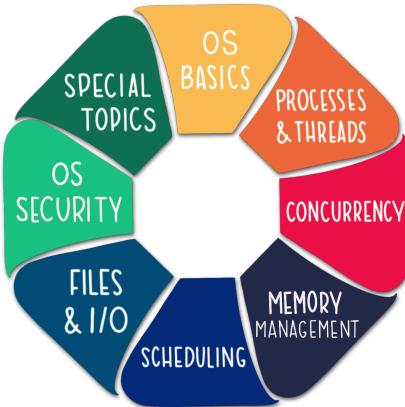
Fall 2020

<https://www.cs.montana.edu/cs460>

Some diagrams and notes used in this slide deck have been adapted from Sean Smith's OS courses @ Dartmouth. Thanks, Sean!

Today

Travis has already helped a few "Free Agents" form team members! Great to have a group to discuss things (e.g., Accountability, PA2, Studying, Yalnix CP1)



- Announcements

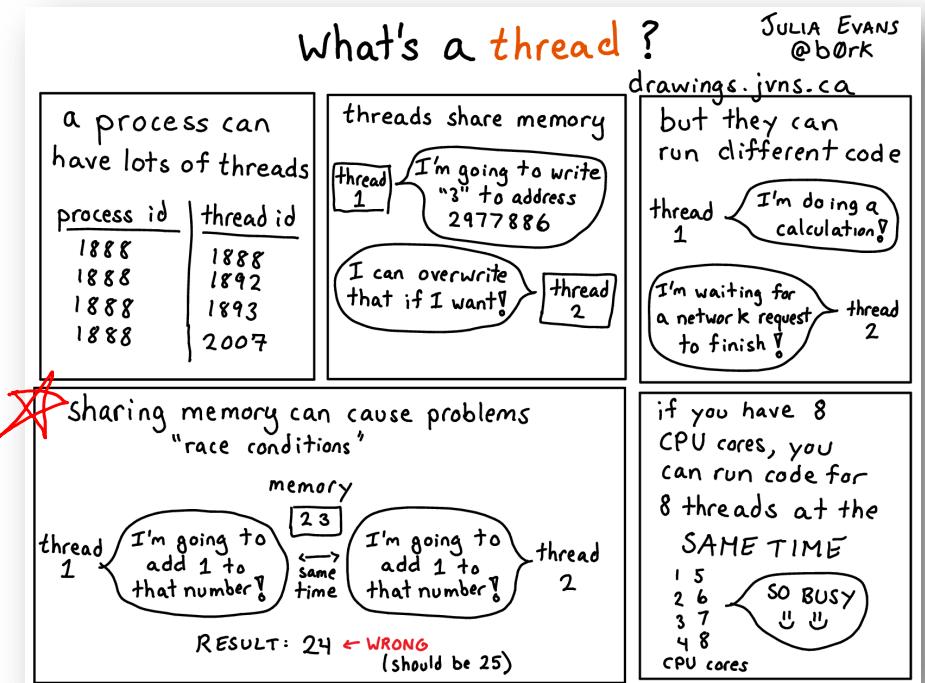
Read over the exam coversheet!

- Heads up.... Exam 1: Friday [10/02/2020] – due @ 11:59 PM (MST)
- Heads up.... PA2 due Sunday [10/04/2020] @ 11:59 PM (MST)
- Heads up.... Yalnix! Form teams ASAP! ☺

Check out the PA2 starter files!

- Learning Objectives

- Threads & Interleaving
- What are "Race Conditions"?
- Examples
- Towards a Solution



Last Time...

- Threads “below the surface”
 - userlevel threads vs. kernellevel threads
- Threads “above the surface” (**pthreads**)
 - `pthread_t mythread;`
 - `pthread_create(...);`
 - `pthread_join(...);`
 - `pthread_cancel(...);`
 - ...

*similar
(not the same)*

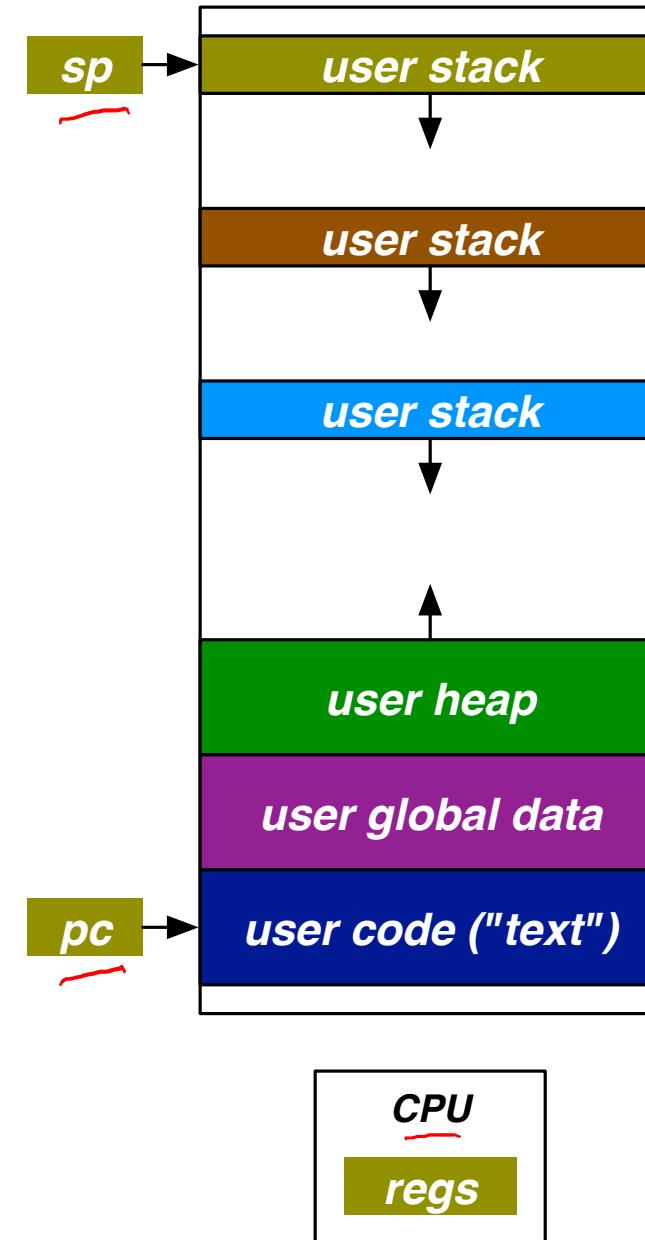
fork

work

exit/abort/kill

thread-related demos:

- `threads0.c`
- `threads1.c`
- `threads2.c`
- `thread_spy.c` -IPC
- `thread_atm.c`
- `ipt.c`



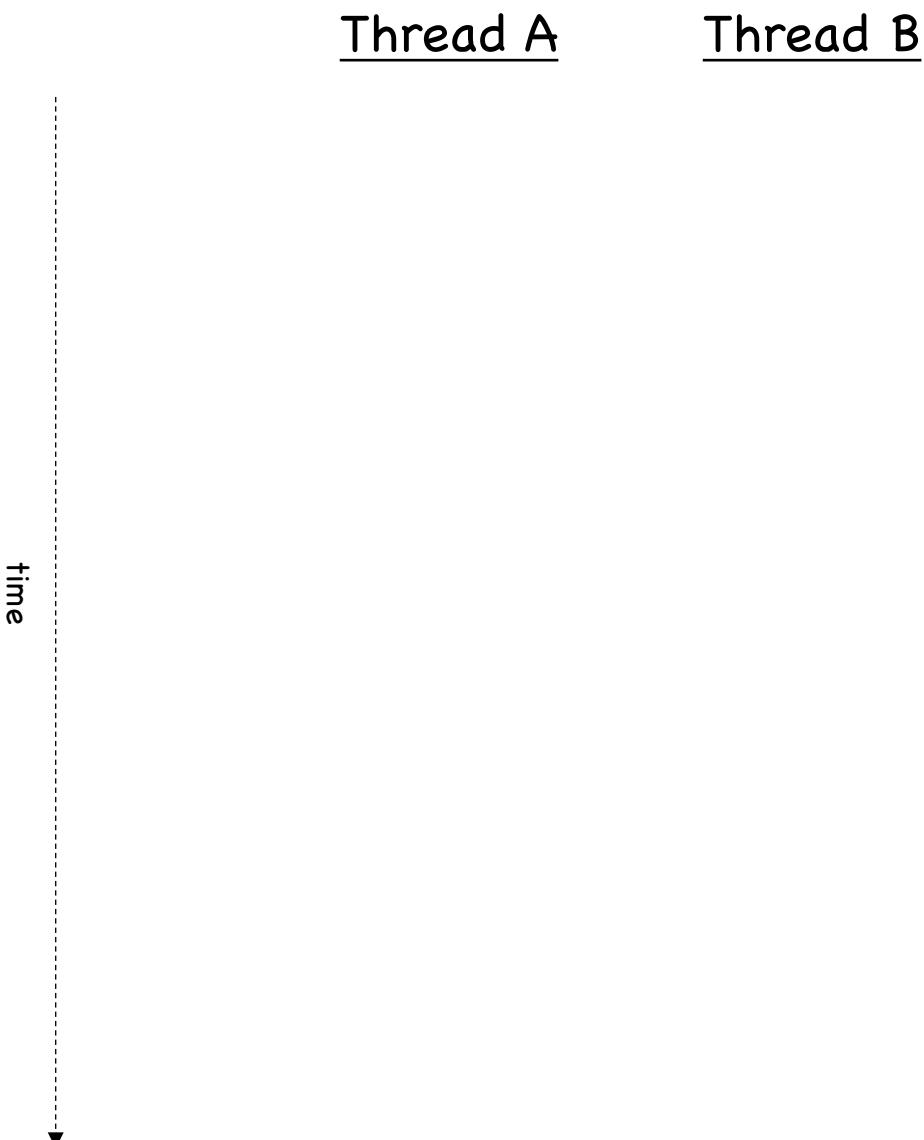
Concurrency – what is it? why is it?

- **Concurrency** is all about communicating & sharing whilst **interleaving** and **overlapping** execution.
 - Note: We'll mostly focus on uniprocessor systems
 - Recall: A **thread** is a single (separately schedulable) unit of execution
- Examples
 - Servers → managing multiple connections
 - Programs w/ UIs → responsiveness
 - Network-/Disk-Bound Programs → hide network/disk latency

Threads & Interleaving

→ See `ipt.c`

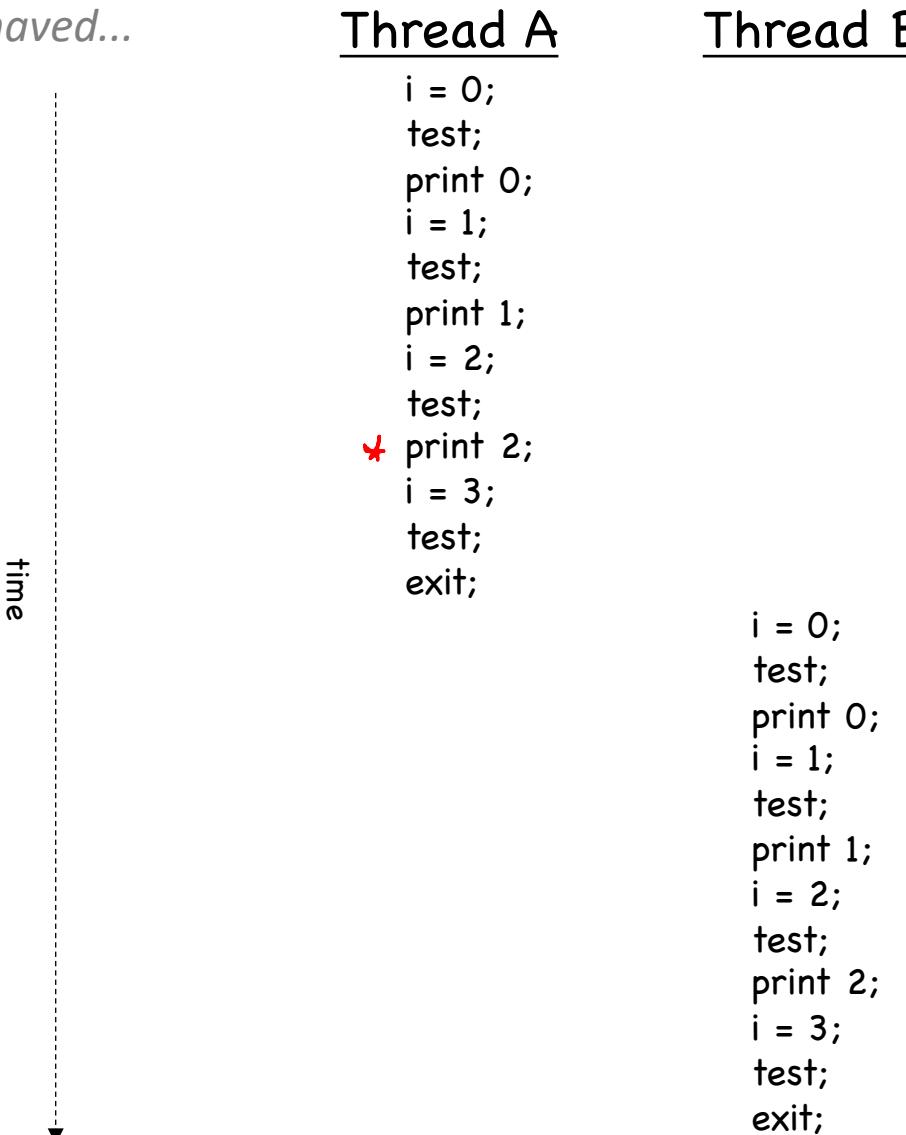
A Simple Threaded Program



Consider code:
int i; //global variable
...
func {
 for (i = 0; i < 3; i++) {
 printf("%d\n", i);
 }
}

A Simple Threaded Program

Maybe they will be well-behaved...

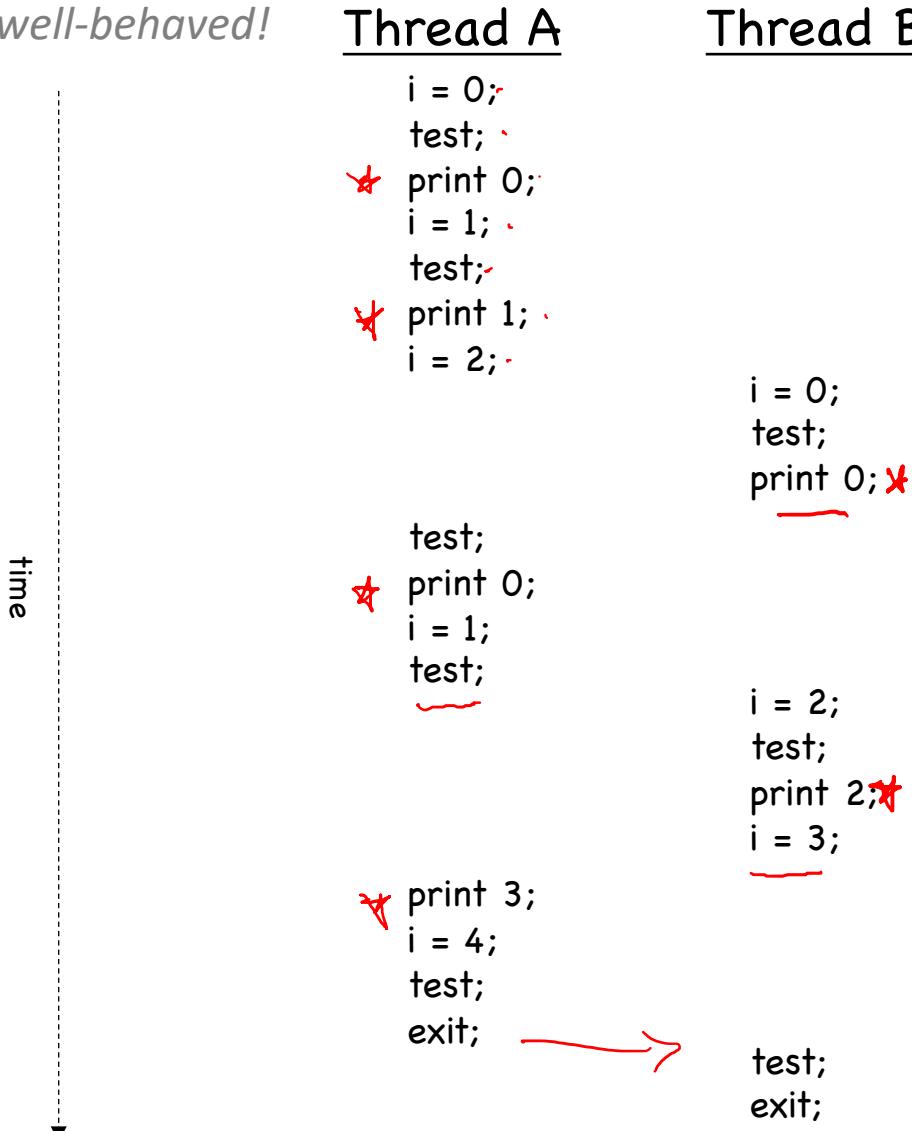


Consider code:
int i; //global variable
...
for (*i* = 0; *i* < 3; *i*++) {
 printf("%d\n", *i*);
}

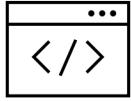
A Simple Threaded Program

...but maybe they won't be well-behaved!

ipx.c



Consider code:
int i; //global variable
...
for (i = 0; i < 3; i++) {
 printf("%d\n", i);
}

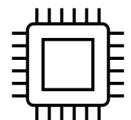


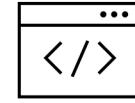
Why Does This Happen?

What happens when a program executes?

- High-level code
 - sequence of assembly instructions
- Execute instructions one *instruction* at a time
- Can be *interrupted / pre-empted at any time!*

Compiler

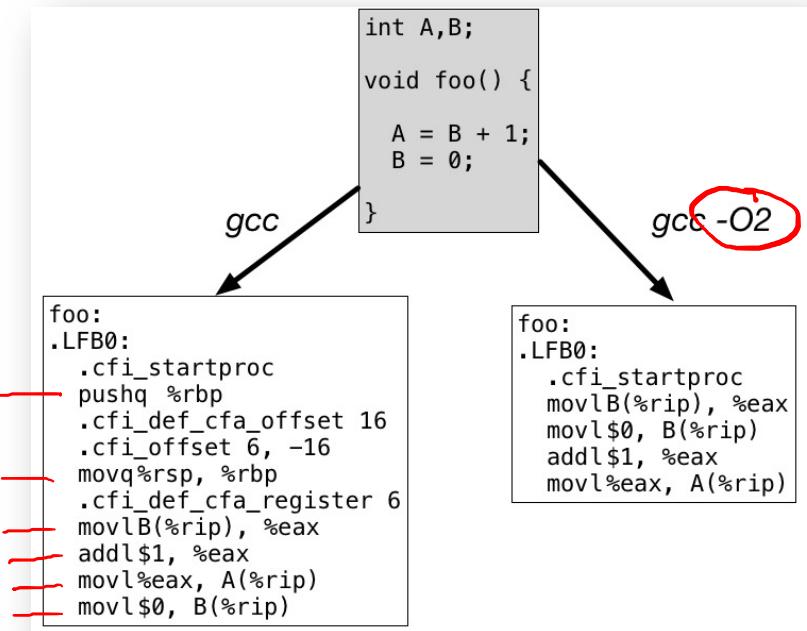




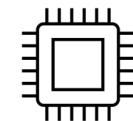
Why Does This Happen?

What happens when a program executes?

- High-level code
→ sequence of assembly instructions
- Execute instructions one *instruction* at a time
- Can be interrupted / pre-empted at any time!



<https://preshing.com/20120625/memory-ordering-at-compile-time/>



Also: compiler re-ordering...

...lots of non-intuitive things, all in the name of better performance!

Programmer vs. Processor View

programmer view:

Programmer's View

```
x = x + 1;
y = y + x;
z = x + 5y;
```

Possible Execution #1

```
x = x + 1;
y = y + x;
z = x + 5y;
```

Possible Execution #2

x = x + 1;
.....
Thread is suspended.
Other thread(s) run.
Thread is resumed.
.....
y = y + x;
z = x + 5y;

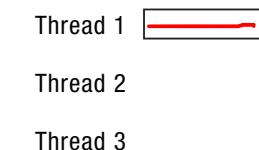
Possible Execution #3

x = x + 1;
y = y + x;
.....
Thread is suspended.
Other thread(s) run.
Thread is resumed.
.....
z = x + 5y;

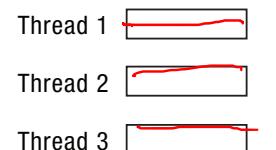


processor view:

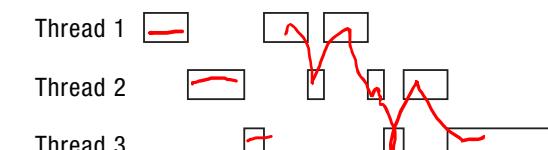
One Execution



Another Execution



Another Execution

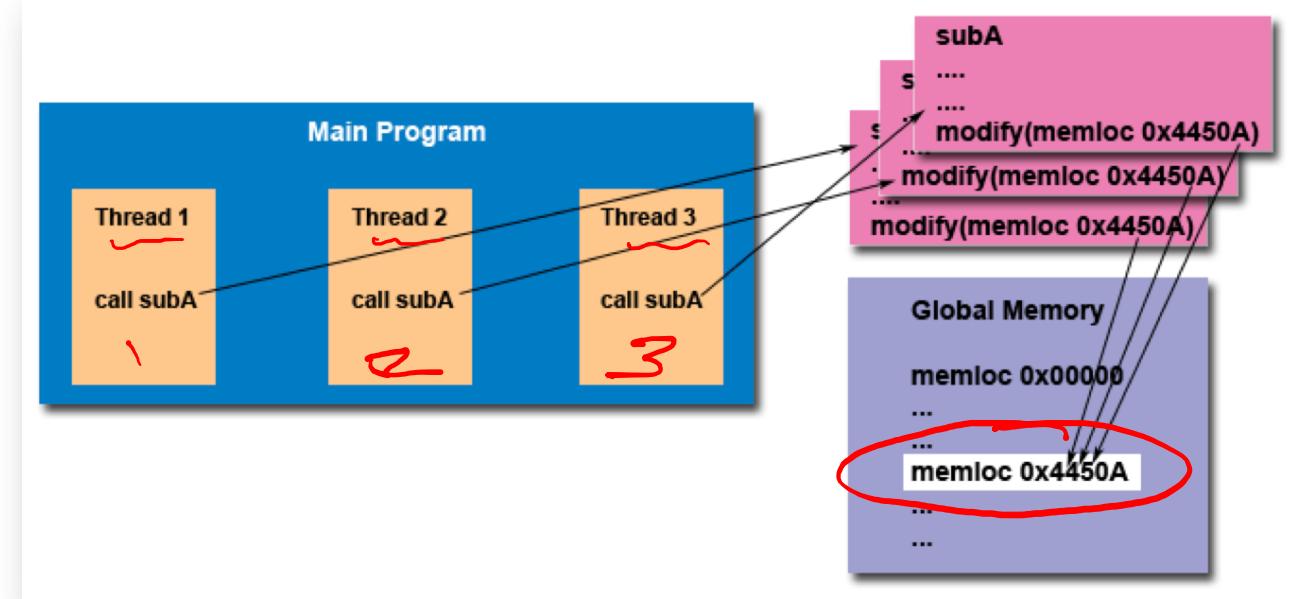


Programmer/program must anticipate all of these possible executions!

Race Conditions

Race Conditions

- What does it mean for a computation to be “correct”?



<https://computing.llnl.gov/tutorials/pthreads/>

- **Race Condition:** when not all observable outcomes are correct.
→ Whether the program works depends on a “race” between concurrent threads or processes.

*Put another way: when **shared data** is corrupted due to **bad interleavings / thread scheduling***

Race Conditions

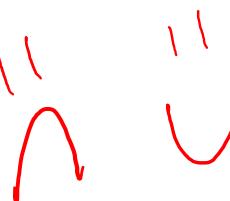
Banking Example

- Alice and Bob share a bank account w/ \$100.
- Both want to make withdrawals...
 - Alice asks for \$90
 - Bob asks for \$80

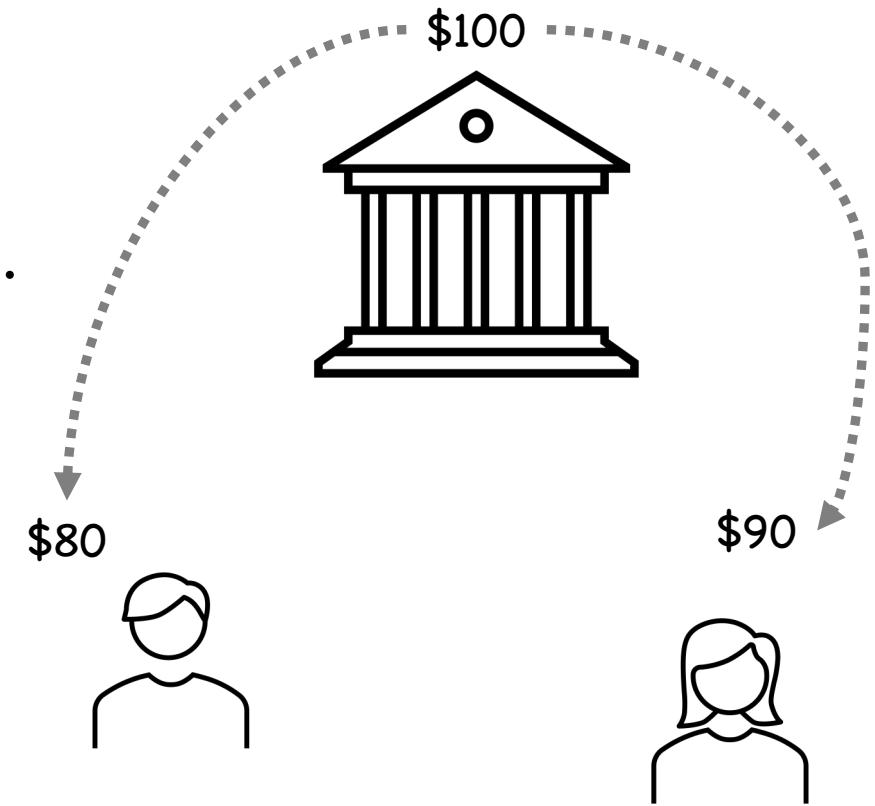
Q: What are correct outcomes?

Alice gets \$90
Bob gets \$80

Q: What might actually happen?



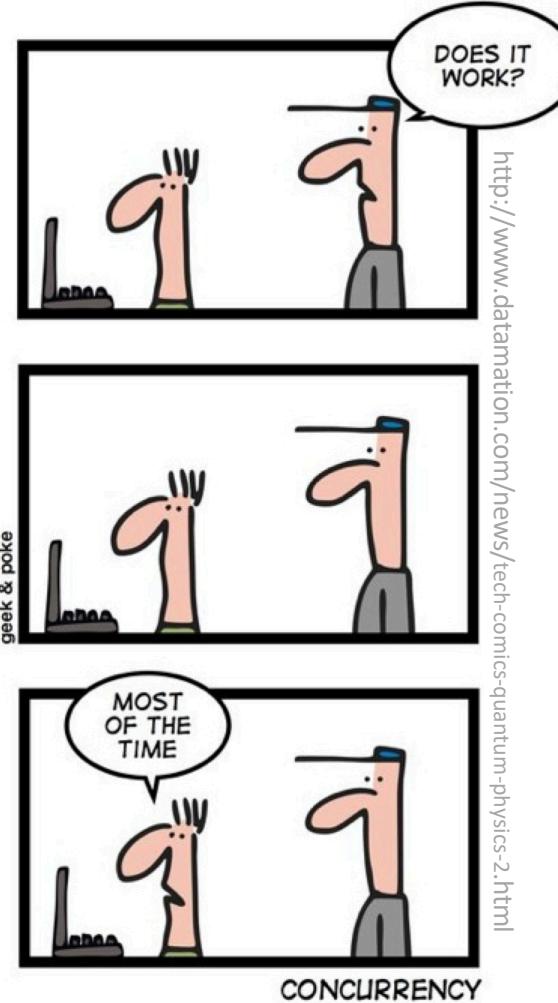
→ See `thread_atm.c`



Debugging Race Conditions

Race conditions are tricky...

- It may be that only some interleavings are problematic
 - Code works fine (most of the time...)
 - try to debug... can't recreate the problem...
- External factors may change interleavings
 - different systems?
 - using a debugger? ("heisenbugs")



Practical Manifestations

When and where should you look for race conditions in your code? Here are some examples...

- When you have **shared data** *
- When you have a data structure with multiple fields
→ changes take several steps
- When you test something, *then* act on the result of the test
* → **TOCTOU** ("Time of Check vs. Time of Use")
- When the action of one thread requires a prior action by another thread
- When you want to wait until two different things are both true

Towards a Solution

Critical Section

How could we
“guard”
a critical section?



A blue rectangular box labeled "Critical Section" with a red asterisk at the top right corner.

Critical Section

```
while(flag) {}
```

```
flag = TRUE;
```

Critical Section

```
flag = FALSE;
```

```
while(flag) {}
```

```
flag = TRUE;
```

Critical Section

```
flag = FALSE;
```

```
while(flag) {}
```

```
flag = TRUE;
```

Critical Section

```
flag = FALSE;
```

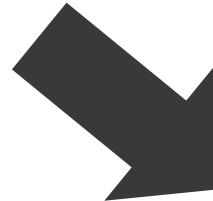
```
while(flag) {}  
flag = TRUE;
```

Critical Section

```
flag = FALSE;
```

} *SW TEST&SET*

Atomic “Test and Set” Instruction



```
while(test-and-set(flag)) {}
```

Critical Section

```
flag = FALSE;
```

Thread Synchronization

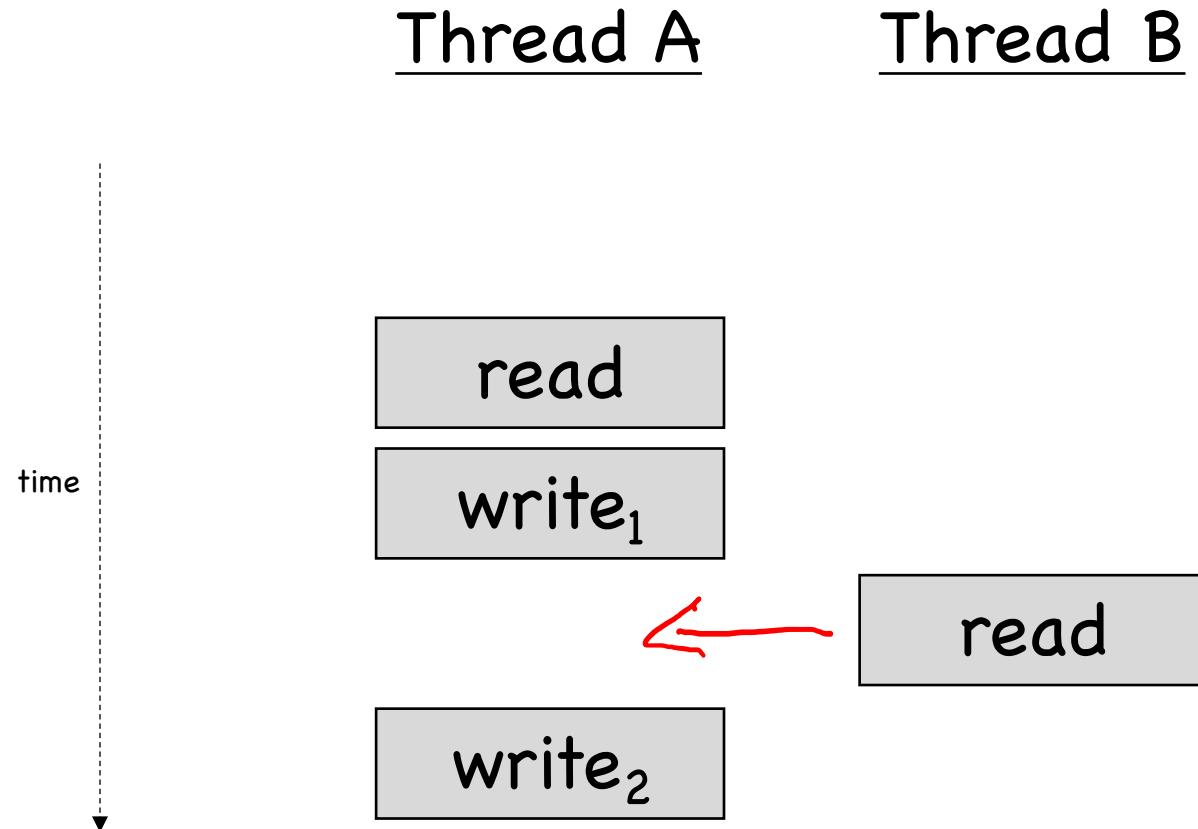
Another Example:

Suppose 2 threads want to access the same variable...

Thread A Thread B



Thread Synchronization

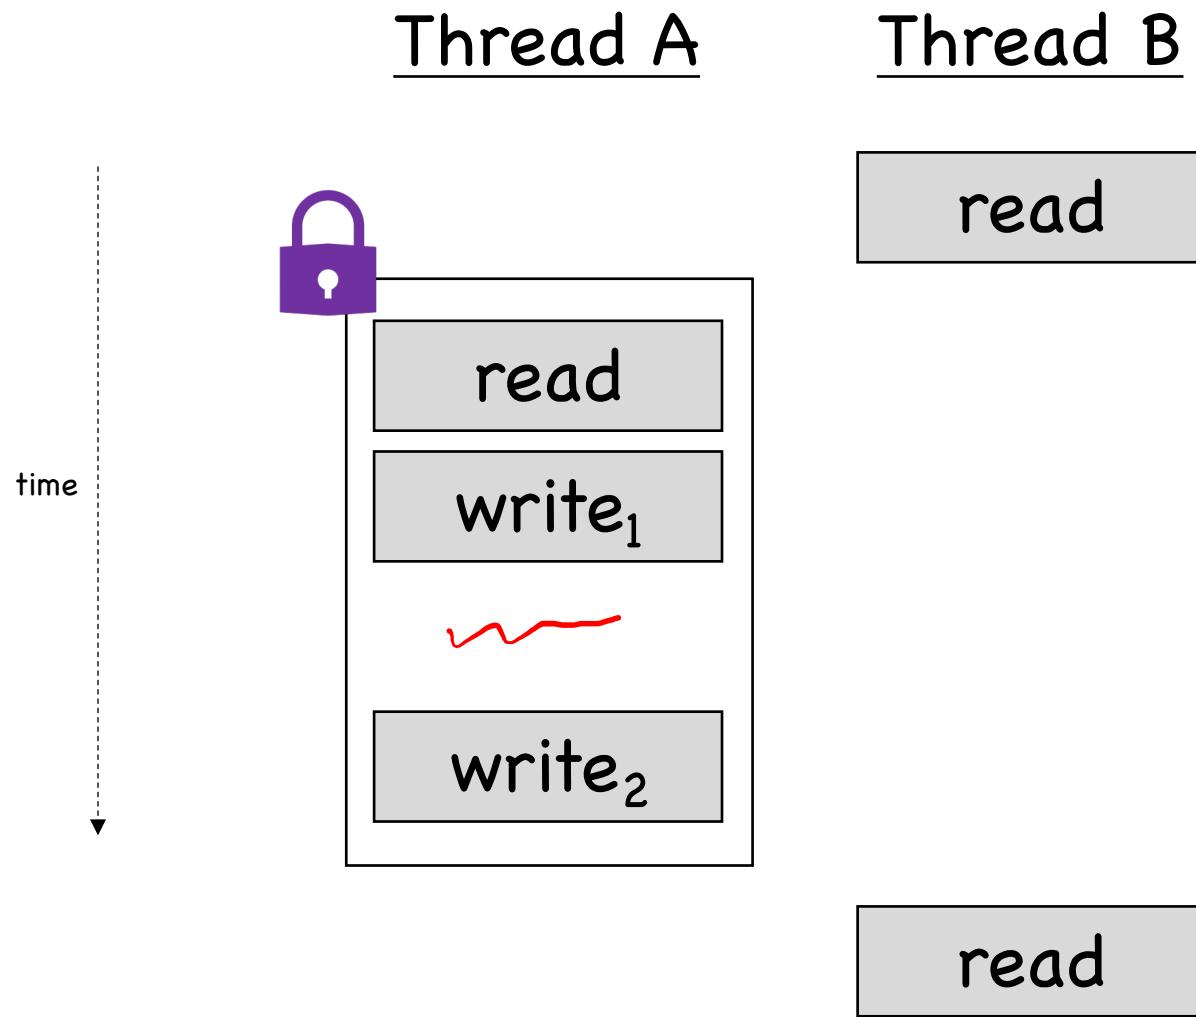


Another Example:

Suppose 2 threads want to access the same variable...

Don't want Thread B to read *while* Thread A is writing...

Thread Synchronization



Another Example:

Suppose 2 threads want to access the same variable...

"Protect" Thread A's multi-step write operation.

(prevents simultaneous read/write of the shared variable)