

Operating Systems!

Memory: Approaches to Memory Management (Part 4)

Prof. Travis Peters

Montana State University

CS 460 - Operating Systems

Fall 2020

<https://www.cs.montana.edu/cs460>

Some diagrams and notes used in this slide deck have been adapted from Sean Smith's OS courses @ Dartmouth. Thanks, Sean!

Today

- Announcements
 - ~~Yalnix~~ How are projects/proposals going?! 😊
 - Exam 1: Nice work! (*Still*) Working on grading...
- Upcoming Deadlines
 - **Sunday [10/18/2020] @ 11:59 PM (MST)** - PA3 (Group Assignment)
Yalnix teams: While PA3 == yalnix checkpoint 1, you should plan to be near completing checkpoint 2 at this point!
 - **Sunday [10/25/2020] @ 11:59 PM (MST)** - Project Proposal
Yalnix teams: you should plan to be near completing checkpoint 3 at this point!



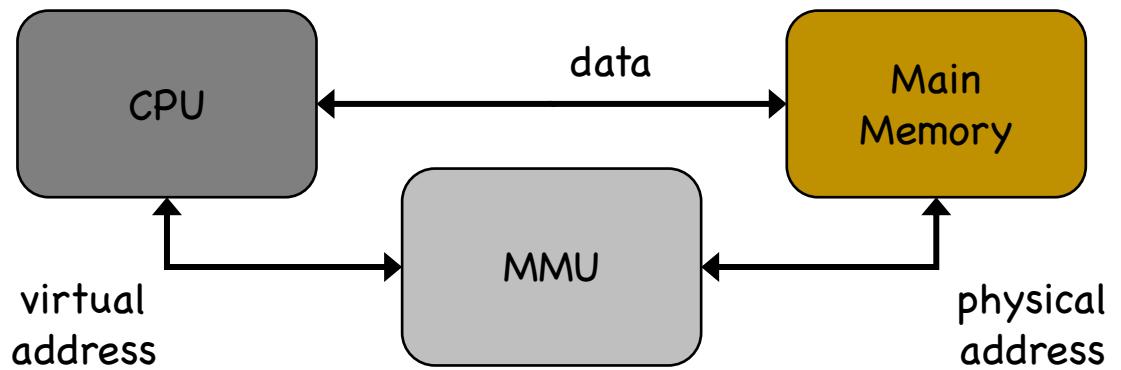


Today (cont.)

Last Week: Logical vs. Physical Addresses,
Paging, Address Translation, ...

- Agenda
 - ~~Page Tables~~
 - ~~Heaps~~ → review fault.c 3,4
 - Stacks → (quickly) review fault.c 5,6,7 + sf.c
 - Page Tables (Practical Issues)
 - The TLB
 - Switching Processes & the TLB

→ Virtual Memory

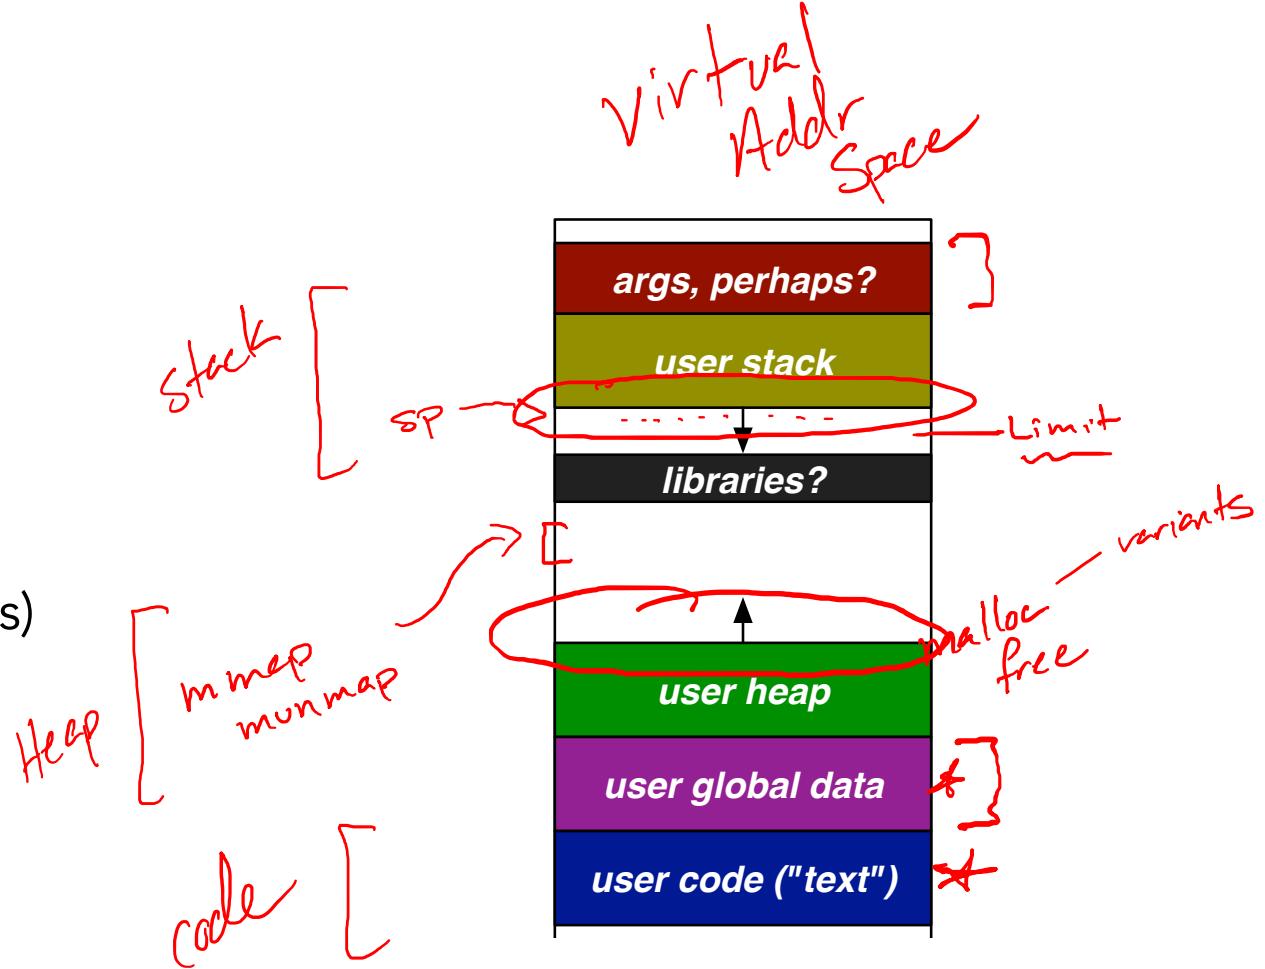


The Stack

Stacks of Mystery

Bottom of the Stack

- Stacks “grow” in a different way...
→ **implicit growth!**
- Basic idea:
 - if access is within max limit, allocate page(s)
 - else send SIGSEGV (which typically kills it)
- How? Page Faults!
 - **major fault**
 - vs.
 - **minor fault**



See:
SE posts on stack allocation in Linux
<https://unix.stackexchange.com/a/239323>
<https://unix.stackexchange.com/a/280865>

see: **fault.c (5,6,7)**

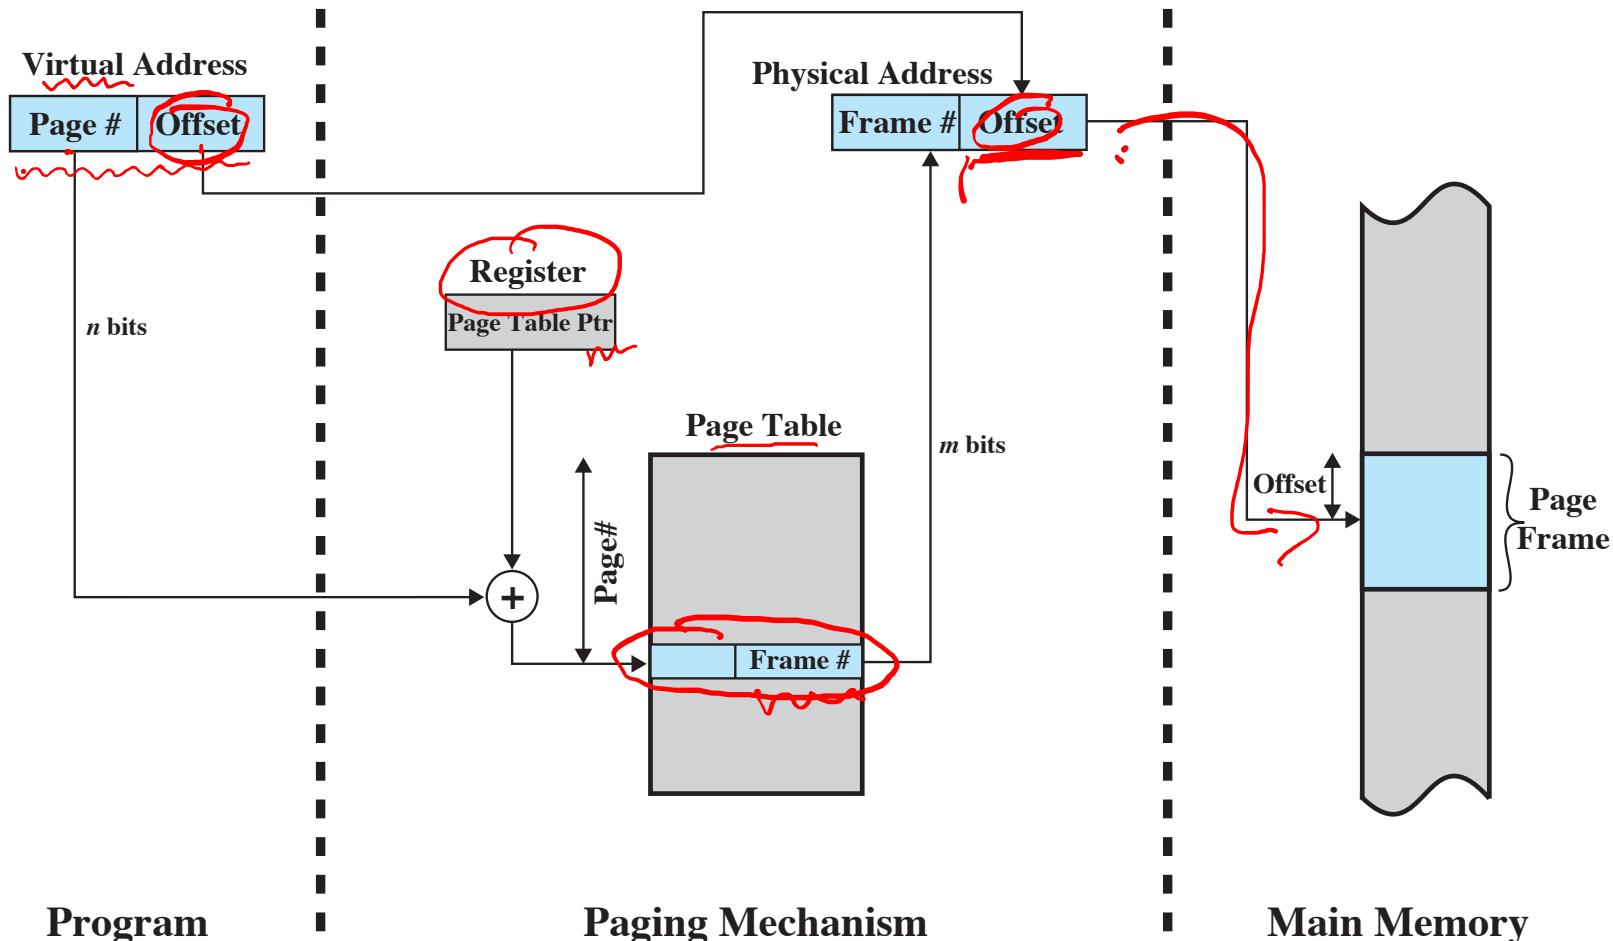
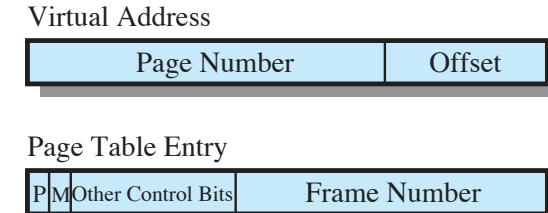
(Why don't we see syscalls when running strace...)

Page Tables

Practical Issues

mapping
 $p \rightarrow f$

Address Translation: 1-Level Paging System



How big are these page tables?

Activity! (Think-Pair-Share)

- Suppose...

- Virtual Memory = 4GB (2^{32})
- Page Size = 4KB (2^{12})
- PTE = 4B (2^2)

4GB RAM

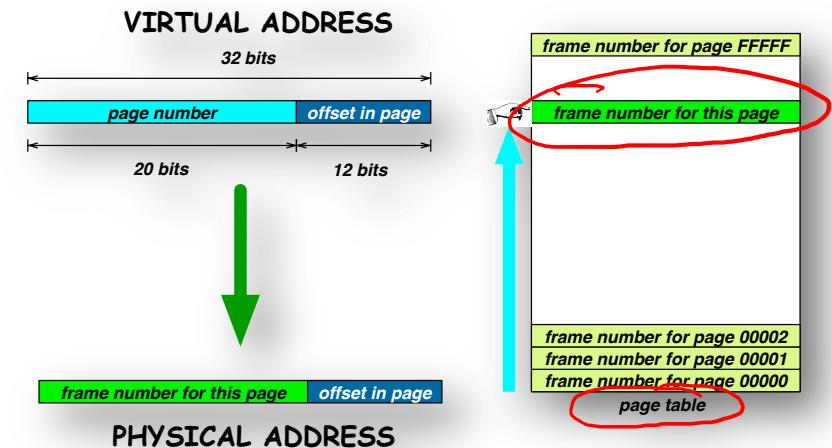
- How many page table entries? $2^{32} / 2^{12} = 2^{20}$

- How much memory is needed to store such a page table?

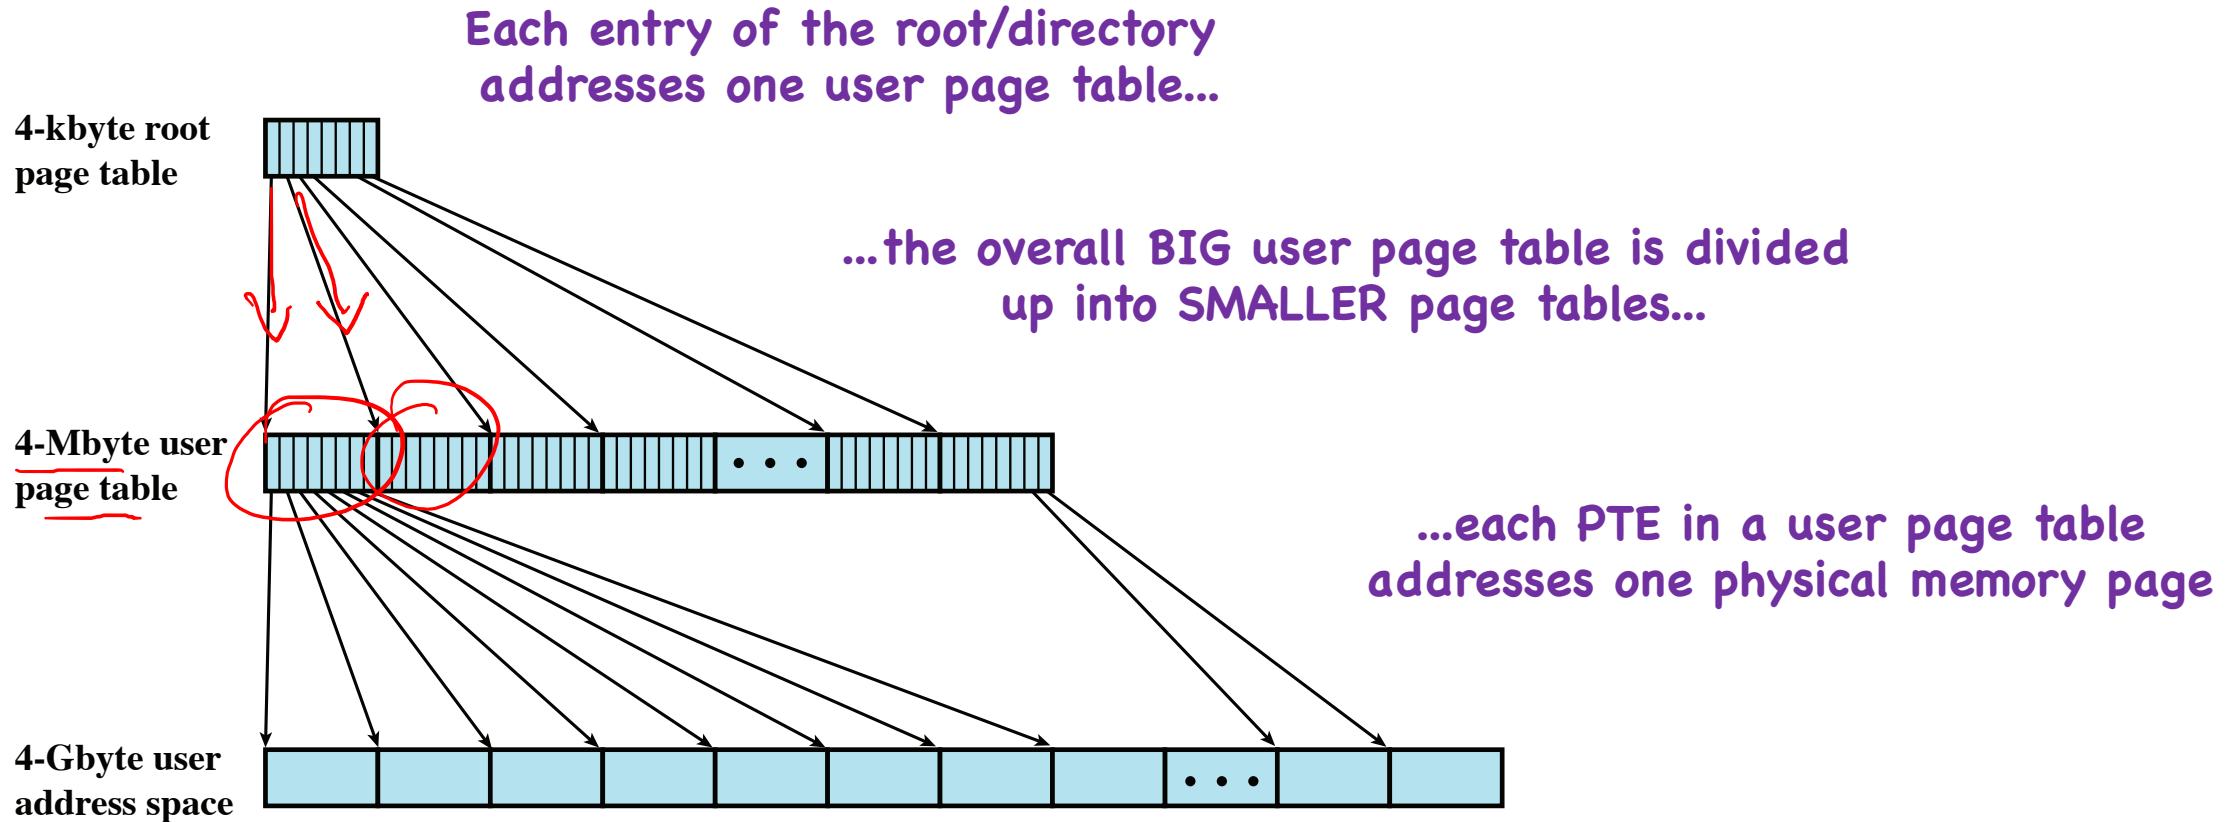
$$2^{20} \cdot 2^2 = 2^{22} \text{ (4MB)}$$

- How many pages are needed to hold the page table?

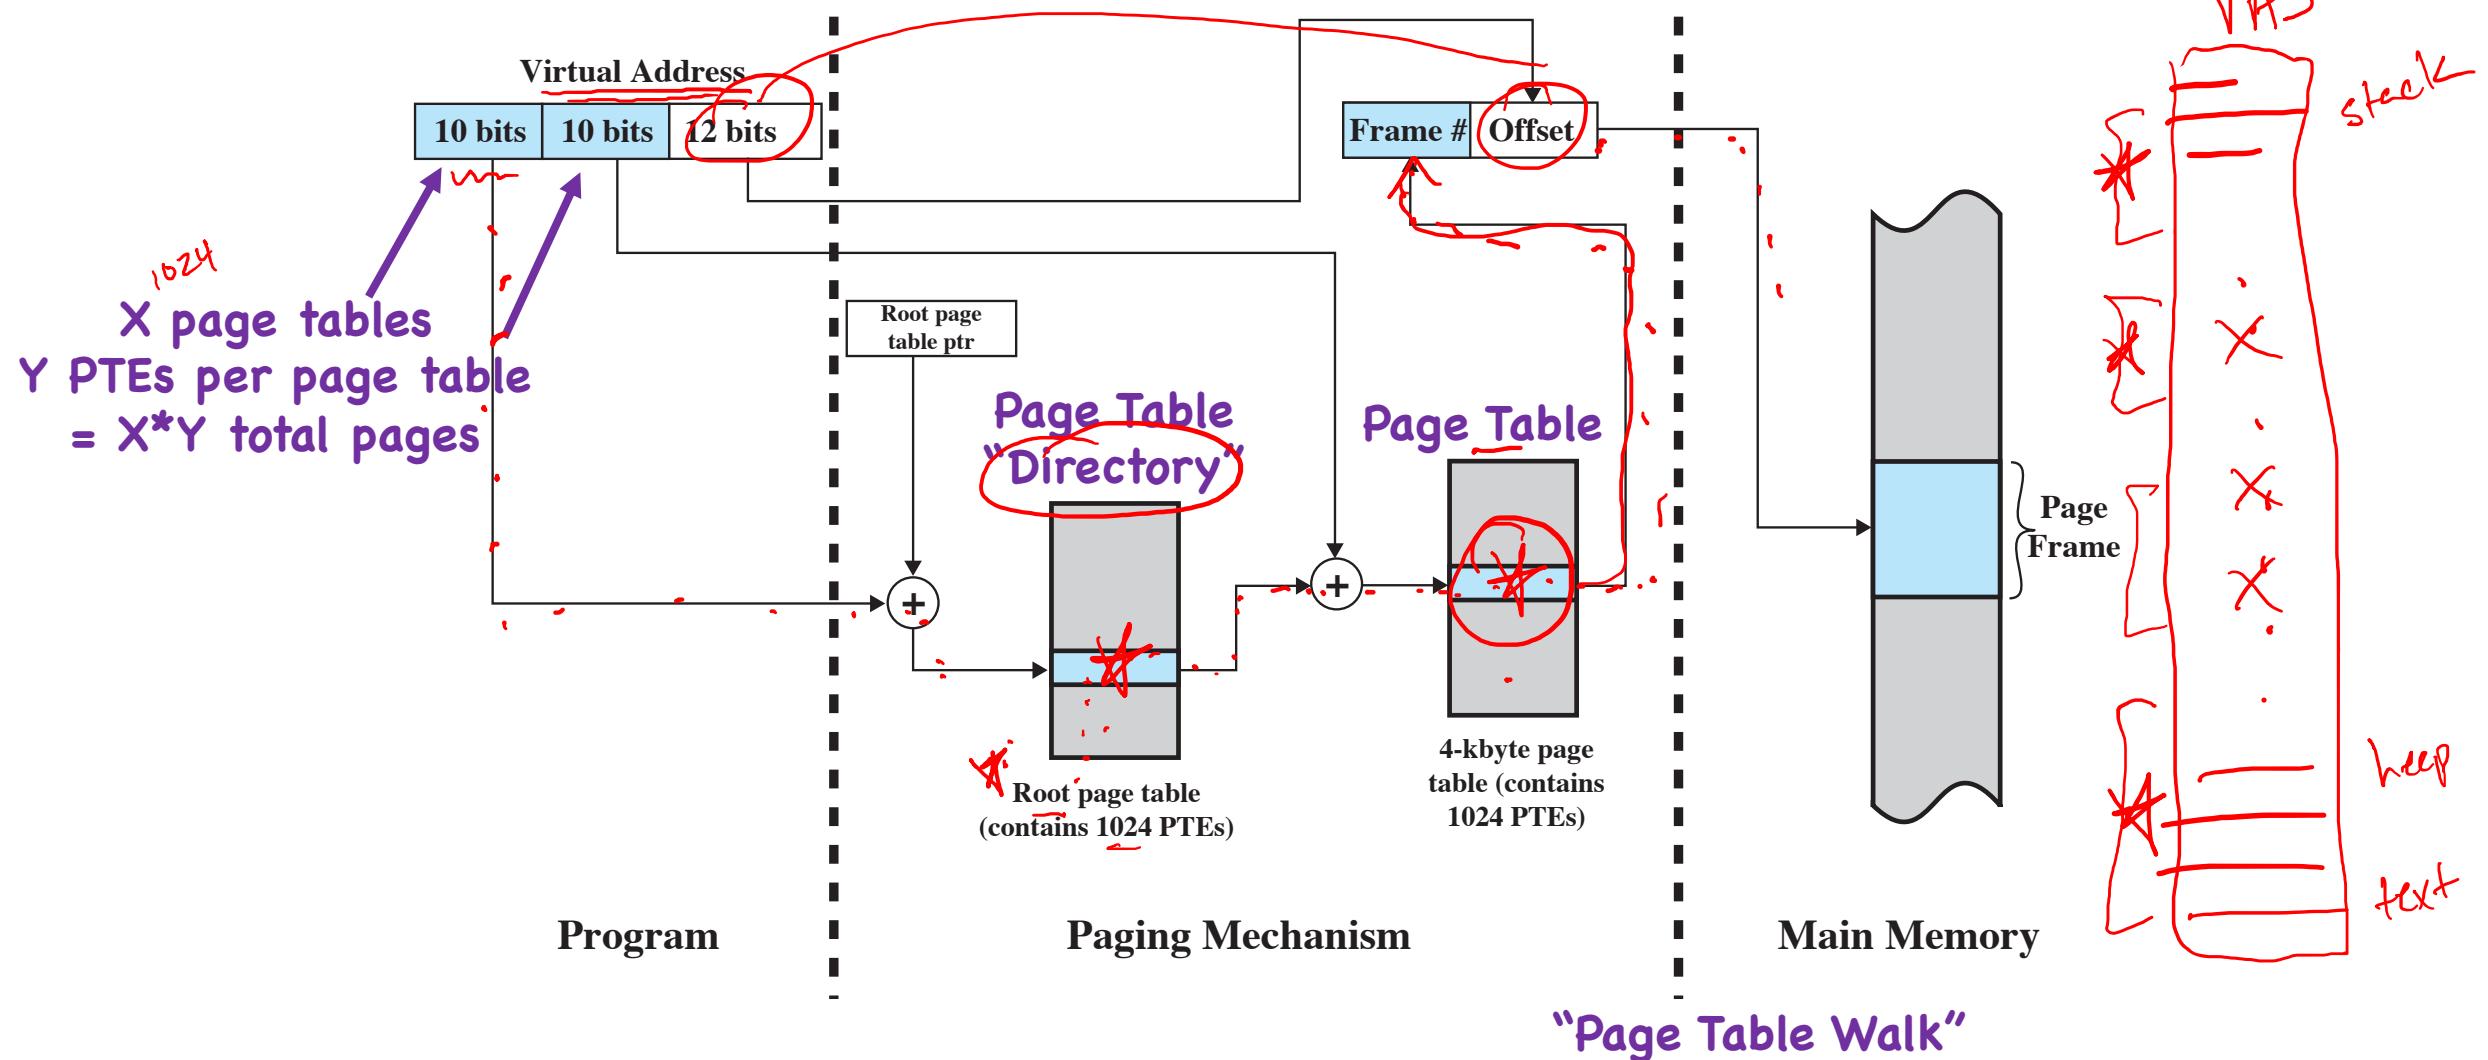
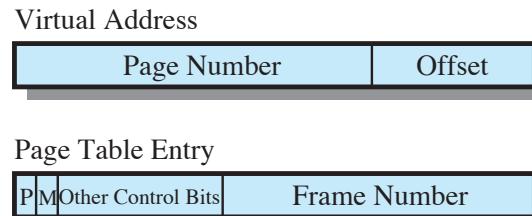
$$2^{22} / 2^{12} = 2^{10}$$



Address Translation: 2-Level Paging System (Overview)



Address Translation: 2-Level Paging System

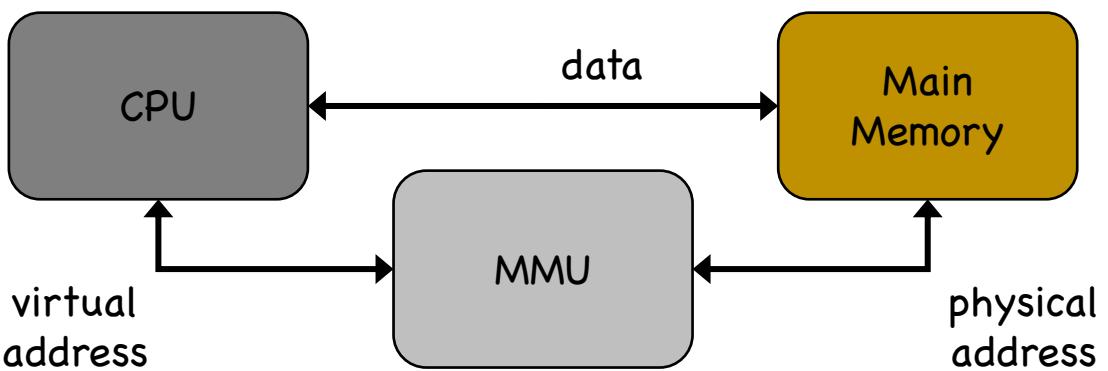
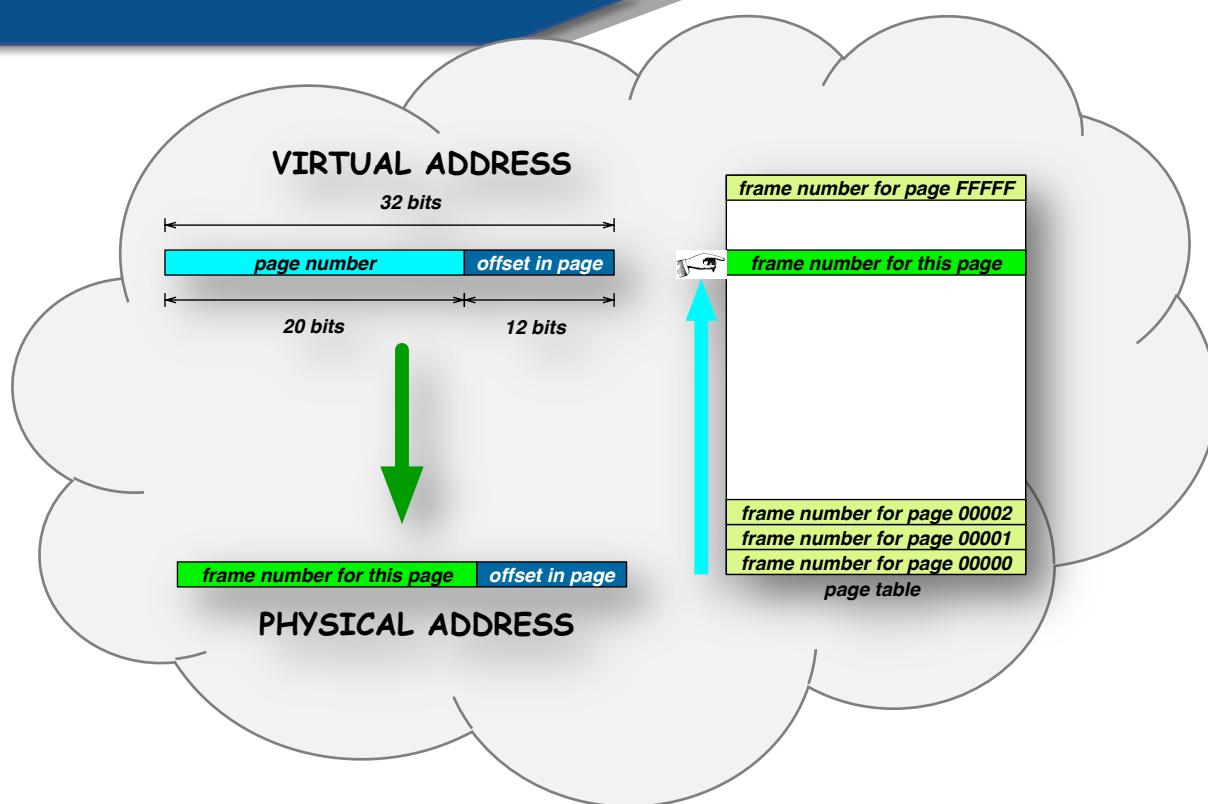


The TLB

An overview of the page table cache

So...

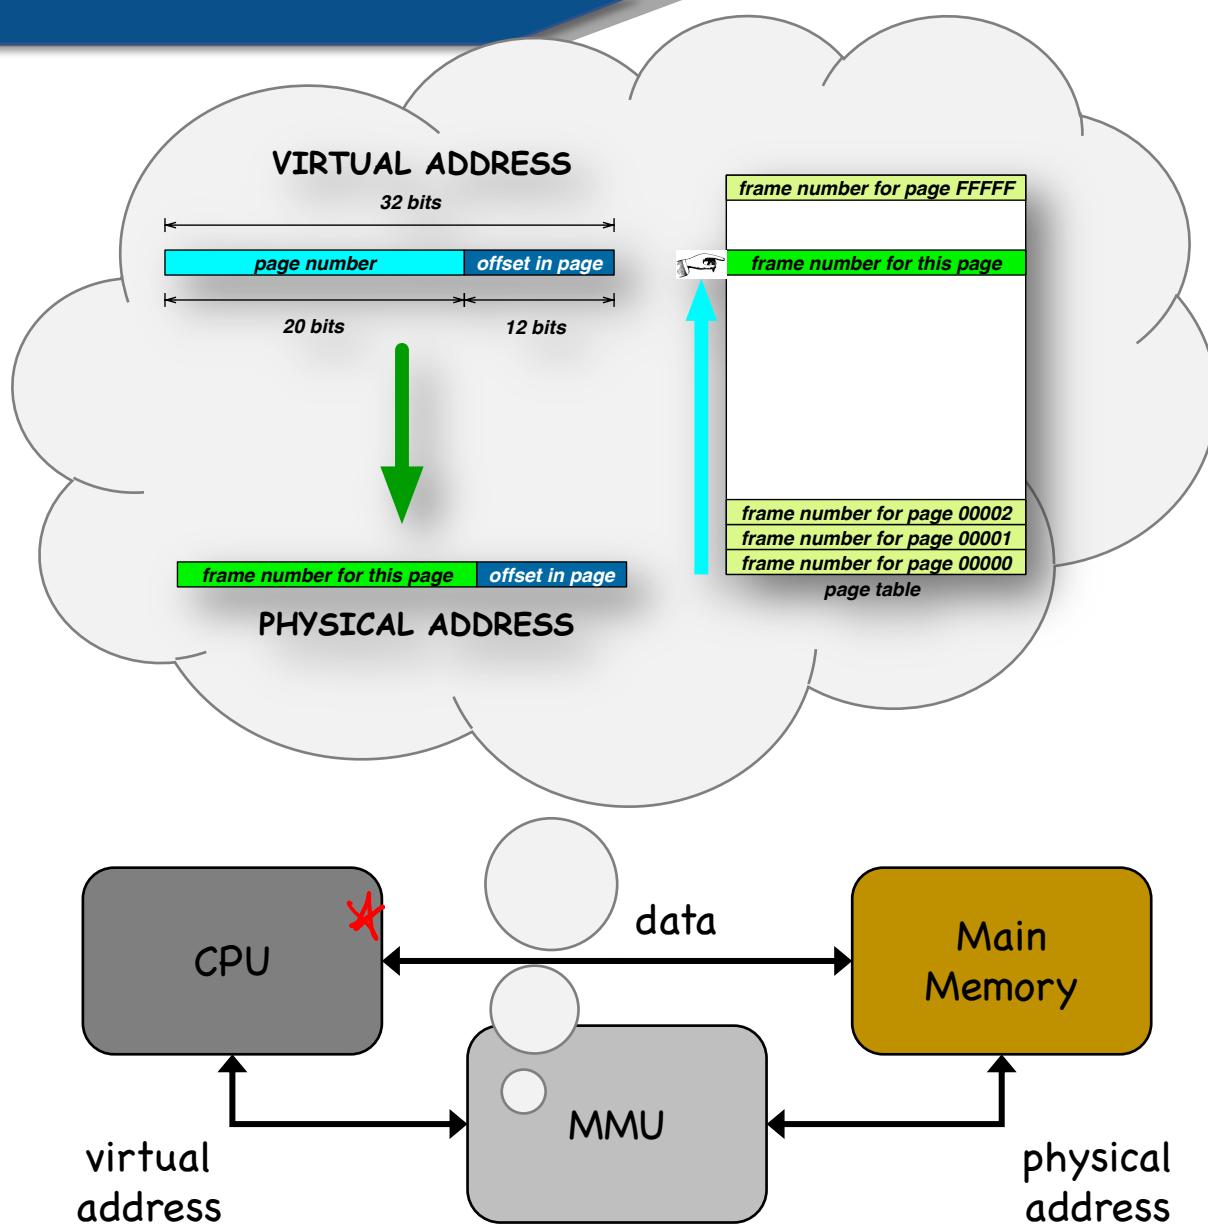
Q: Where should the page table live?



Option 1: CPU/MMU

Q: Where should the page table live?

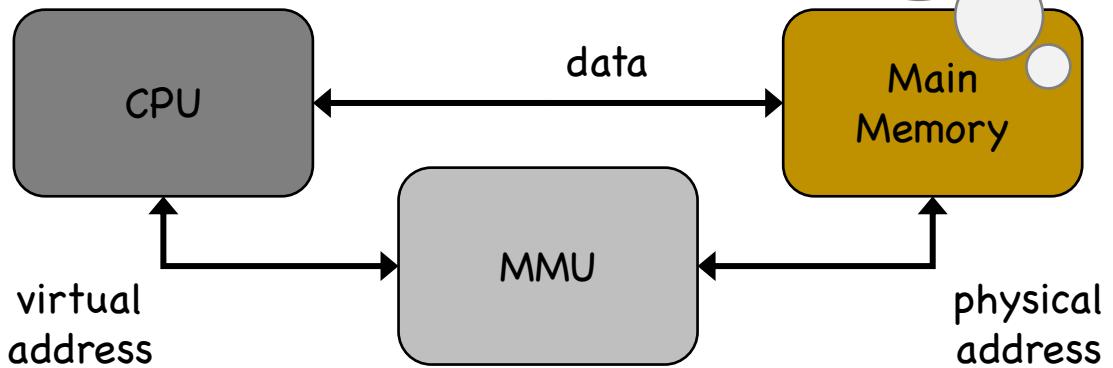
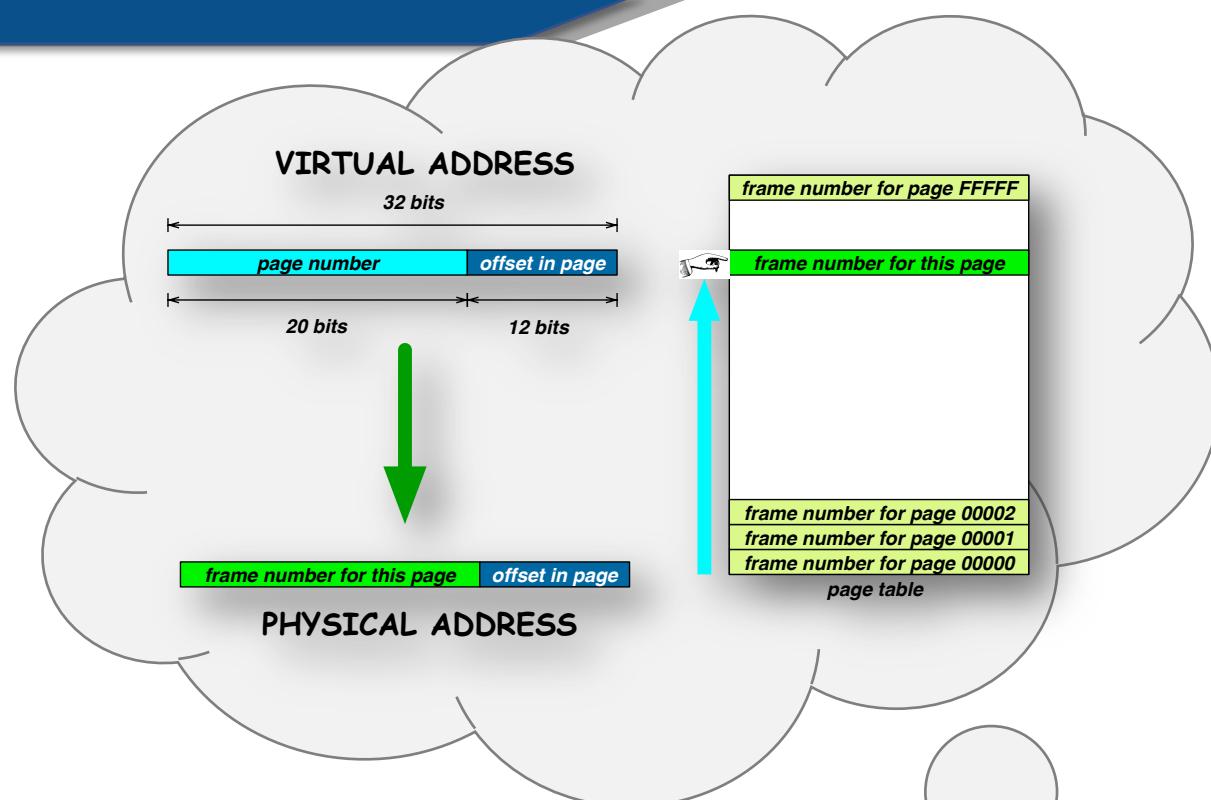
- Within special register(s) inside the MMU?
 - Too small... 😞



Option 2: MM

Q: Where should the page table live?

- Within special register(s) inside the MMU?
 - Too small... 😞
- Outside in physical memory?
 - How does MMU know where the page table is?
 - What about performance...?

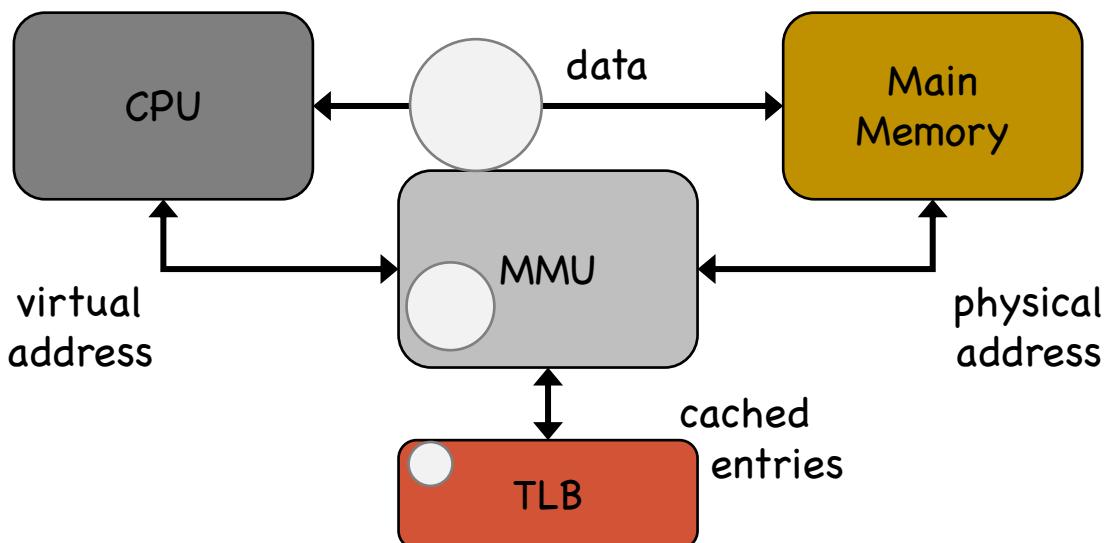
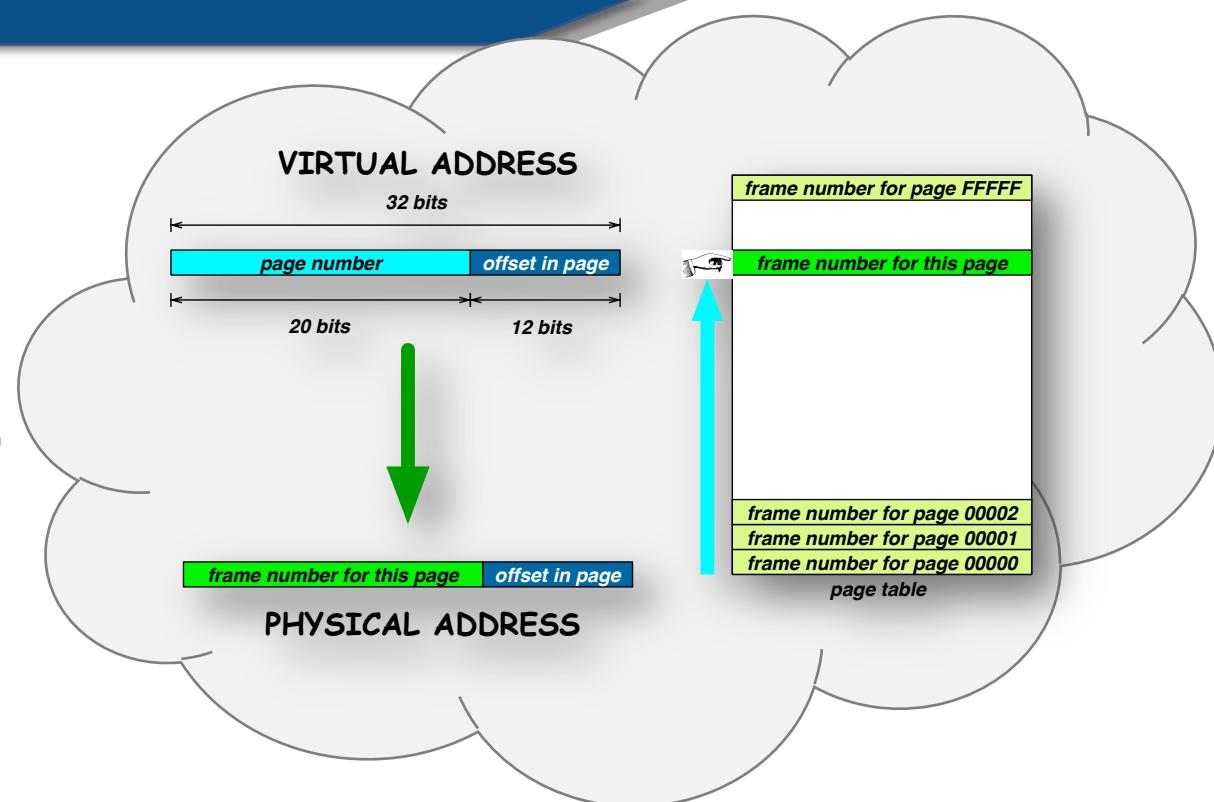


The TLB

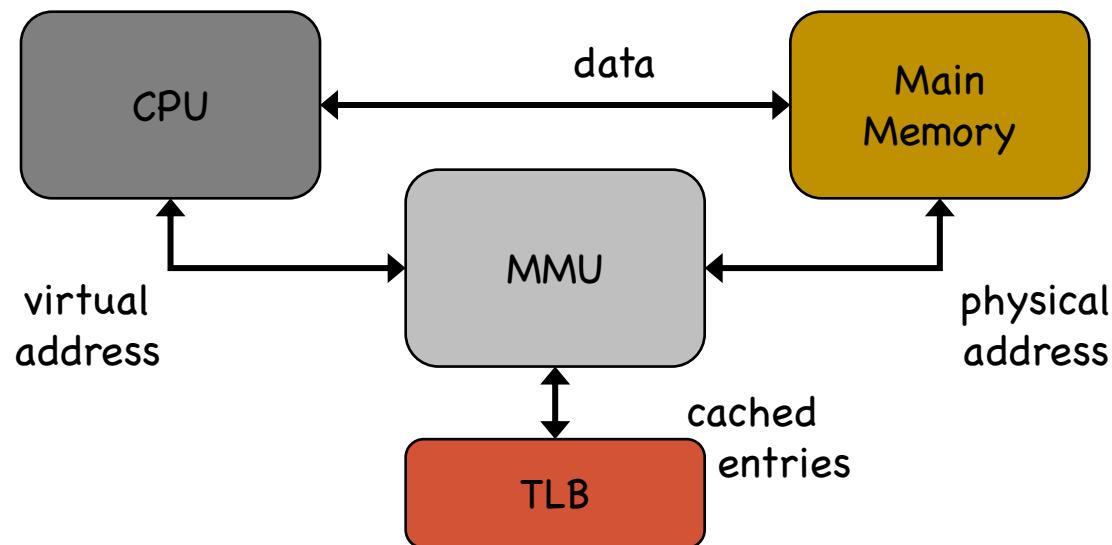
Q: Where should the page table live?

- Within special register(s) inside the MMU?
 - Too small... 😞
- Outside in physical memory?
 - How does MMU know where the page table is?
 - What about performance...?

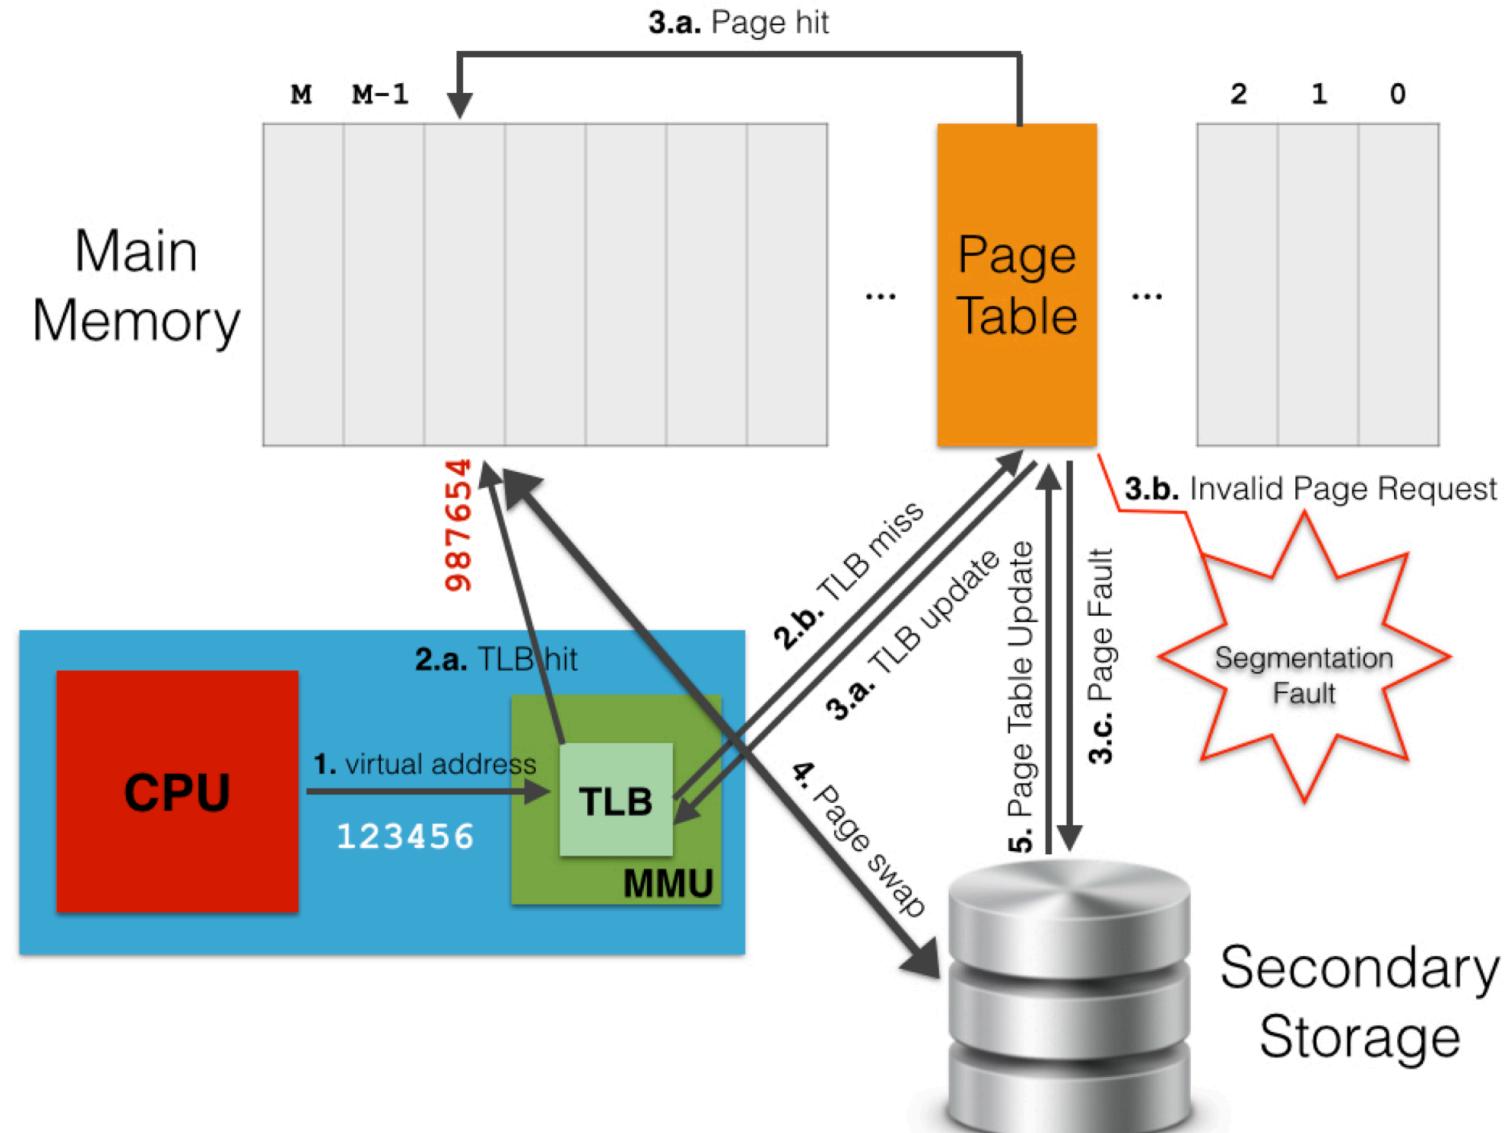
*The
Translation Lookaside Buffer (TLB)
to the rescue!*



Using a TLB in general...



- **TLB as a cache for PTEs**
 - TLB saves N <p,pte> pairs
 - check TLB first...
- **TLB Hit**
 - p exists in TLB
 - quickly returns pte (frame #)
- **TLB Miss**
 - p does not exist in TLB
 - search page table in MM
 - update TLB accordingly
 - *page fault* if pte not in MM!
 - effective access time – see: strider.c



Further Reading:

http://alasir.com/articles/cache_principles/tlb_virtual_memory.html