

# *Operating Systems!*

# **Memory: Approaches to Memory Management (Part 6)**

Prof. Travis Peters

Montana State University

CS 460 - Operating Systems

Fall 2020

<https://www.cs.montana.edu/cs460>

*Some diagrams and notes used in this slide deck have been adapted from Sean Smith's OS courses @ Dartmouth. Thanks, Sean!*

# Today

- Announcements
  - ~~Yalnix~~ How are projects/proposals going?! ☺
  - Congrats on submitting PA3!
- Upcoming Deadlines
  - **Sunday [10/25/2020] @ 11:59 PM (MST)** - Project Proposal



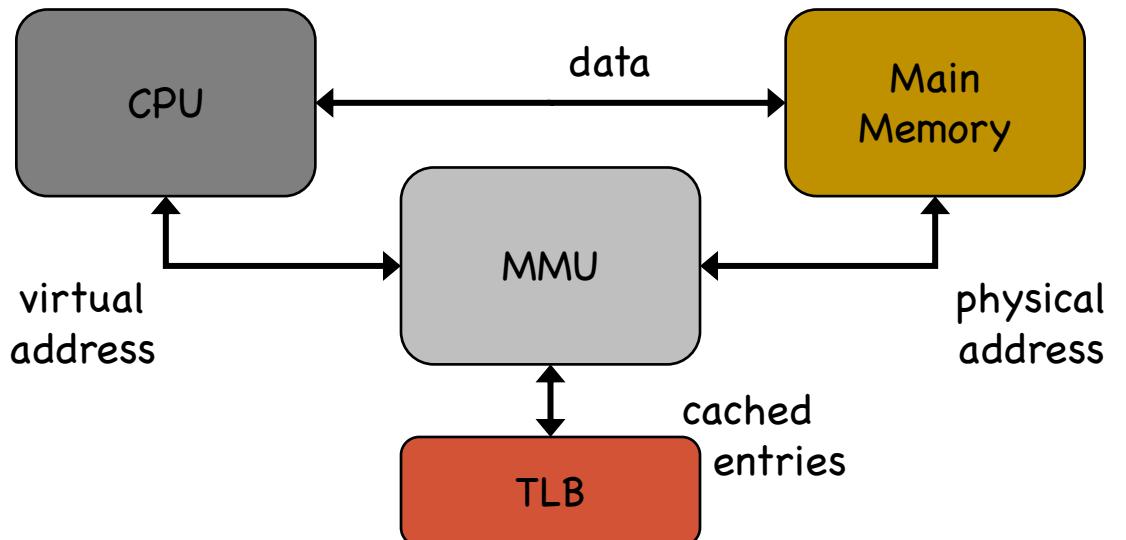
*Yalnix teams:  
you should plan to be near completing checkpoint 3 at this point!*

*Non-Yalnix teams:  
let me know if you want to discuss projects  
before submitting your proposal!*



# Today (cont.)

- Agenda
  - ~~Page Tables~~
  - ~~Heaps~~ → review fault.c 3,4
  - ~~Stacks~~ → (quickly) review fault.c 5,6,7 + sf.c
  - ~~Page Tables (Practical Issues)~~
  - ~~The TLB~~
  - ~~Switching Processes & the TLB~~
  - ~~Revisiting Syscalls~~
  - Memory Variations
  - Paging to Disk
  - ➔ Virtual Memory



# Memory Variations

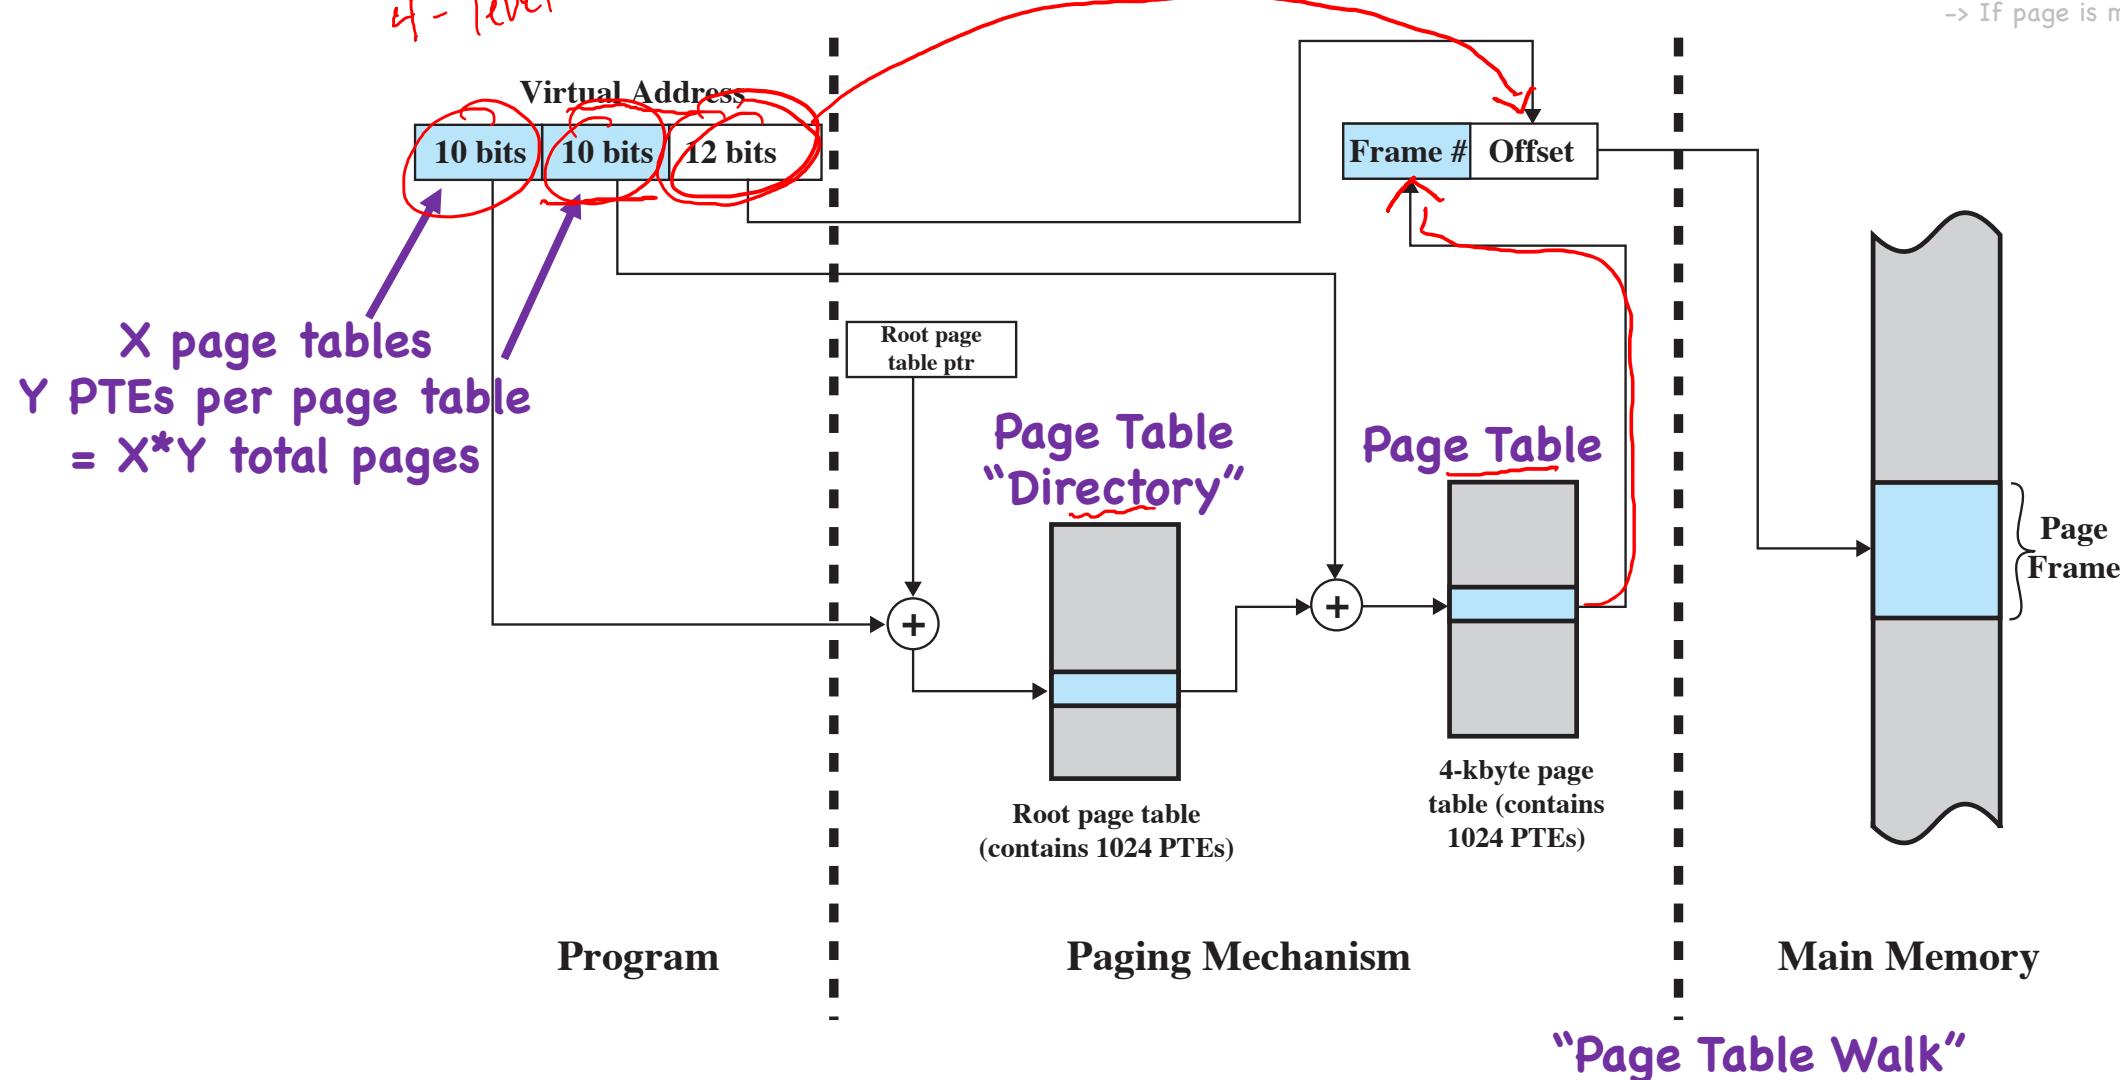
A quick look at other variations in memory management

Virtual Address

Page Number	Offset
-------------	--------

# Recall: 2-Level Paging System

1 - Level  
4 - Level

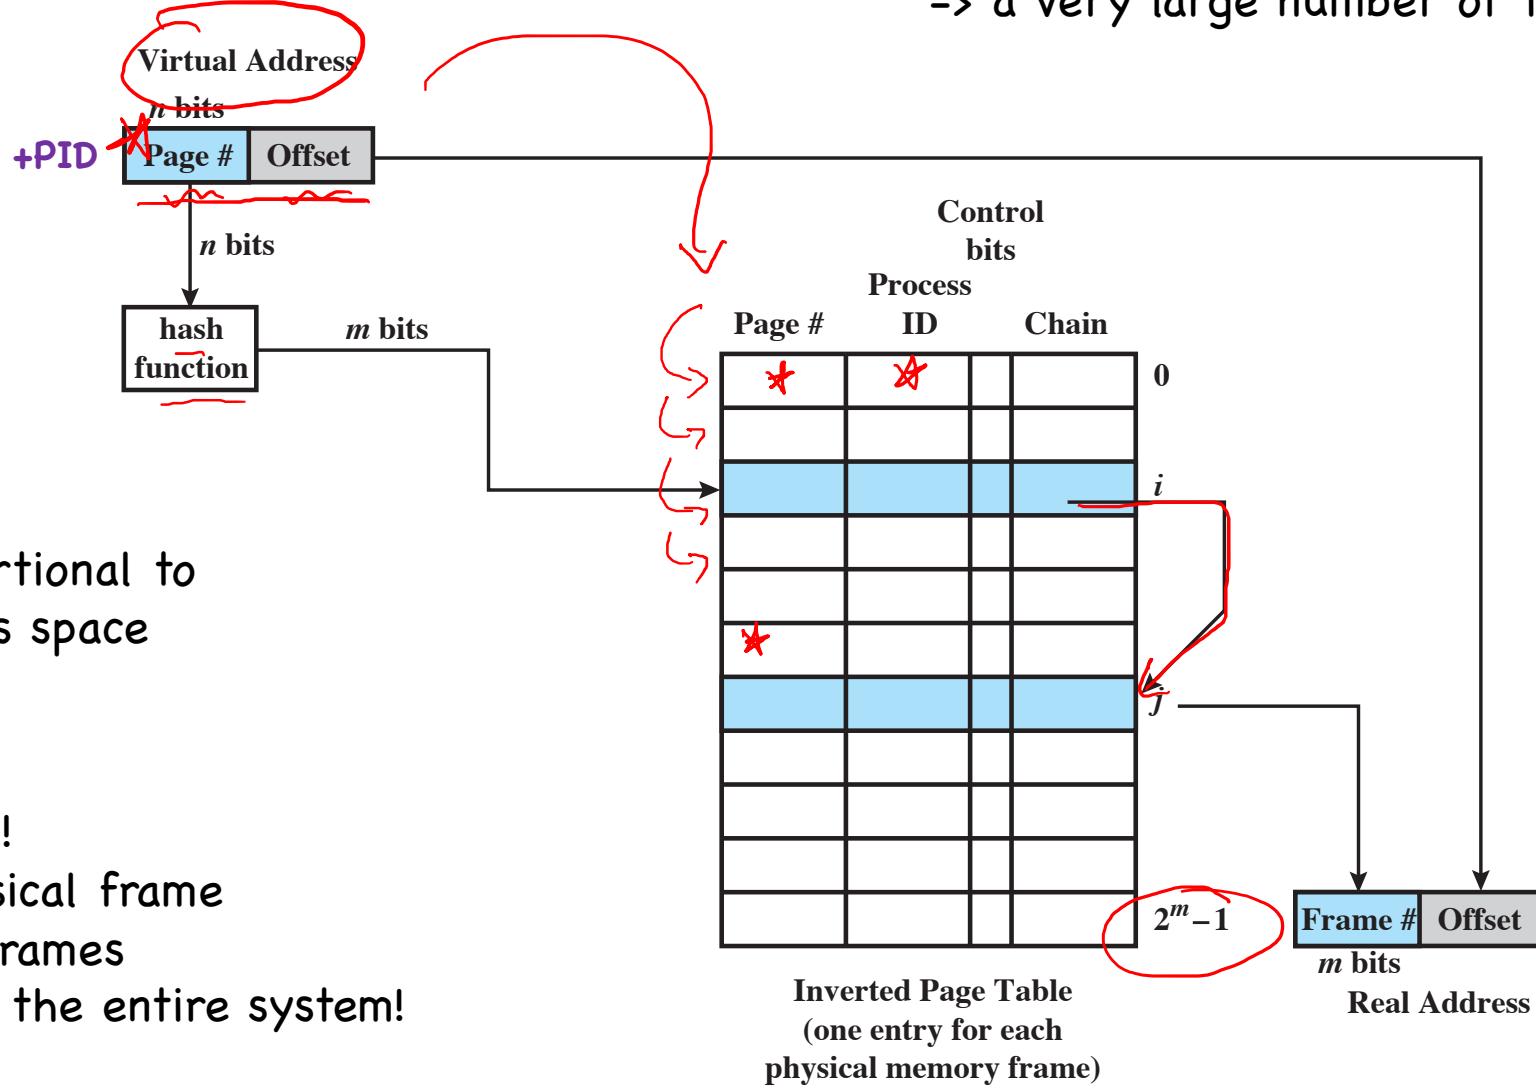


Page Table Entry

P	Other Control Bits	Frame Number
---	--------------------	--------------

- > If page is present (P), use frame number
- > If page is modified (M), write out

# Inverted Page Tables



# Segmentation

## Observation:

Programmers don't think about "pages" and "frames"...

... they think about regions (**segments**) such as "code" and "stack" and "heap"

Virtual Address

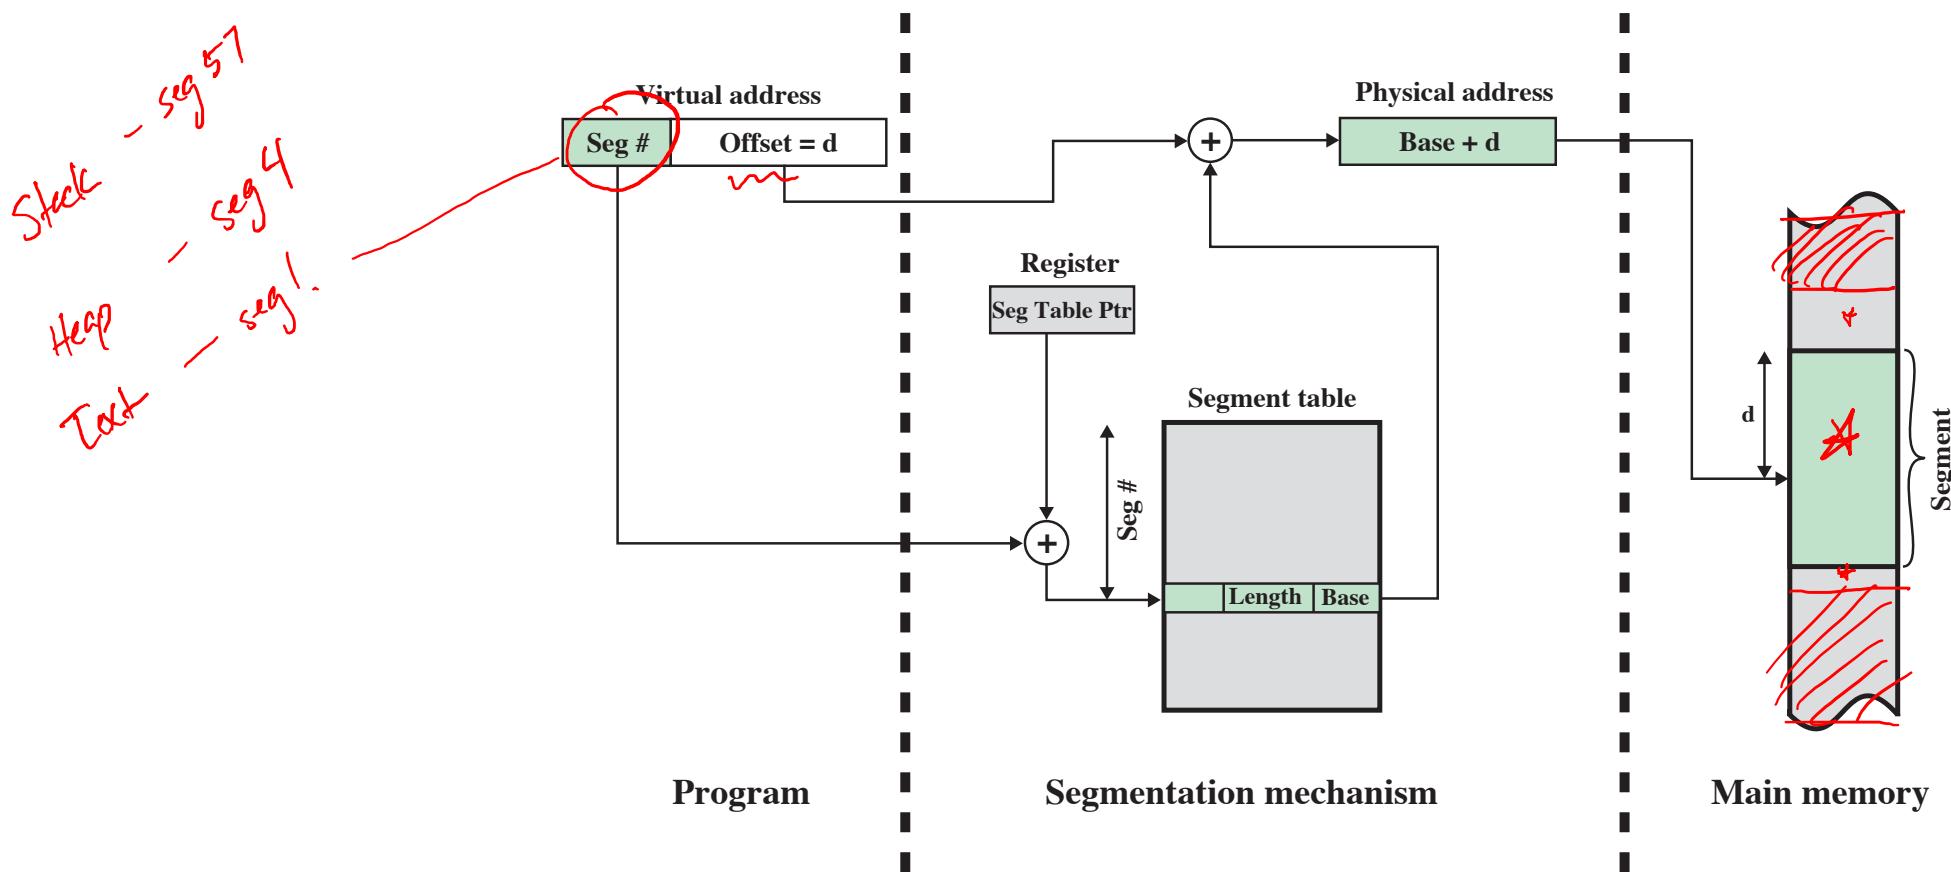
Segment Number	Offset
----------------	--------

Segment Table Entry

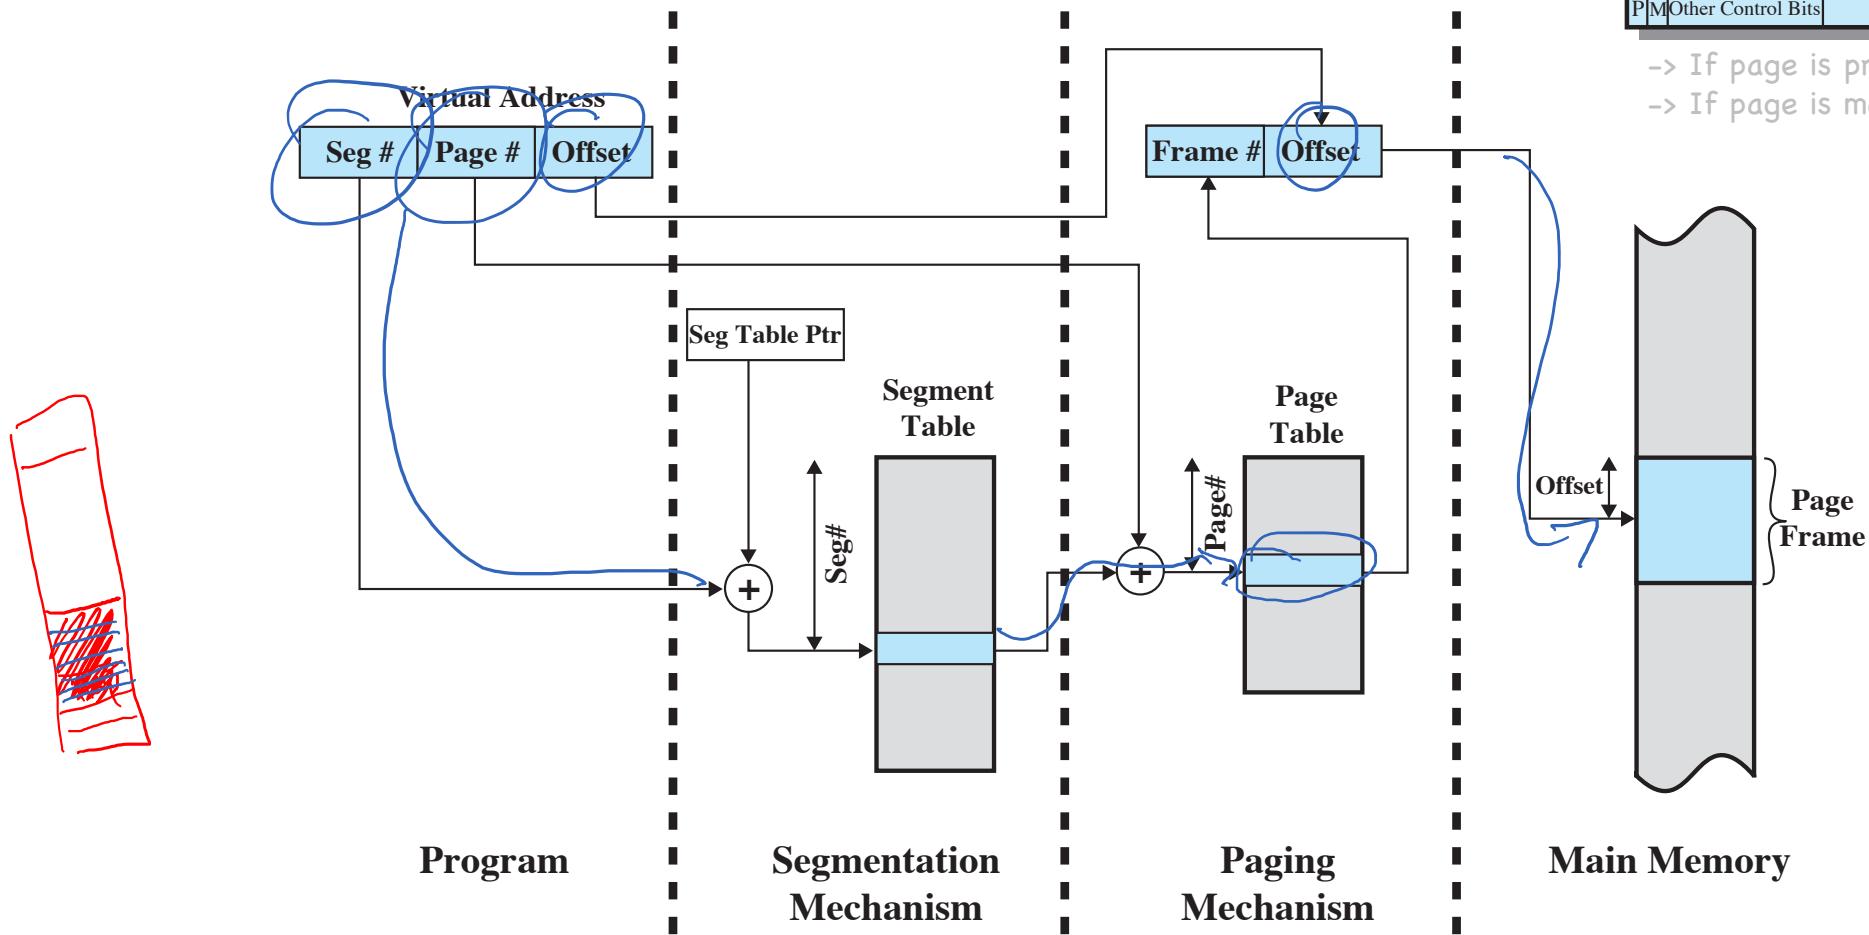
P   M   Other Control Bits	Length	Segment Base
----------------------------	--------	--------------

-> If seg. is present (P), use seg. number

-> If seg. is modified (M), write out



# Paging & Segmentation



Virtual Address

Segment Number	Page Number	Offset
----------------	-------------	--------

Segment Table Entry

Control Bits	Length	Segment Base
--------------	--------	--------------

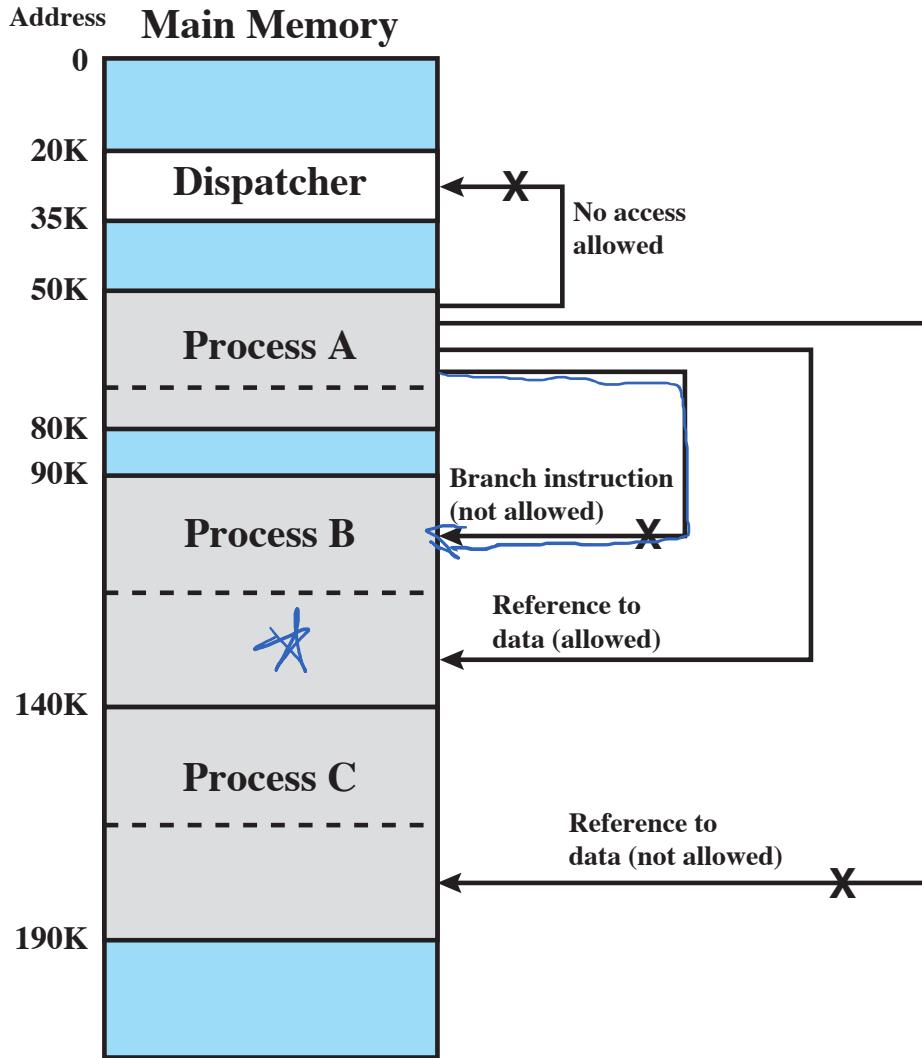
Page Table Entry

P	M	Other Control Bits	Frame Number
---	---	--------------------	--------------

- > If page is present (P), use frame number
- > If page is modified (M), write out

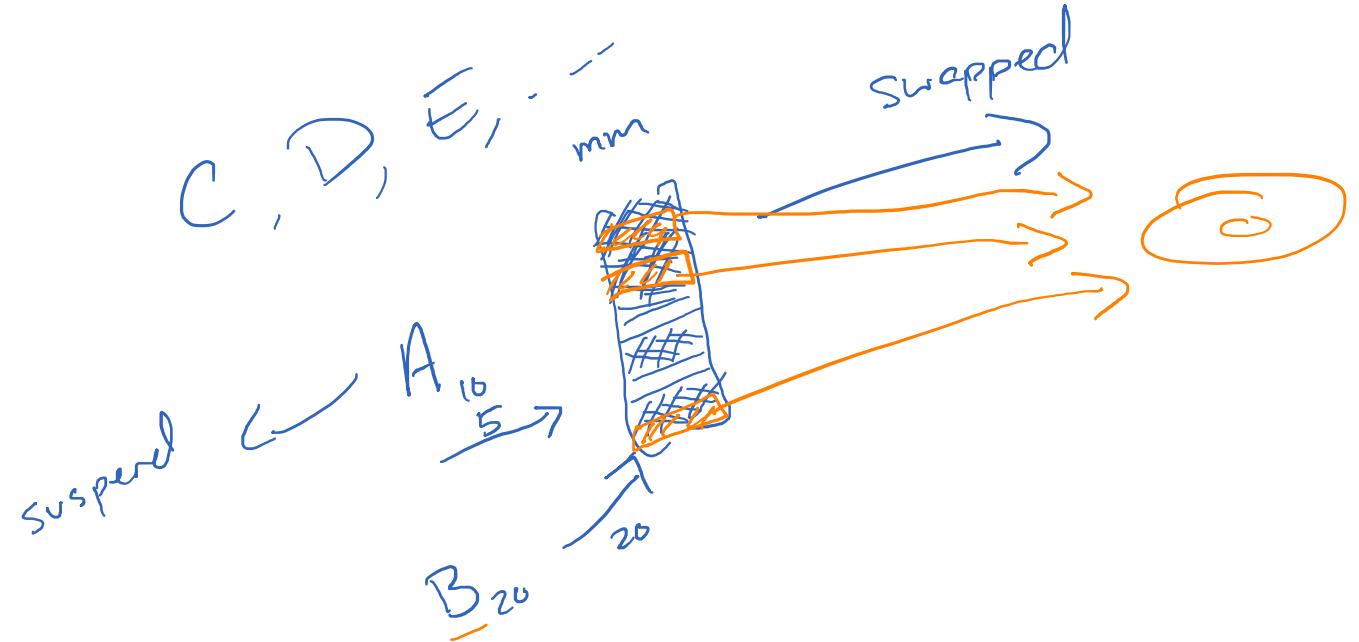
# Paging & Segmentation

(cont.)



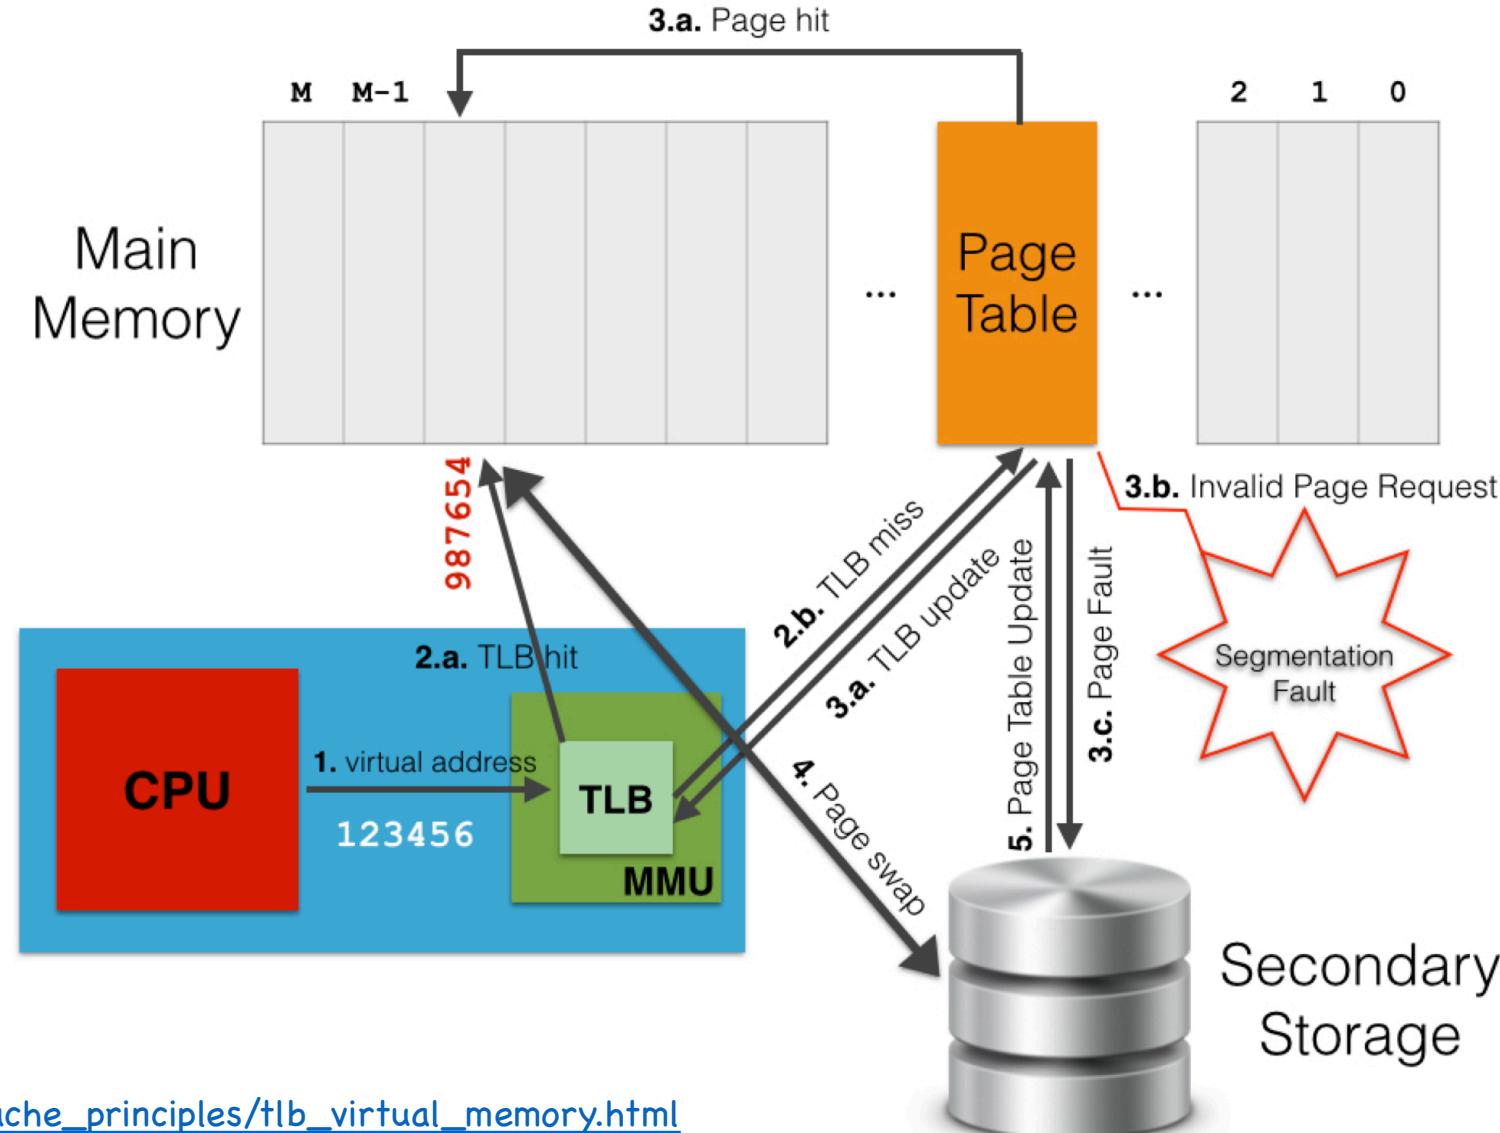
-> Configure settings  
for **segments** (r/w/x)

{  
-> Sure we can do this  
on a per-page level,  
but it is less intuitive



# Paging to Disk

# Paging to Disk



## Further Reading:

[http://alasir.com/articles/cache\\_principles/tlb\\_virtual\\_memory.html](http://alasir.com/articles/cache_principles/tlb_virtual_memory.html)

<https://gabrieletolomei.files.wordpress.com/2013/10/mmuv.jpg>

# Designing an OS with Virtual Memory

- Goal: Reduce Frequency of Page Faults
- *Why? What's the cost?*

r/w    to/from disk

# Designing an OS with Virtual Memory

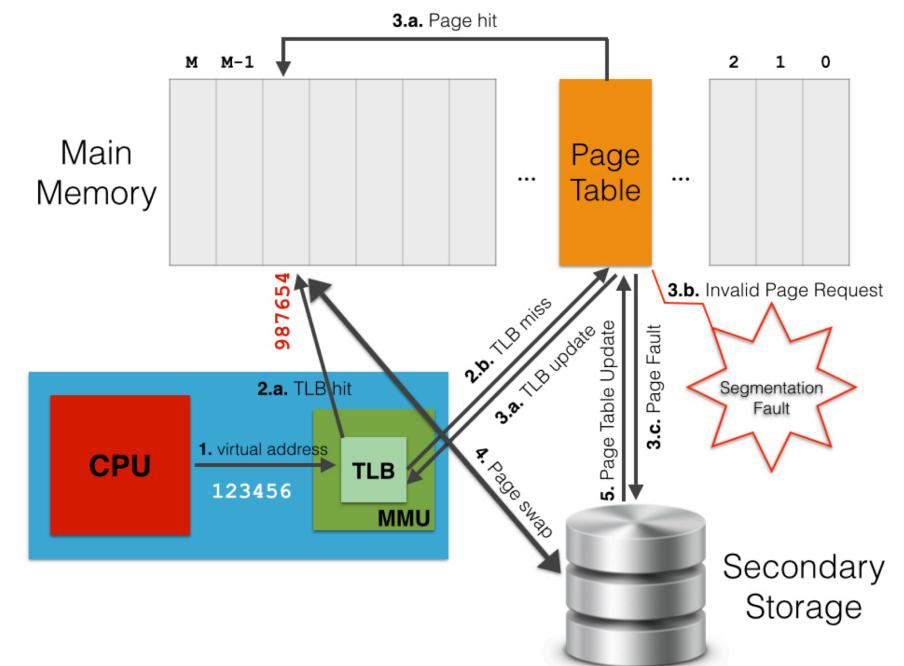
- Goal: Reduce Frequency of Page Faults
- ***Why? What's the cost?***
  - Must make decision for which page(s) to swap in/out + associated I/O
  - Process switch

***What issues must we consider to address this goal?***

# Fetch Policy

*When should a page be brought into main memory?*

- Demand Paging
  - Bring pages in when a referenced
  - as needed / on demand
- Pre-paging
  - Bring in several pages at once
  - Ex: other nearby pages on disk



<https://gabrieletolomei.files.wordpress.com/2013/10/mmuv.jpg>

# Placement Policy

*Where should a page be put in main memory?*

- Segmentation?
  - Best-Fit, First-Fit, Etc.
- Paging?
  - Not really an issue: page-frame mappings are equally efficient
  - Just need a free frame...

# Replacement Policy

*Which page\* should be replaced when a new page must be brought into main memory\*\*?*

- Strategies:
  - Optimal (OPT) --- sometimes called “MIN”
  - LRU
  - FIFO
  - Clock (“2nd Chance”)
  - + Random Placement, ...

\* Want to replace an “unlocked” page that is least likely to be referenced in the near future.

\*\*Most interesting if we assume that all frames in MM are occupied, then a page fault occurs.

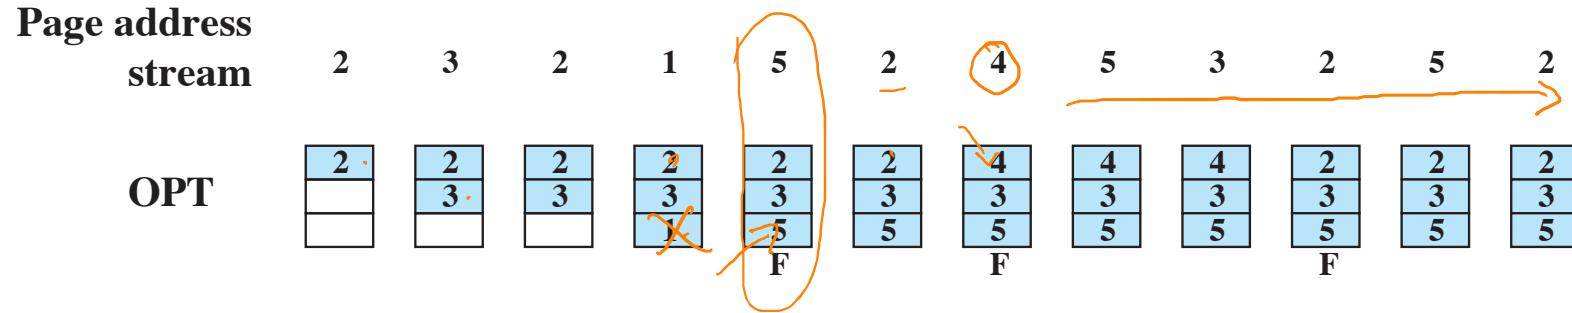
# **Example:** Page-Replacement Algorithms

*Reference 5 distinct pages...*

Page address stream	2	3	2	1	5	2	4	5	3	2	5	2
------------------------	---	---	---	---	---	---	---	---	---	---	---	---

# **Example:** Page-Replacement Algorithms

*Reference 5 distinct pages...*



# Example: Page-Replacement Algorithms

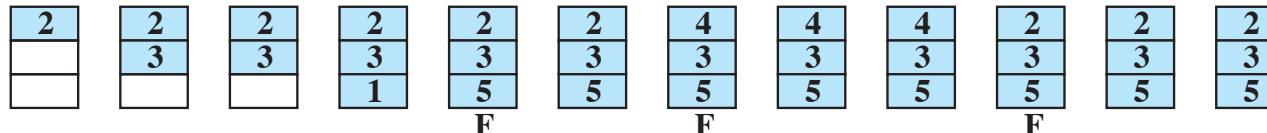
Reference 5 distinct pages...

Page address

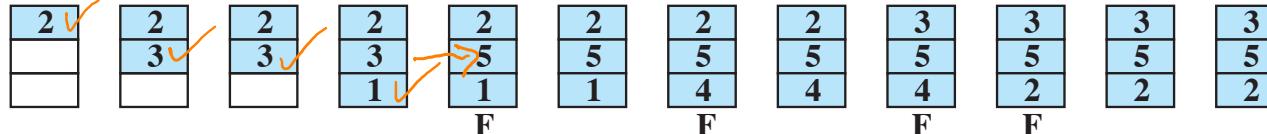
stream

2      3      2      1      5      2      4      5      3      2      5      2

OPT



LRU



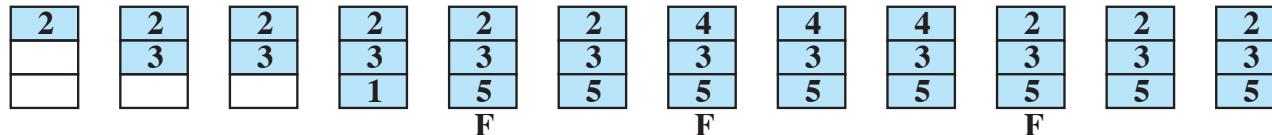
# **Example:** Page-Replacement Algorithms

*Reference 5 distinct pages...*

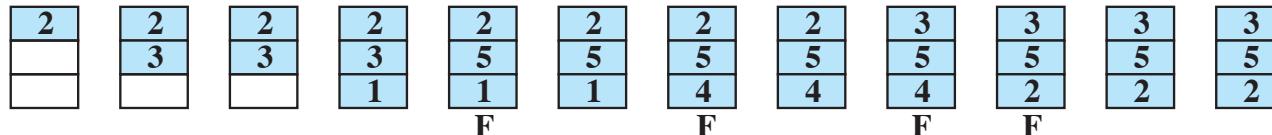
Page address  
stream

2    3    2    1    5    2    4    5    3    2    5    2

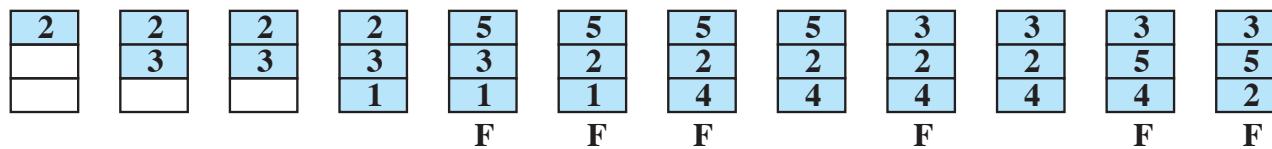
OPT



LRU



FIFO



# Example: Page-Replacement Algorithms

*Reference 5 distinct pages...*

Page address  
stream

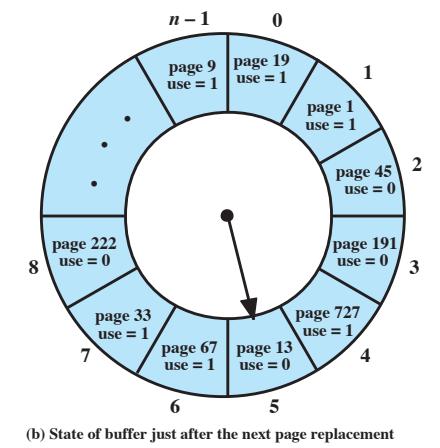
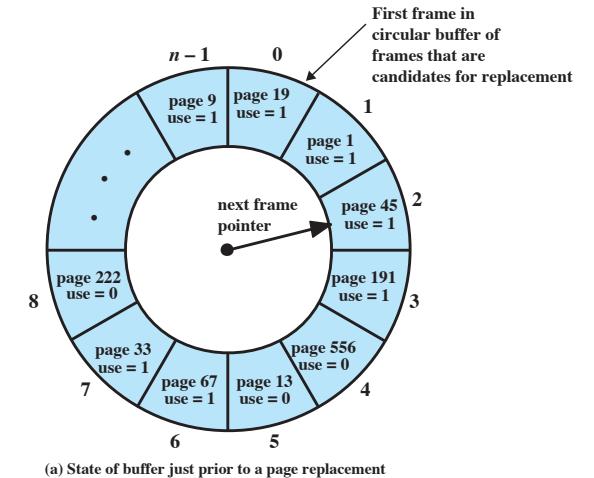
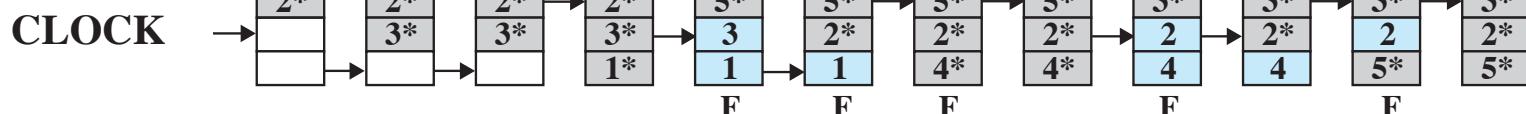
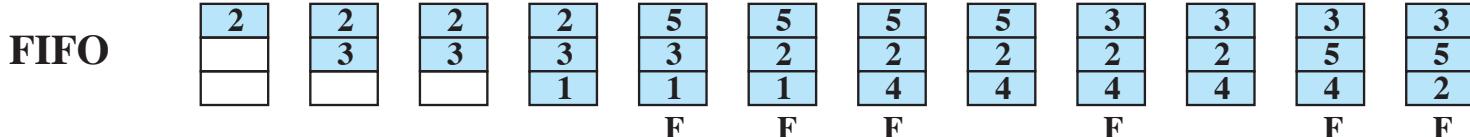
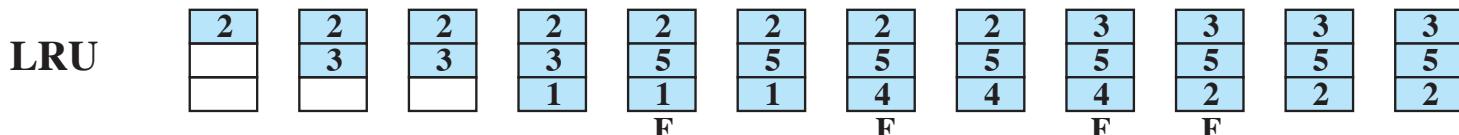
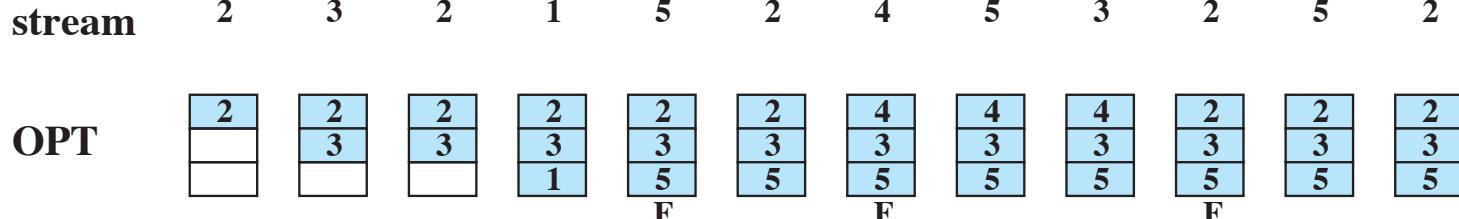


Figure 8.15 Example of Clock Policy Operation

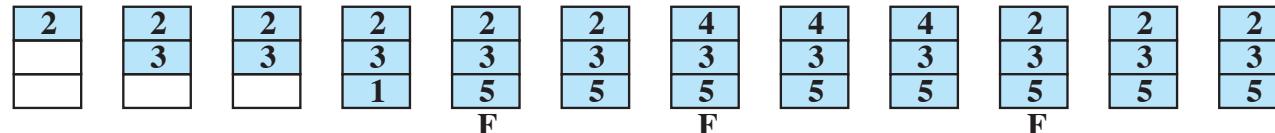
# Example: Page-Replacement Algorithms

Reference 5 distinct pages...

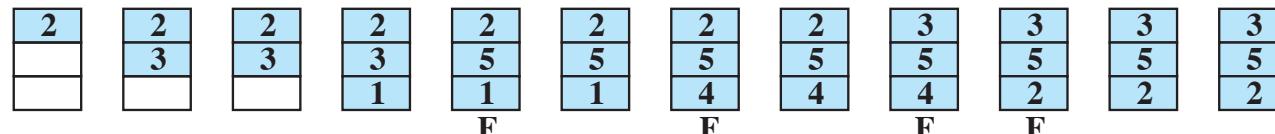
Page address  
stream

2    3    2    1    5    2    4    5    3    2    5    2

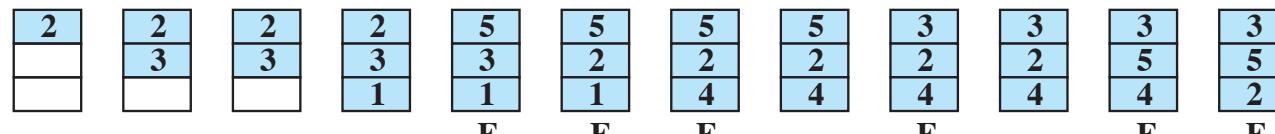
OPT



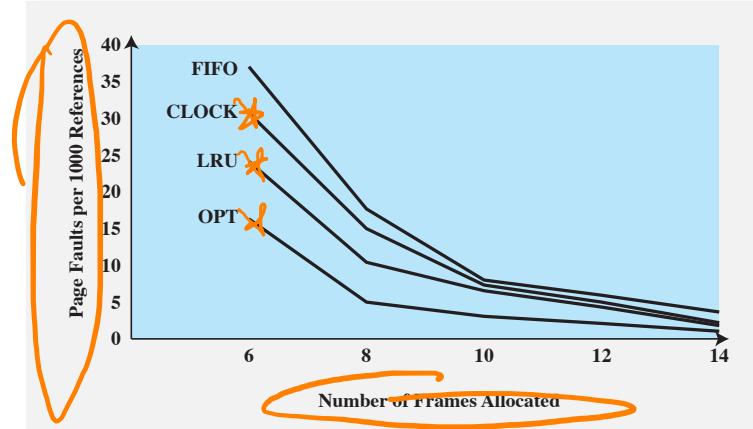
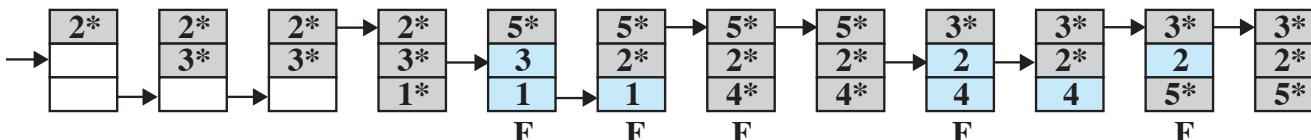
LRU



FIFO



CLOCK



# Going Further...

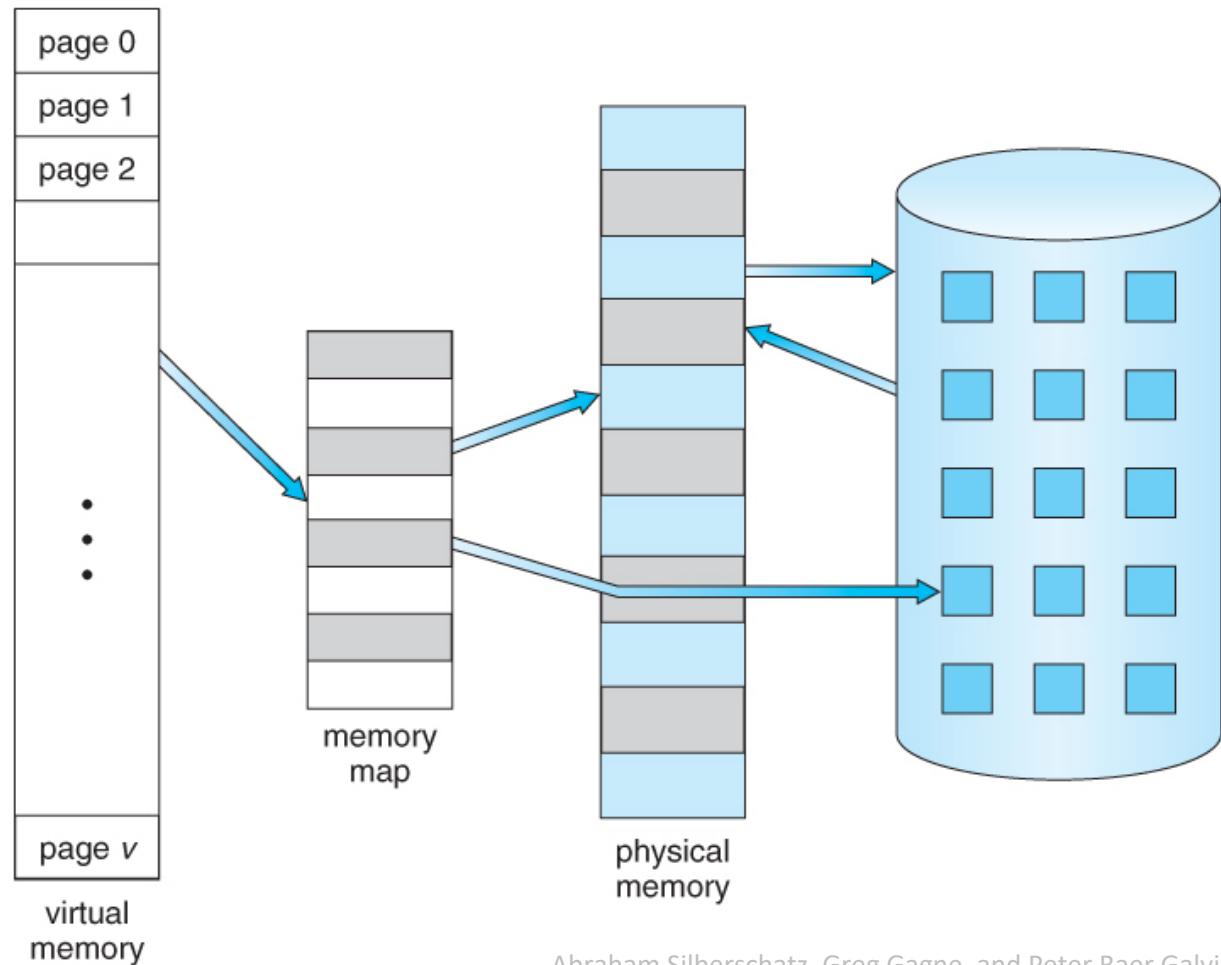
- Resident Set Management
  - *how many pages should be in main memory?*
    - Resident Set Size: **Fixed** vs. **Variable**
    - Replacement Scope: **Local** (current process) vs. **Global** (everything)
- Cleaning Policy
  - *when should a page be written out from main memory?*  
(opposite of "Fetch Policy")
- Load Control
  - *how many processes should be in main memory?*  
(i.e., the level/degree of multiprogramming)

# Virtual Memory

Putting it all together!

# Virtual Memory

1. Bring bare minimum of process into main memory (MM)
2. Run process in the normal way
3. Use page tables to keep track of pages
4. Page not in MM?
  - > process moved to "blocked" state
  - > OS issues "read" from disk (re: page table)
  - > schedule other processes...
  - > I/O interrupt indicates when page is in MM
  - > blocked process moved back to "ready" state
5. Use replacement strategy to make space as needed



- Abraham Silberschatz, Greg Gagne, and Peter Baer Galvin,  
"Operating System Concepts," Ninth Edition, Chapter 9

# Summary

## **Big Ideas**

- Programs broken into non-contiguous pieces (pages or segments)
- Logical addresses dynamically translated into physical addresses at run-time
- Not all memory needs to be in MM during execution!
- Special HW/caches help to speed up memory access

## **So what is Virtual Memory?**

- All memory can be addressed as if it were part of main memory
- References to memory are dynamically translated
- Limited only by the addressing scheme of the system & amount of secondary memory

