

# *Operating Systems!*

# Scheduling: **Uniprocessor Scheduling** **(Part 1)**

Prof. Travis Peters

Montana State University

CS 460 - Operating Systems

Fall 2020

<https://www.cs.montana.edu/cs460>

Some diagrams and notes used in this slide deck have been adapted from Sean Smith's OS courses @ Dartmouth. Thanks, Sean!

# CPU scheduling

JULIA EVANS  
@bork

once upon a time...



I want to run a program!

Sorry I can only do 1 program at a time from the 60s or so

COMP UTARR

your computer today

program

no, me!

ME

operating system

I want to run!

no, me!

I have stuff to do too you know!

every CPU core can only run 1 program at a time

ok time to stop it's Jimmy's turn to use the CPU now

operating system

every program gets a few ms at a time

steps when we switch the running process  
"context switch"

→ save:

- registers
- stack pointer
- which CPU instruction to start at next time

→ set up memory for new process

→ load new registers and stuff

all this takes time  
(2 microseconds?).  
It's ok to do but  
you don't want to be switching processes constantly

you don't use the CPU when you're waiting



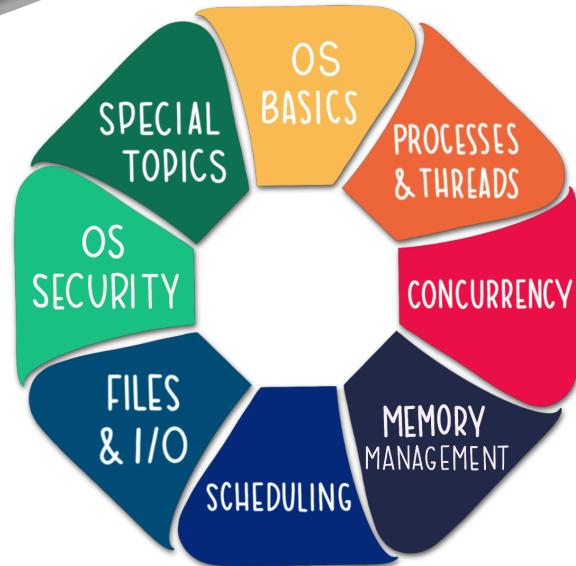
hey I'm waiting for a network response

cool! I'll run other stuff until that comes back.

<https://drawings.jvns.ca/scheduling/>

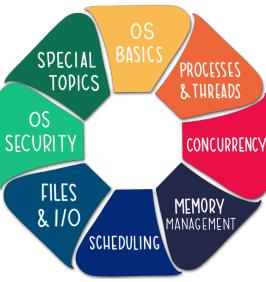
# Today

- Announcements
  - ~~Yalnix~~ How are projects/proposals going?! 😊
  - Congrats on submitting PA3!
- Upcoming Deadlines
  - **Project Proposal (HARD DEADLINE)**  
**Sunday [10/25/2020] @ 11:59 PM (MST)**



***Yalnix Teams:**  
You should plan to be near completing  
checkpoint 3 at this point!*

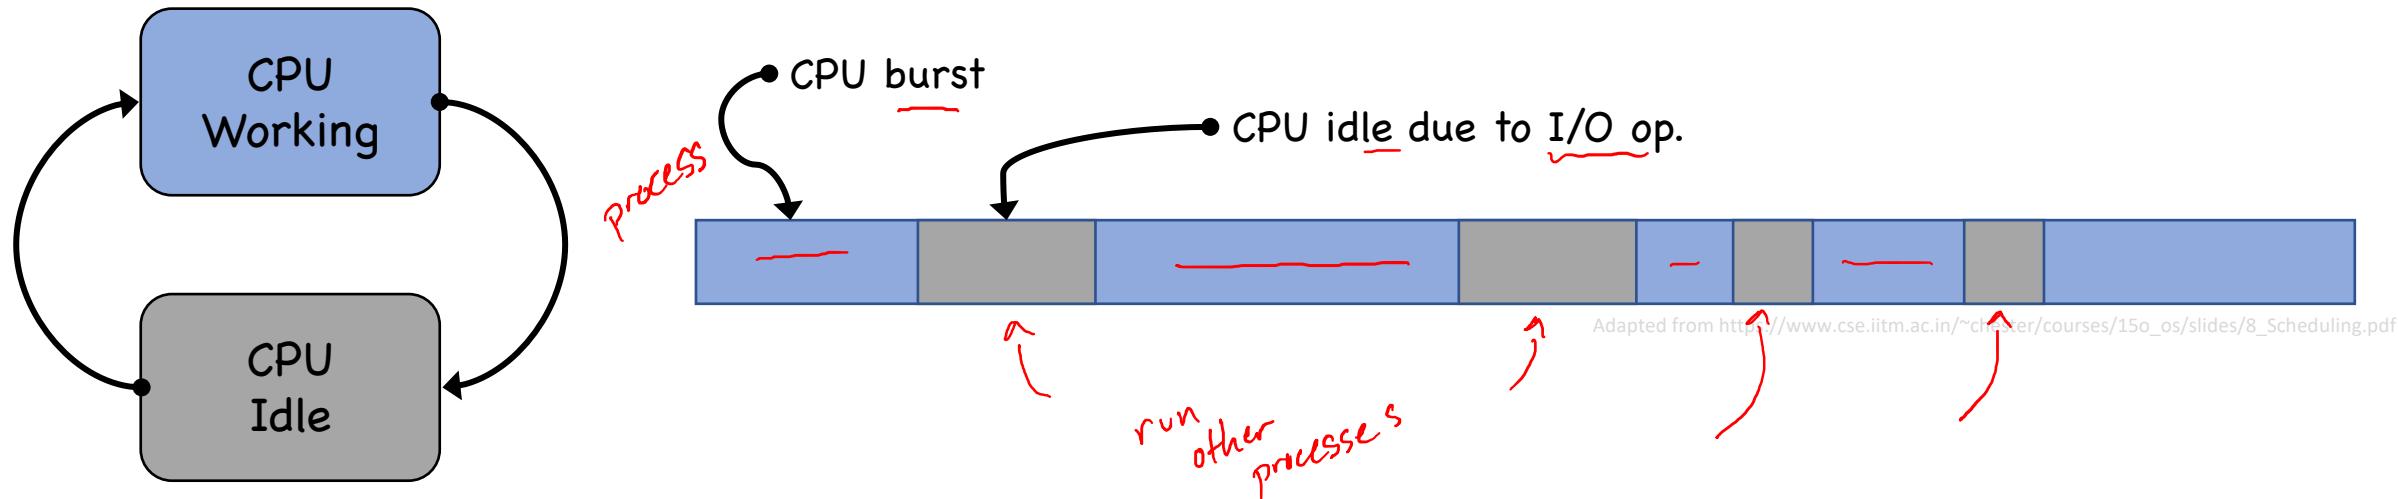
***Non-Yalnix Teams:**  
Let me know if you want to discuss projects  
before submitting your proposal!*



# Today (cont.)

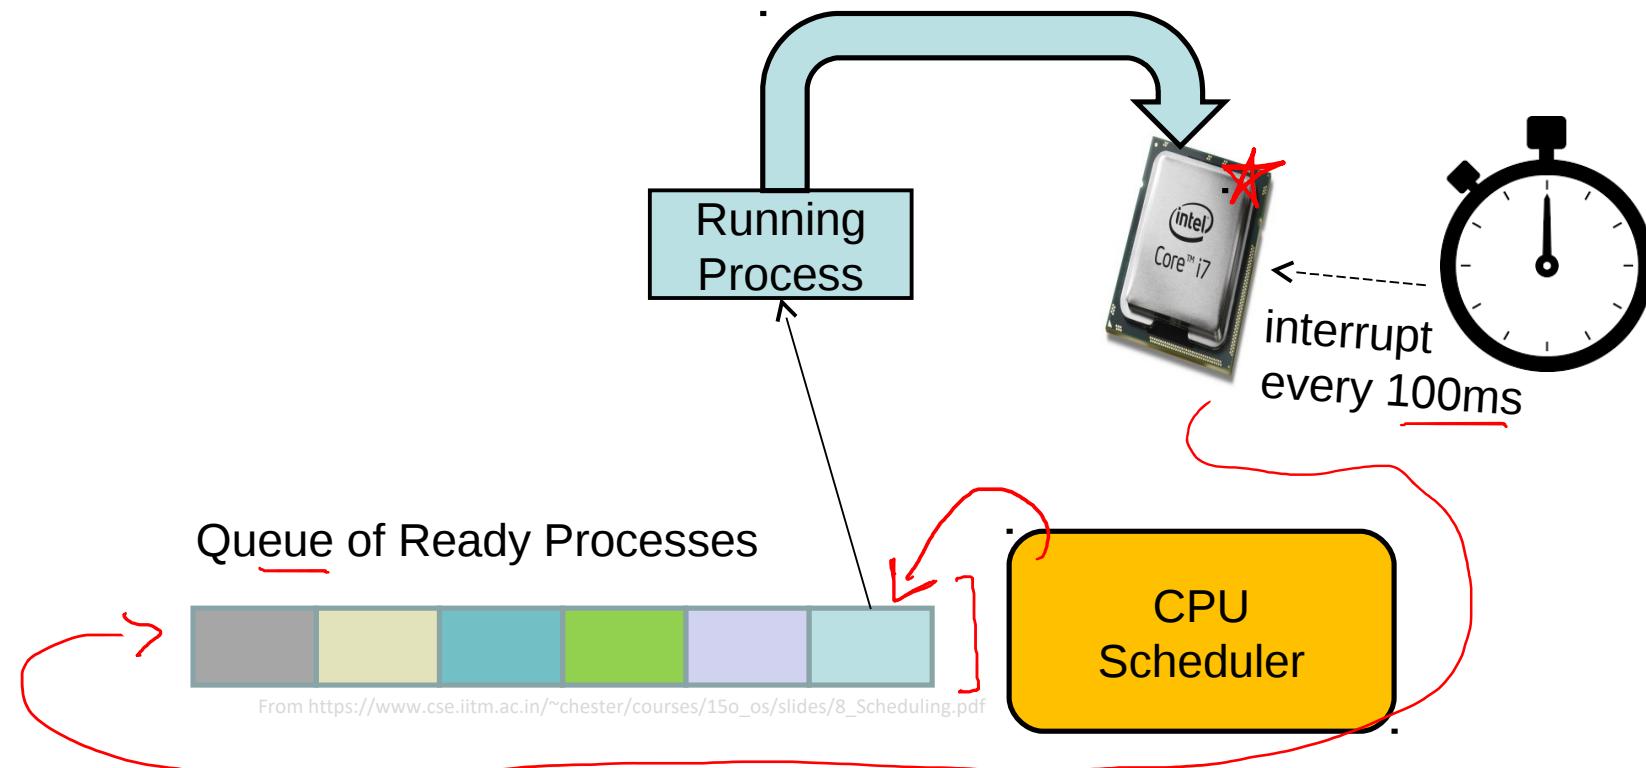
- Agenda

- Types of Processes
- Long-, medium-, and *short-term* scheduling
- Performance of different scheduling policies



Adapted from [http://www.cse.iitm.ac.in/~chatter/courses/150\\_os/slides/8\\_Scheduling.pdf](http://www.cse.iitm.ac.in/~chatter/courses/150_os/slides/8_Scheduling.pdf)

# CPU Scheduler

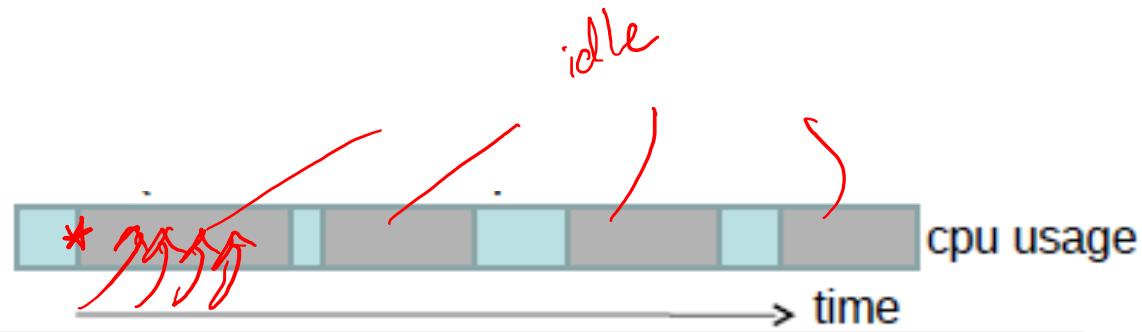


- Scheduler triggered to run when timer-based interrupt occurs, or when process is blocked on I/O
- Scheduler picks another process from the Ready Queue (RQ)
- Switch to new process via context switch

# Types of Processes

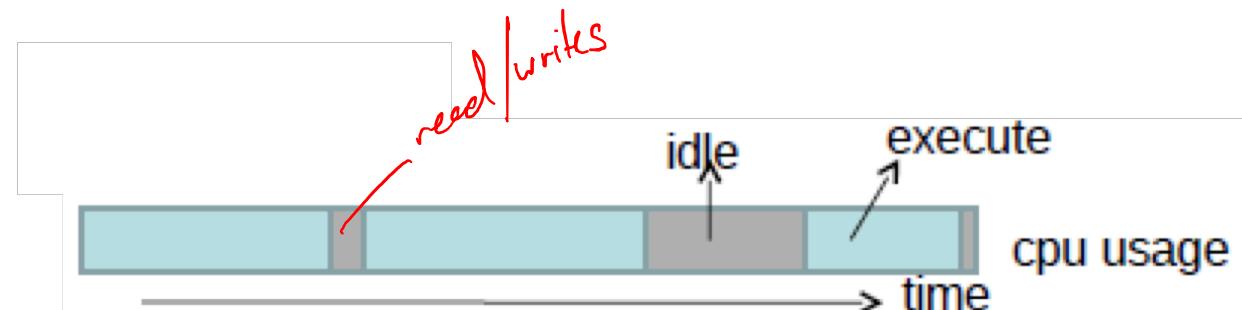
## I/O Bound

- most work is I/O-related
- small bursts of CPU activity, then waits for I/O
- a process that primarily depends (and wait on) I/O operations
- affects user interaction should run with high-priority

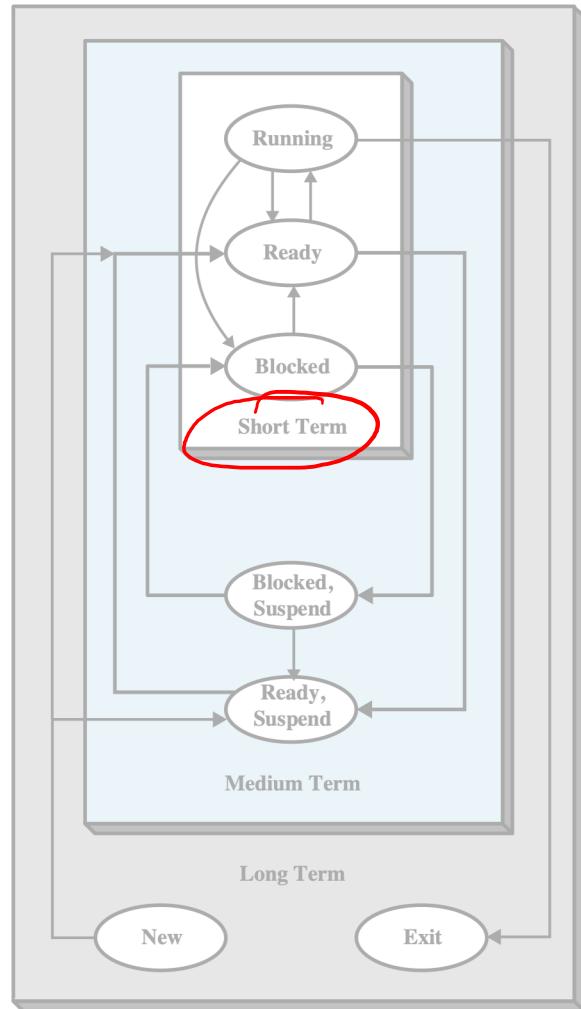


## Processor Bound

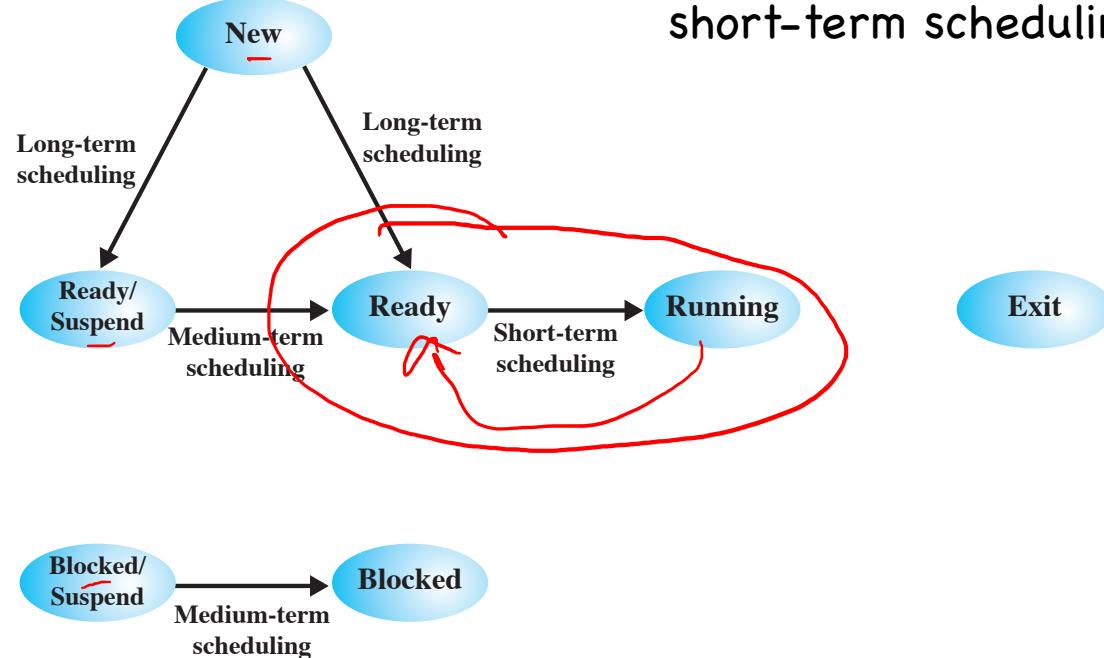
- most work is CPU-related
- hardly any I/O
- a process that primarily performs computational work
- e.g., gcc, scientific modeling, 3D rendering
- suitable for running with lower-priority



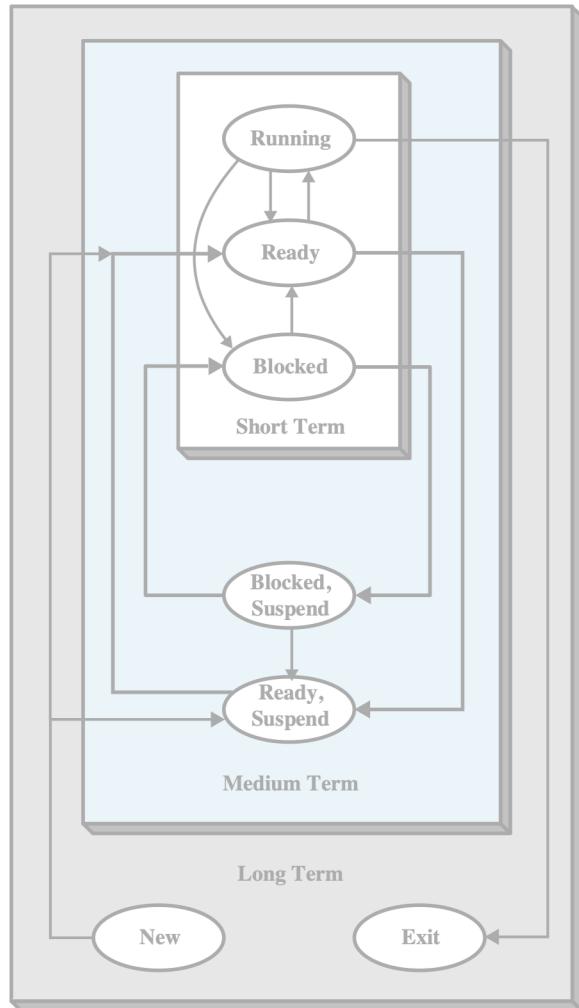
# Levels of Scheduling



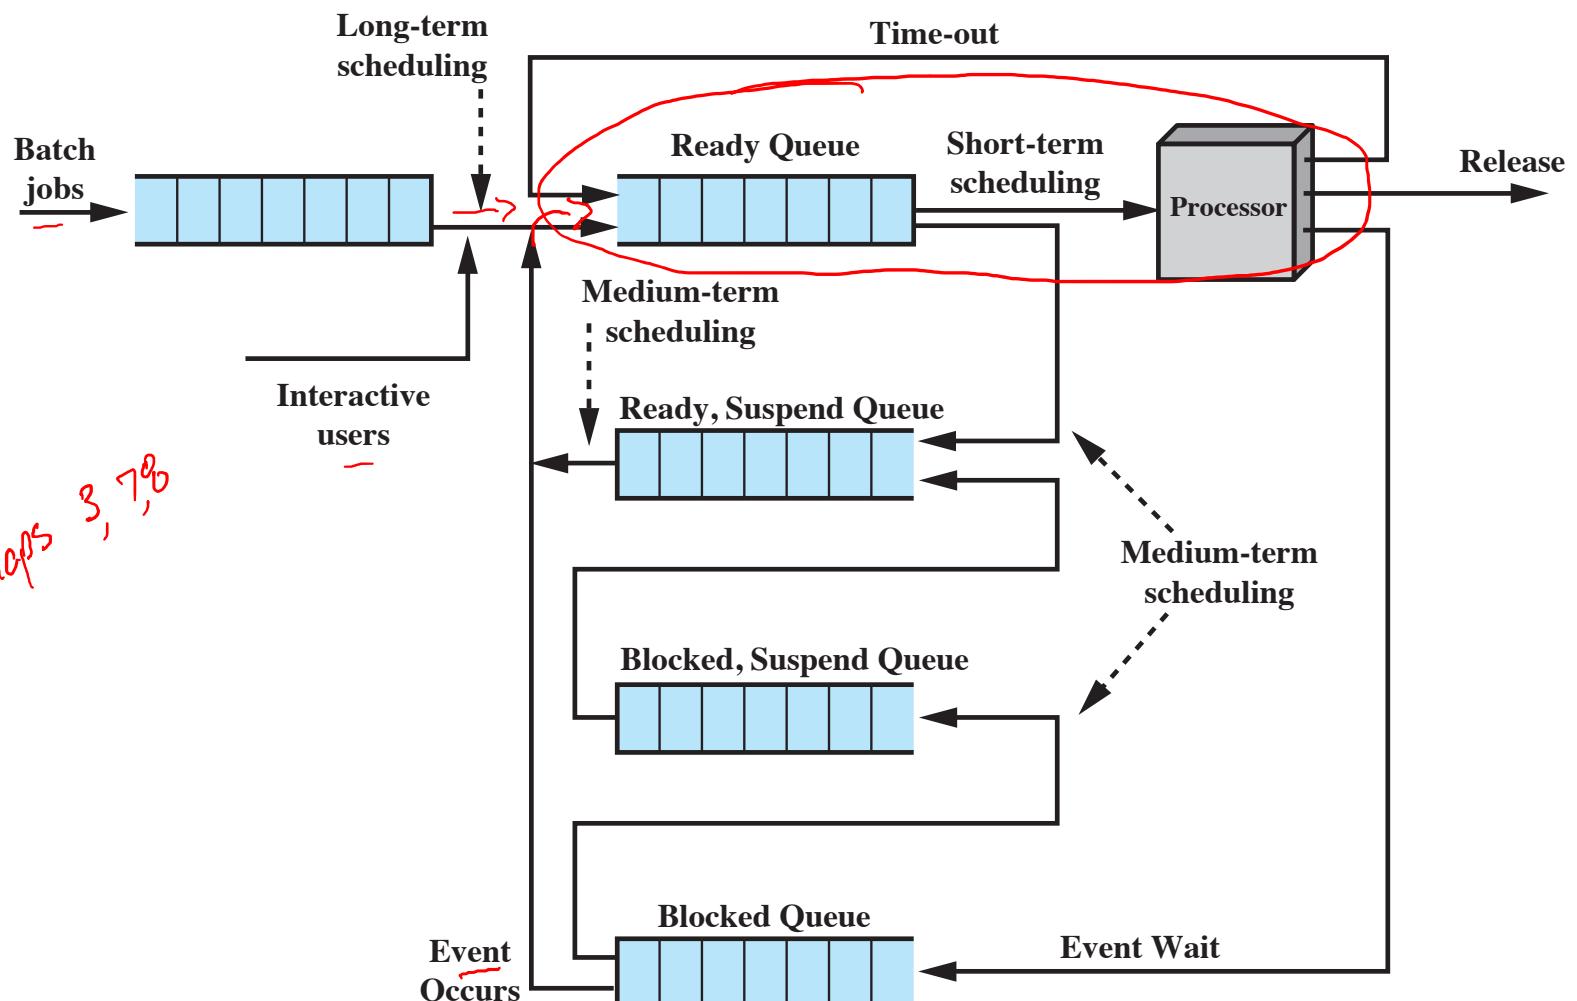
Recall process state transition diagram, now presented to illustrate long-, medium-, and short-term scheduling decisions



# Levels of Scheduling



A revised queuing diagram  
for scheduling



# Scheduling Algorithms & Criteria

- Criteria for scheduling policies
  - We'll concentrate on short-term scheduling
  - Goal: allocate processor time to optimize 1+ aspects of system behavior



- **User**-oriented criteria vs. **System**-oriented criteria



- Performance-oriented criteria vs. Non-performance-oriented criteria



# Scheduling Algorithms & Criteria (cont.)

## Summary of Scheduling Metrics

### CPU Utilization

Amount of time the CPU spends doing productive work ( $utilization = 1 - idle\_time$ )

### Throughput

# of processes that are completed per time unit (e.g., 54 procs per second)

### Turnaround Time

Amount of time from process arrival to process completion ( $TT = completion - arrival$ )

### Response Time

How quickly a process starts running after submission ( $response\_time = start - arrival$ )

## Summary of Key Scheduling Criteria

**Turnaround time** This is the interval of time between the submission of a process and its completion. Includes actual execution time plus time spent waiting for resources, including the processor. This is an appropriate measure for a batch job.

**Response time** For an interactive process, this is the time from the submission of a request until the response begins to be received. Often a process can begin producing some output to the user while continuing to process the request. Thus, this is a better measure than turnaround time from the user's point of view. The scheduling discipline should attempt to achieve low response time and to maximize the number of interactive users receiving acceptable response time.

**Deadlines** When process completion deadlines can be specified, the scheduling discipline should subordinate other goals to that of maximizing the percentage of deadlines met.

**Predictability** A given job should run in about the same amount of time and at about the same cost regardless of the load on the system. A wide variation in response time or turnaround time is distracting to users. It may signal a wide swing in system workloads or the need for system tuning to cure instabilities.

**Throughput** The scheduling policy should attempt to maximize the number of processes completed per unit of time. This is a measure of how much work is being performed. This clearly depends on the average length of a process but is also influenced by the scheduling policy, which may affect utilization.

**Processor utilization** This is the percentage of time that the processor is busy. For an expensive shared system, this is a significant criterion. In single-user systems and in some other systems, such as real-time systems, this criterion is less important than some of the others.

**Fairness** In the absence of guidance from the user or other system-supplied guidance, processes should be treated the same, and no process should suffer starvation.

**Enforcing priorities** When processes are assigned priorities, the scheduling policy should favor higher-priority processes.

**Balancing resources** The scheduling policy should keep the resources of the system busy. Processes that will underutilize stressed resources should be favored. This criterion also involves medium-term and long-term scheduling.

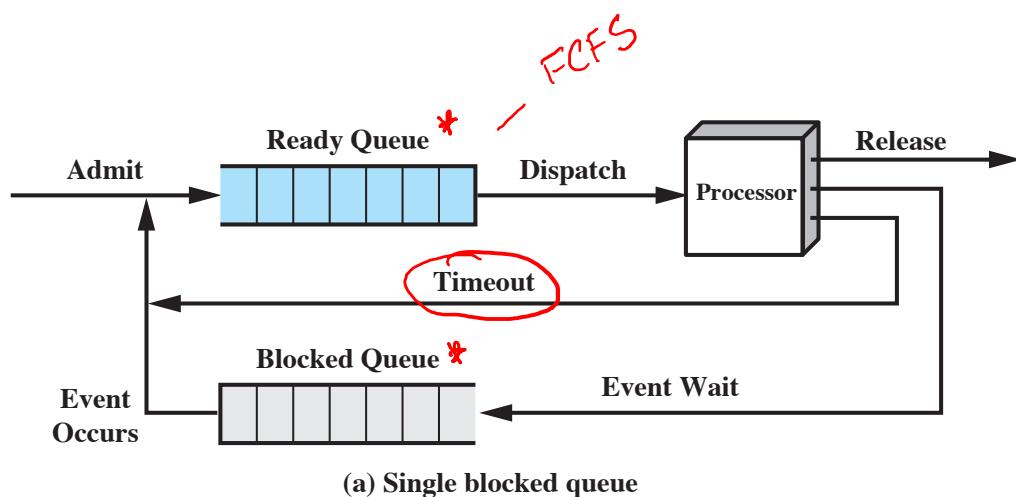
# How should the scheduler select a process to run?

# Ready. Or Ready Not.

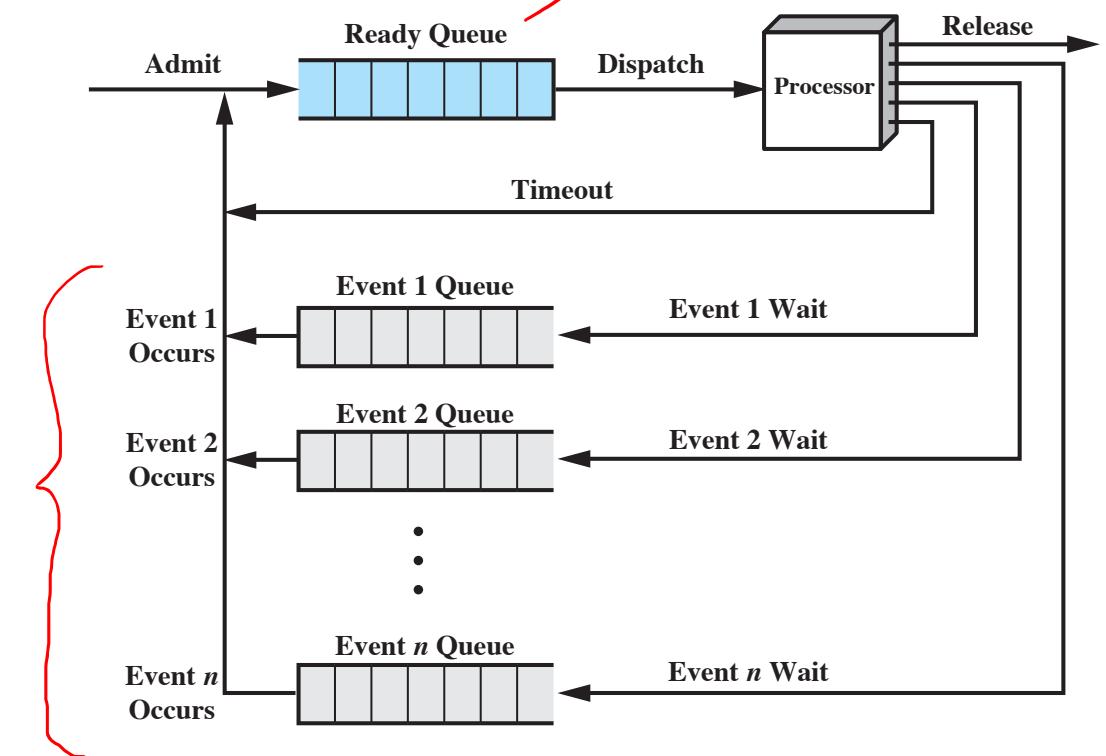


no clear way to

prioritize



(a) Single blocked queue



(b) Multiple blocked queues

# Using Priorities for Scheduling

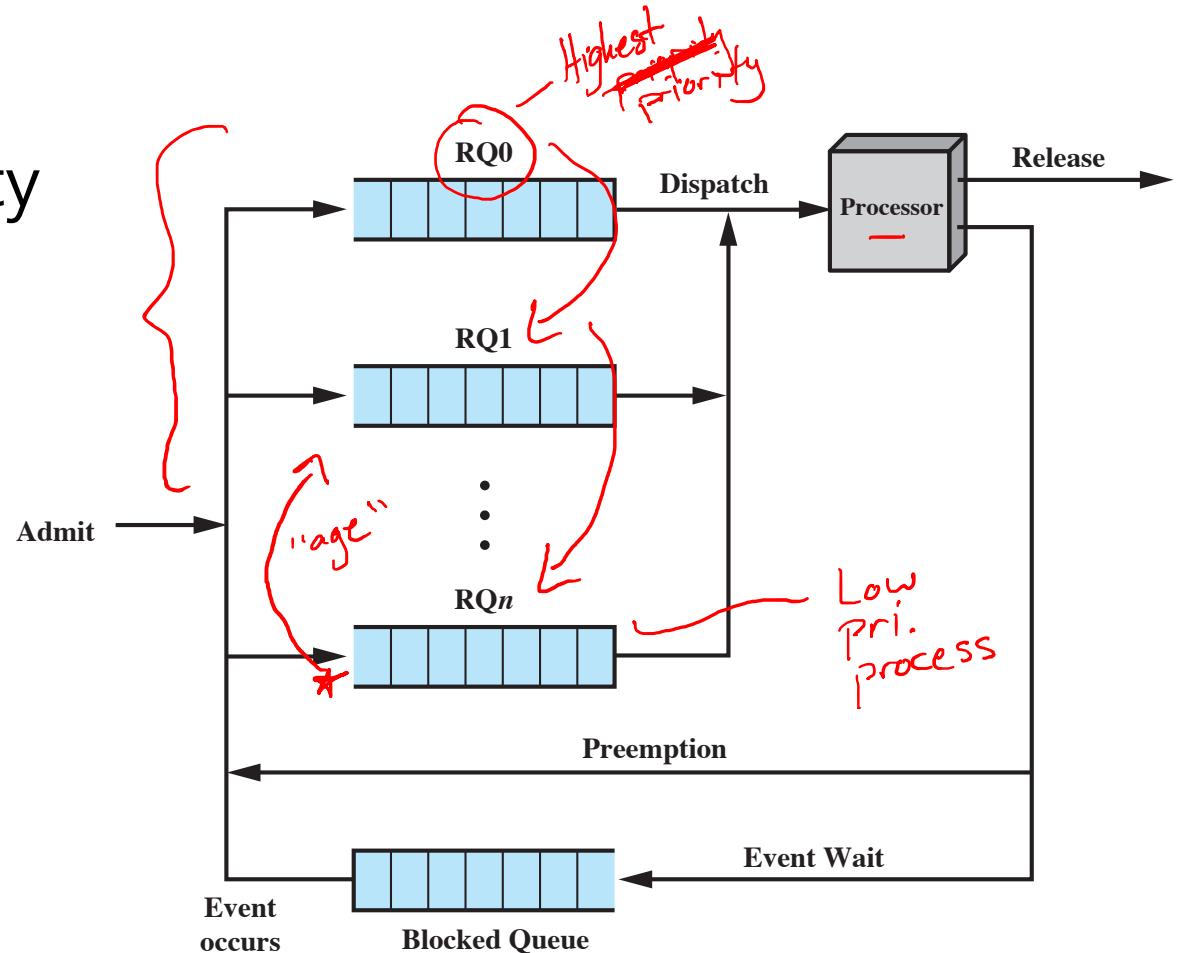
- Each process is assigned a priority

Smaller # = higher priority, Bigger # = lower priority (typically...)  
-20 - 0                                  > 0 - 20

- Scheduler selects process from higher priority queues

*Check lower priority queues successively*

- Q: Any Potential Issues?



NOTE: multiple blocked queues and suspended states are not shown here to simplify the diagram... but you'd want those too! ☺

# Comparing Scheduling Algorithms

# Characteristics of Various Scheduling Policies

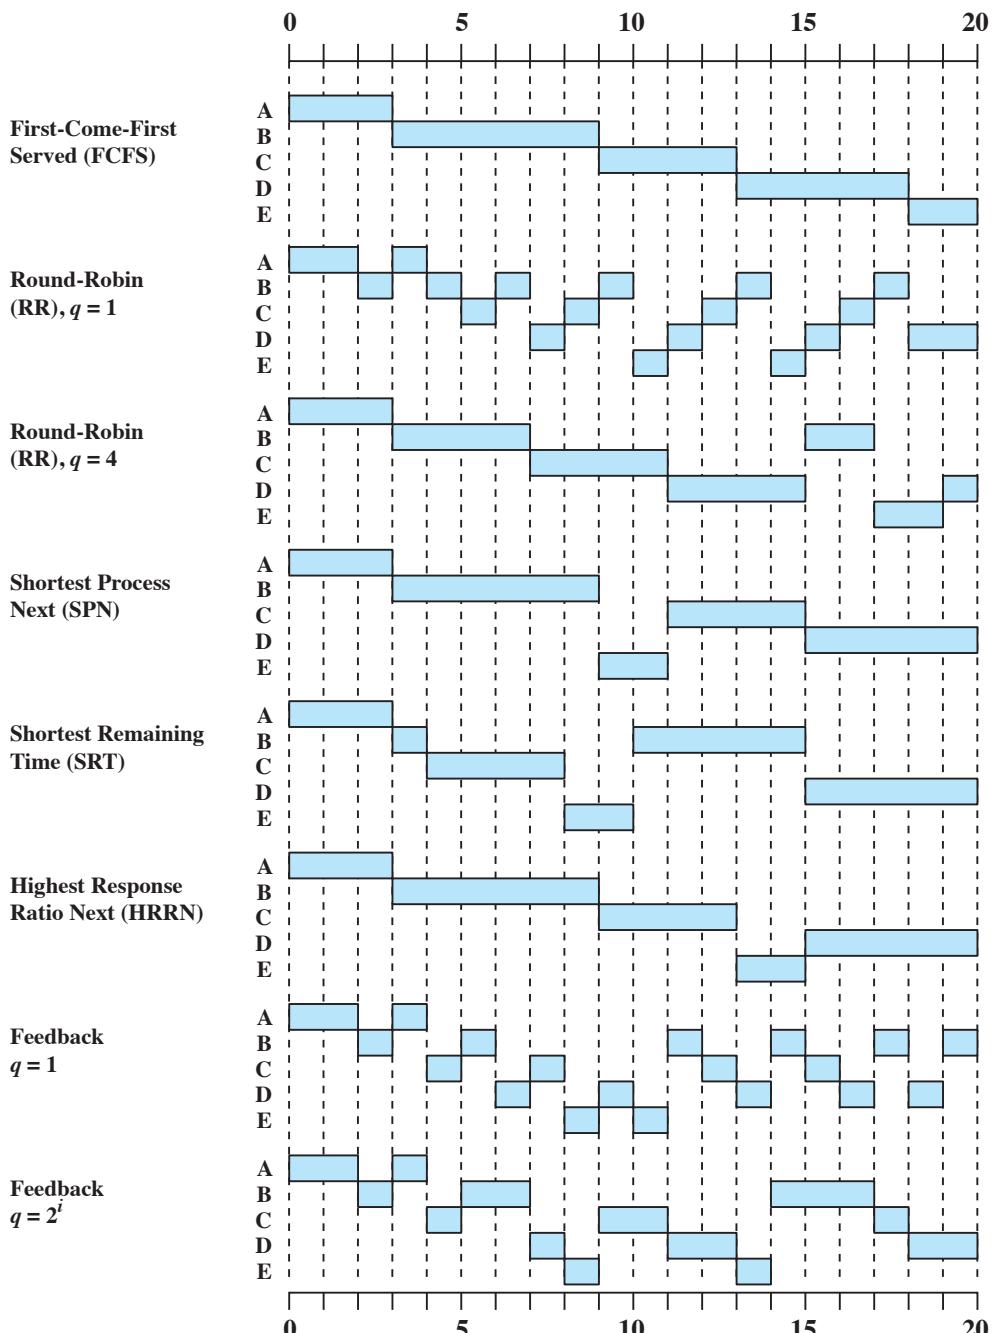
|                            | <b>FCFS</b>                                                                     | <b>Round robin</b>                              | <b>SPN</b>                                      | <b>SRT</b>                  | <b>HRRN</b>                        | <b>Feedback</b>               |
|----------------------------|---------------------------------------------------------------------------------|-------------------------------------------------|-------------------------------------------------|-----------------------------|------------------------------------|-------------------------------|
| <b>Selection Function</b>  | $\max[w]$                                                                       | constant                                        | $\min[s]$                                       | $\min[s - e]$               | $\max\left(\frac{w + s}{s}\right)$ | (see text)                    |
| <b>Decision Mode</b>       | Non-preemptive                                                                  | Preemptive (at time quantum)                    | Non-preemptive                                  | Preemptive (at arrival)     | Non-preemptive                     | Preemptive (at time quantum)  |
| <b>Throughput</b>          | Not emphasized                                                                  | May be low if quantum is too small              | High                                            | High                        | High                               | Not emphasized                |
| <b>Response Time</b>       | May be high, especially if there is a large variance in process execution times | Provides good response time for short processes | Provides good response time for short processes | Provides good response time | Provides good response time        | Not emphasized                |
| <b>Overhead</b>            | Minimum                                                                         | Minimum                                         | Can be high                                     | Can be high                 | Can be high                        | Can be high                   |
| <b>Effect on Processes</b> | Penalizes short processes; penalizes I/O bound processes                        | Fair treatment                                  | Penalizes long processes                        | Penalizes long processes    | Good balance                       | May favor I/O bound processes |
| <b>Starvation</b>          | No                                                                              | No                                              | Possible                                        | Possible                    | No                                 | Possible                      |

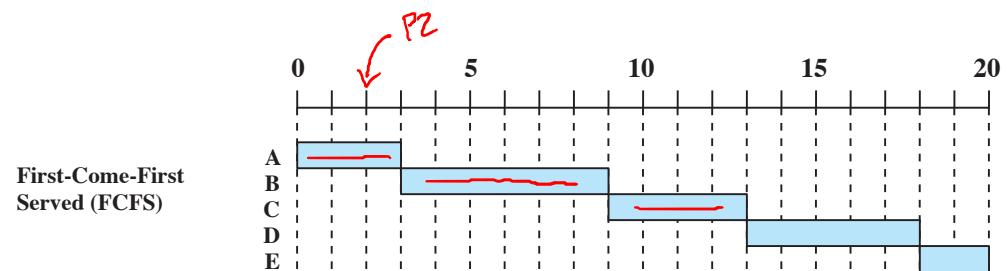
$w$  = time spent WAITING,  $e$  = time spent EXECUTING,  $s$  = TOTAL service time

# Comparing Scheduling Algorithms

## An Example

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| A       | 0            | 3            |
| B       | 2            | 6            |
| C       | 4            | 4            |
| D       | 6            | 5            |
| E       | 8            | 2            |





| Process                   | A    | B    | C    | D    | E    |      |
|---------------------------|------|------|------|------|------|------|
| Arrival Time              | 0    | 2    | 4    | 6    | 8    |      |
| Service Time ( $T_s$ )    | 3    | 6    | 4    | 5    | 2    | Mean |
| FCFS                      |      |      |      |      |      |      |
| Finish Time               | 3    | 9    | 13   | 18   | 20   |      |
| Turnaround Time ( $T_r$ ) | 3    | 7    | 9    | 12   | 12   | 8.60 |
| $T_r/T_s$                 | 1.00 | 1.17 | 2.25 | 2.40 | 6.00 | 2.56 |

# First-Come-First-Served (FCFS)

- Strict queueing scheme (a.k.a. FIFO)
- The process that has been in the ready queue the longest is selected next
- Q: What happens when a short process arrives just after a long process?
- Not a great solution on its own for a uniprocessor system.  
Q: Why?

