

Operating Systems!

Scheduling: **Uniprocessor Scheduling** **(Part 2)**

Prof. Travis Peters

Montana State University

CS 460 - Operating Systems

Fall 2020

<https://www.cs.montana.edu/cs460>

Some diagrams and notes used in this slide deck have been adapted from Sean Smith's OS courses @ Dartmouth. Thanks, Sean!

Today

- Announcements
 - ~~Yalnix~~ How are projects/proposals going?! 😊
- Upcoming Deadlines
 - **Project Proposal (HARD DEADLINE)**
Sunday [10/25/2020] @ 11:59 PM (MST)

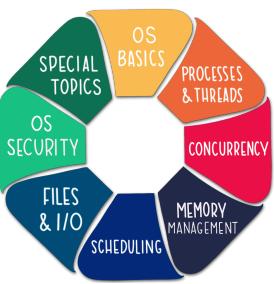


Yalnix Teams:

You should plan to be near completing checkpoint 3 at this point!

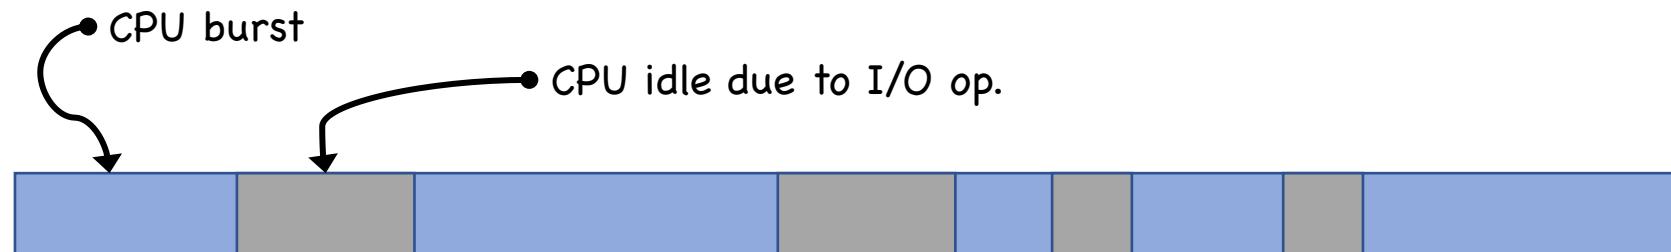
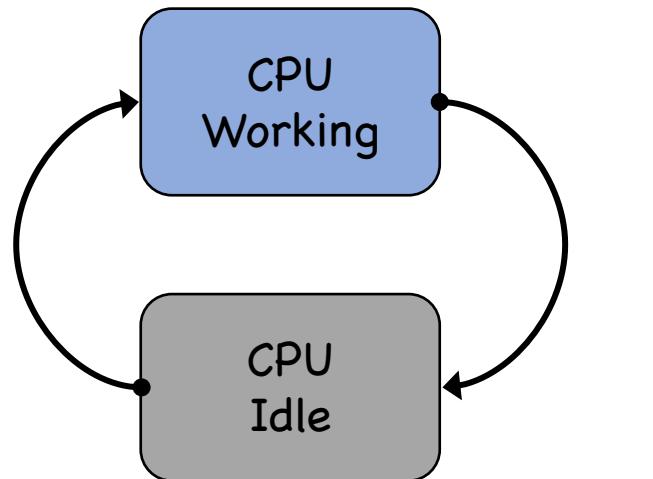
Non-Yalnix Teams:

Let me know if you want to discuss projects before submitting your proposal!



Today (cont.)

- Agenda
 - ~~Types of Processes~~
 - ~~Long-, medium-, and short-term~~ scheduling
 - Performance of different scheduling policies

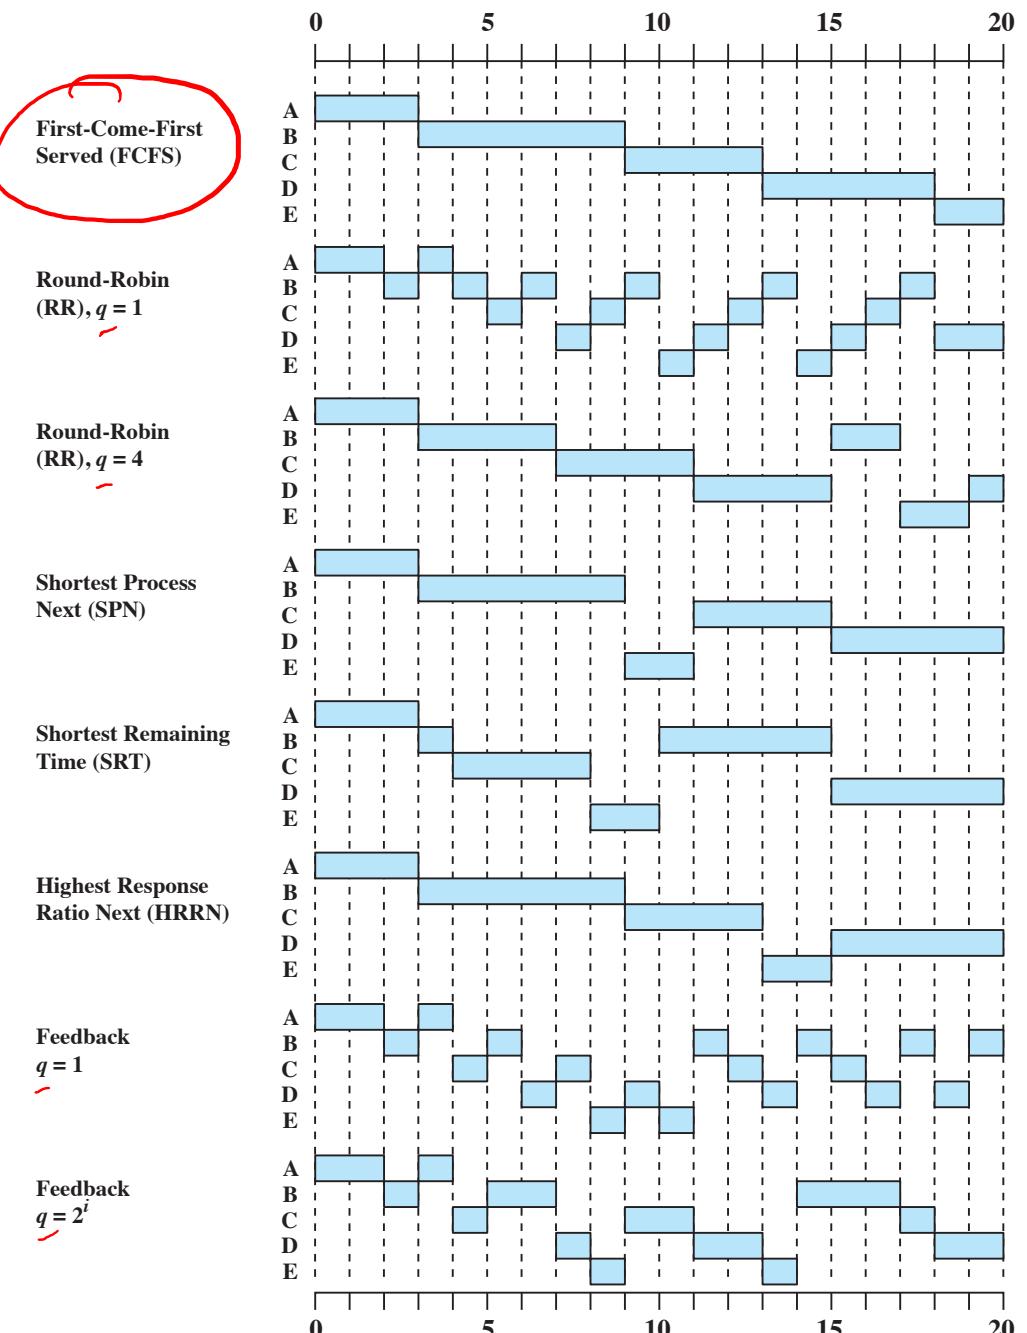


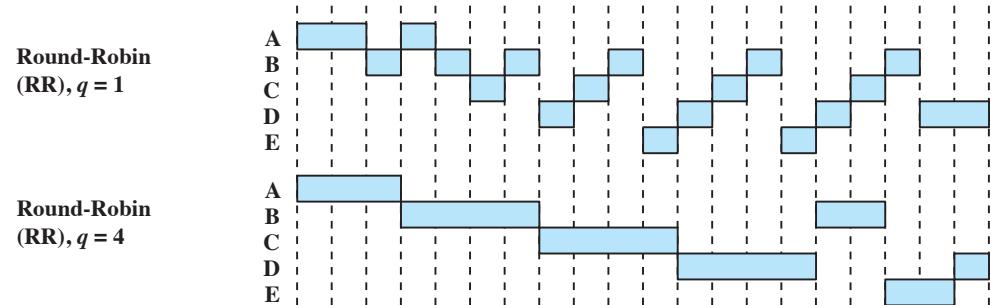
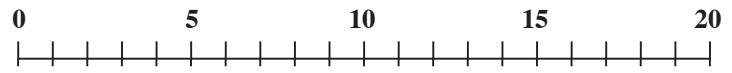
Adapted from https://www.cse.iitm.ac.in/~chester/courses/15o_os/slides/8_Scheduling.pdf

Comparing Scheduling Algorithms

An Example

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2





Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time (T_s)	3	6	4	5	2	Mean

Round Robin (RR)

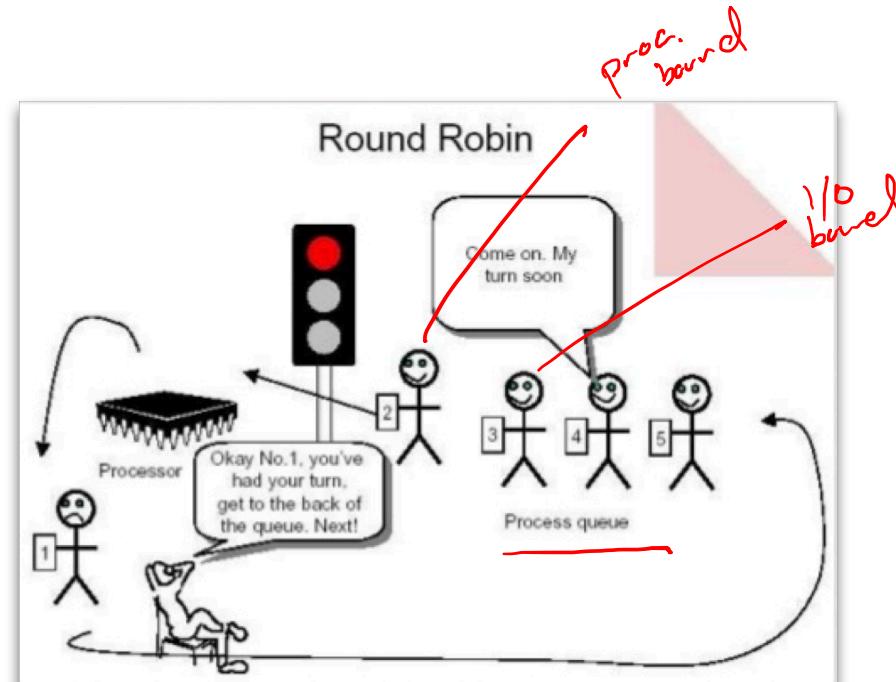
RR $q = 1$

Finish Time	4	18	17	20	15	
Turnaround Time (T_r)	4	16	13	14	7	10.80
T_r/T_s	1.33	2.67	3.25	2.80	3.50	2.71

RR $q = 4$

Finish Time	3	17	11	20	19	
Turnaround Time (T_r)	3	15	7	14	11	10.00
T_r/T_s	1.00	2.5	1.75	2.80	5.50	2.71

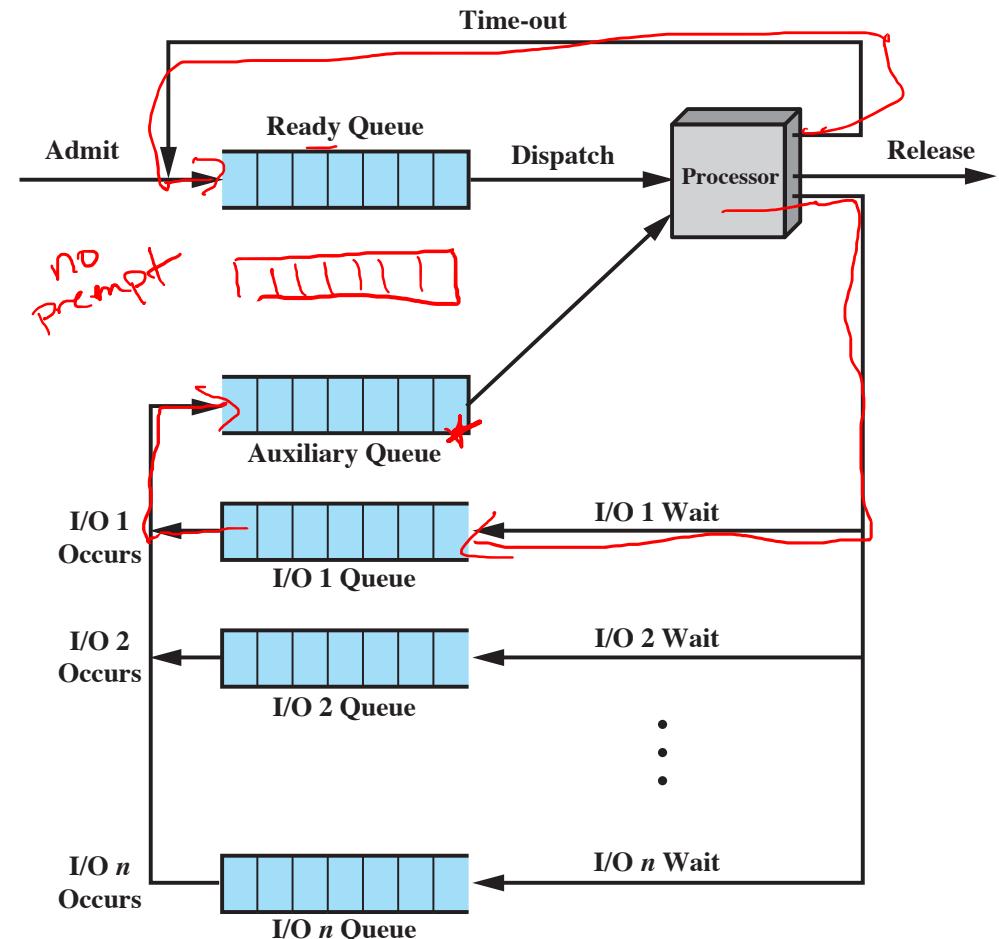
- Use preemption based on a clock!
 - clock-based interrupt (periodic intervals of execution)
 - preempted process goes back on the ready queue
 - “Time Slicing”
- Principle design choice: length of time quanta (slice)
 - Q: Should we prefer small slices or big slices?
- Q: Any drawbacks?

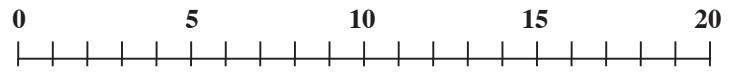


Refinement: *Virtual Round Robin (VRR)*

Same as RR... plus...

- Each process gets the opportunity to complete its time quanta
- Use **Auxiliary Queue (AQ)** to hold recently unblocked I/O-bound processes
- Processes in AQ get preference over processes in normal **Ready Queue (RQ)**

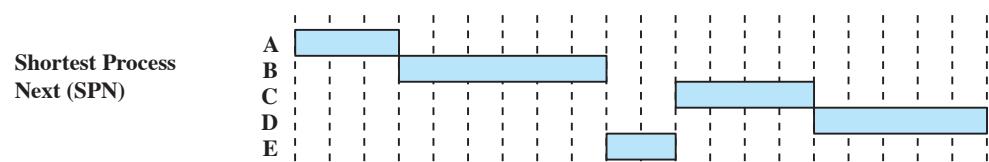




Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time (T_s)	3	6	4	5	2	Mean

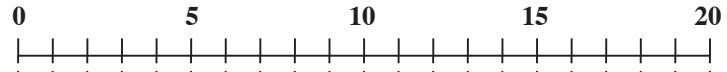
Shortest Process Next (SPN)

- No preemption allowed \rightarrow processes run to completion
- Process w/ shortest (expected) processing time is selected next \rightarrow short processes can "cut the line"
- Need to know (or estimate) the required processing time \rightarrow How?



SPN						
Finish Time	3	9	15	20	11	
Turnaround Time (T_r)	3	7	11	14	3	7.60
T_r/T_s	1.00	1.17	2.75	2.80	1.50	1.84

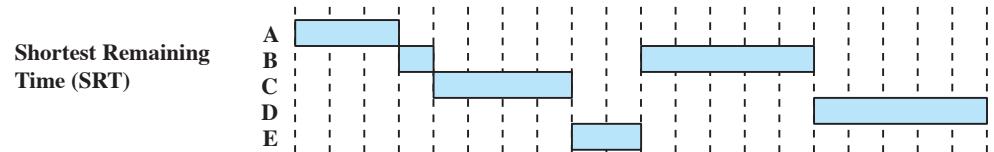
- Q: Any issues with this approach?
Longer proc could starve
- Q: Any undesirable characteristics?
Escalate of proc priority have been waiting a long time



Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time (T_s)	3	6	4	5	2	Mean

Shortest Remaining Time (SRT)

- Preemption allowed!
 - > processes can be interrupted
- Process w/ shortest (expected) **remaining** processing time is selected next
 - > processes w/ less remaining time can "cut the line" and preempt those w/ longer remaining times

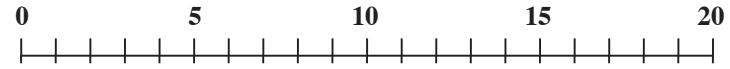


SRT						
Finish Time	3	15	8	20	10	
Turnaround Time (Tr)	3	13	4	14	2	7.20
Tr/T_s	1.00	2.17	1.00	2.80	1.00	1.59

- long procs could starve

switch proc if proc arrives

- Q: Advantages? Disadvantages?



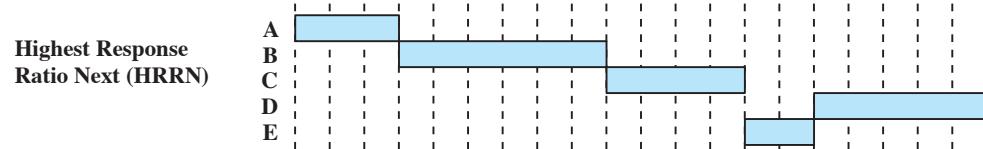
Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time (T_s)	3	6	4	5	2	
						Mean

Highest Response Ratio Next (HRRN)

- When the current process blocks
-> choose process w/ the greatest **normalized turnaround time (R)**

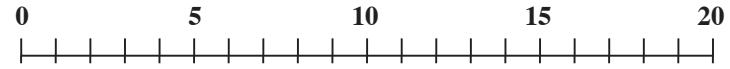
$$R = \frac{w + s}{s}$$

R = response ratio, w = time spent waiting, s = expected service time



HRRN						
Finish Time	3	9	13	20	15	
Turnaround Time (Tr)	3	7	9	14	7	8.00
Tr/T_s	1.00	1.17	2.25	2.80	3.5	2.14

- Accounts for the age of the process!
-> shorter jobs are favored...
-> ... but aging process without service increases the ratio,
so a longer process will eventually get past competing shorter jobs.



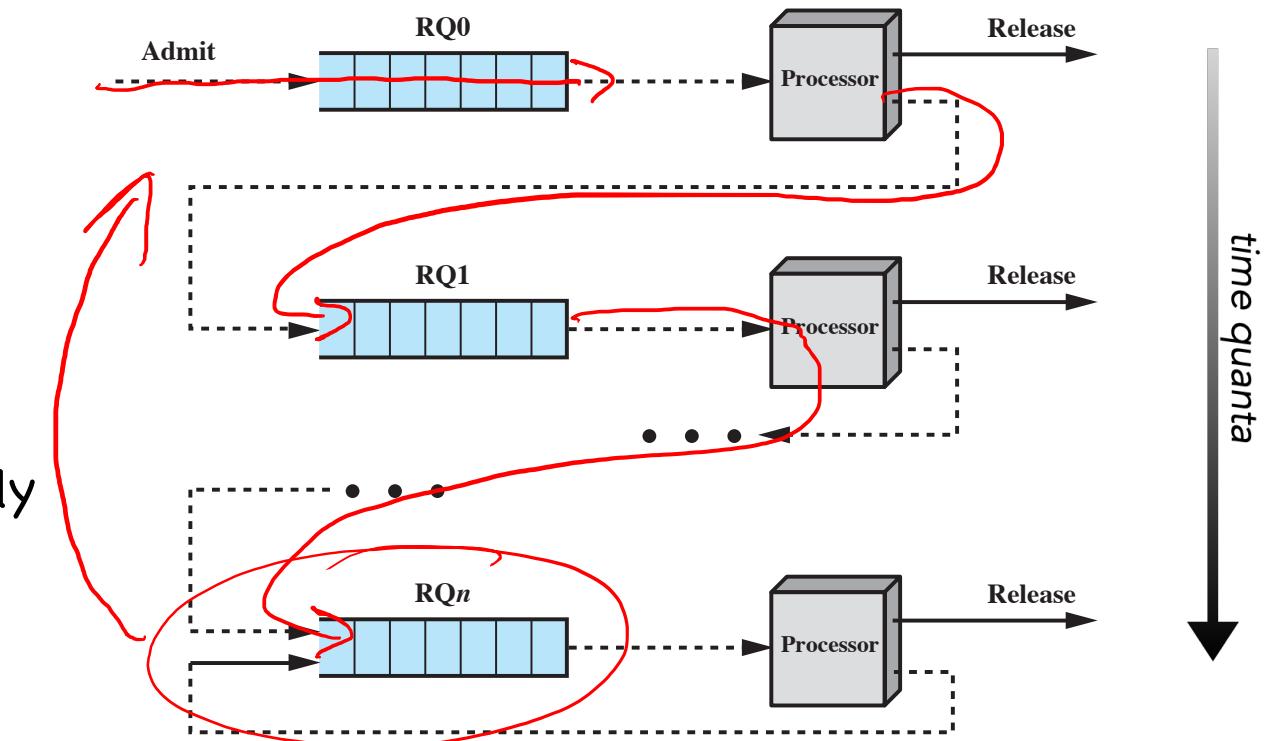
Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time (T_s)	3	6	4	5	2	Mean

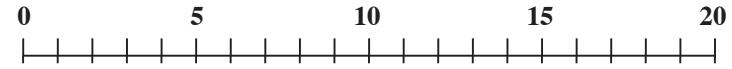
Multilevel Feedback

- SPN, SRT, HRRN are tricky
-> *need reasonable estimates for execution time*
- What if we instead “penalize” jobs that have been running longer?
- If we can’t get easy, reliable measures of
time remaining to execute
let us focus on
time spent executing so far!

Multilevel Feedback

- Use some time quanta (re: RR/VRR) to periodically interrupt a running process
- “Demote” process to lower PQs each time it returns to the RQ
- Shortest processes still finish relatively quickly
- Use FCFS within each Q (Except lowest PQ, which uses RR)

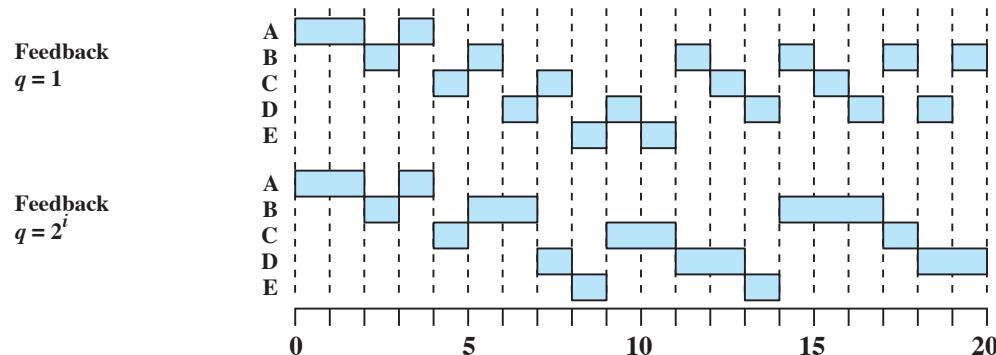




Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time (T_s)	3	6	4	5	2	
						Mean

Multilevel Feedback

- SPN, SRT, HRRN are tricky
-> need reasonable estimates for execution time
- What if we instead “penalize” jobs that have been running longer?
- If we can’t get easy, reliable measures of
time remaining to execute
let us focus on
time spent executing so far!



FB $q = 1$						
Finish Time	4	20	16	19	11	
Turnaround Time (Tr)	4	18	12	13	3	10.00
Tr/T_s	1.33	3.00	3.00	2.60	1.5	2.29
FB $q = 2^i$						
Finish Time	4	17	18	20	14	
Turnaround Time (Tr)	4	15	14	14	6	10.60
Tr/T_s	1.33	2.50	3.50	2.80	3.00	2.63

Fair-Share Scheduling

- Approaches so far all treat processes as a fairly homogeneous collection of ready-to-execute items
 - Many independent processes
 - Can be sub-divided by priority
 - Otherwise, treated the same
- Q: Can we see a potential issue with this?
 - What if there are 2 users, Alice and Bob, where each run 1 “app”
 - ... and Alice’s app is made up of 1 process
 - ... and Bob’s app is made up of 100 processes
- Fair-Share Scheduling
 - assign each user/group a “share” of the processor
 - monitor resource usage...
 - give **fewer resources** to user/group that has had **more than their fair share**
 - give **more resources** to user/group that has had **less than their fair share**

NOTE:
“fair-share”
doesn't necessarily mean
“equal share”

So...
which is best?

Performance Comparisons & General Take-Aways

- Difficult to develop closed analytic models for comparisons*

All approaches (except RR and FCFS) depend on expected service times, which really complicates things

- Relative performance depends on a variety of factors

- **probability distribution** of service times of processes
- **efficiency** of scheduling mechanisms
- **efficiency** of context switching mechanisms
- **demand on**, and efficiency of, I/O subsystem(s)

→ See textbook and slides on “Queueing Analysis” and “Simulation Modeling”

Summary

Big Ideas

- Scheduling granularity
-> long-, medium-, and short-term*
- In scheduling, there are many objectives to balance
-> e.g., process type, priority, (normalized) response time, no starvation.
- Various algorithms have been developed to address short-term scheduling
-> FCFS, RR/VRR, SPN, SRT, HRRN, MLFQ, ...

What next?

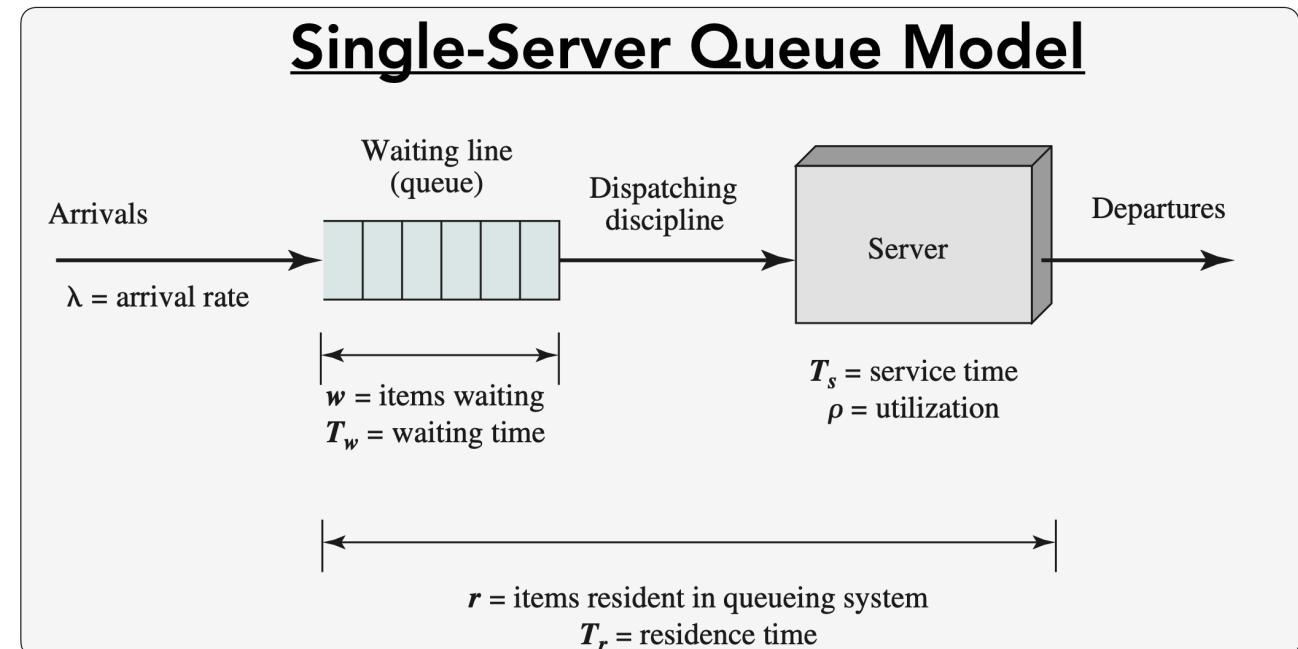
- Approaches to performance analysis and comparisons
- Extensions into multiprocessor and multicore scheduling
- Real-time scheduling
- Linux scheduling & the Completely Fair Scheduler (CFS)
- *and so much more!*

What Next?

Topics to explore when going beyond the basics of scheduling

Approach: Queuing Analysis

- Idea: Use analytic models
- Assume: Poisson arrivals
→ **random arrival times**
- Assume: exponential service times
→ **more demand == longer service times**
- Assume: priority-based scheduling...
→ **higher priority items served first**
→ **FCFS for items of equal priority**



Extensions also exist for
multiple servers and multiple queues

Approach: Simulation Modeling

- Pros:** discrete event simulations allow a variety of models to be simulated & compared
 - Able to visualize direct comparisons between policies
 - Compare wait times, normalized turnaround times, ...
- Cons:** results only really hold for a given “run”
 - a particular collection of processes...
 - a particular set of assumptions & constraints...

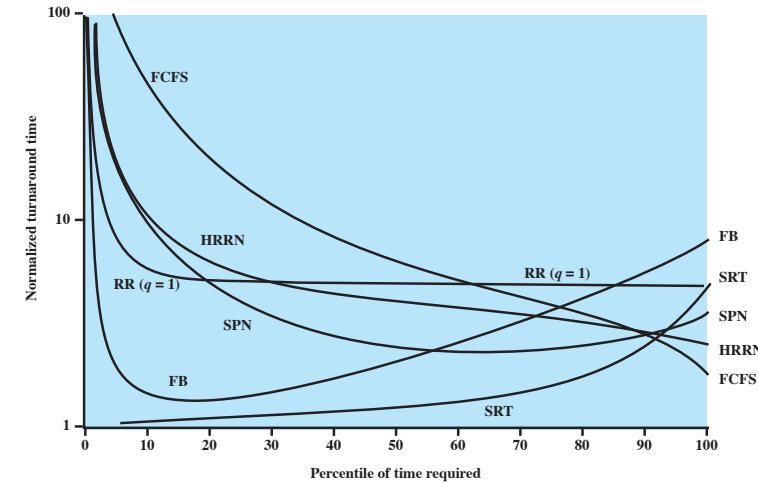


Figure 9.14 Simulation Results for Normalized Turnaround Time

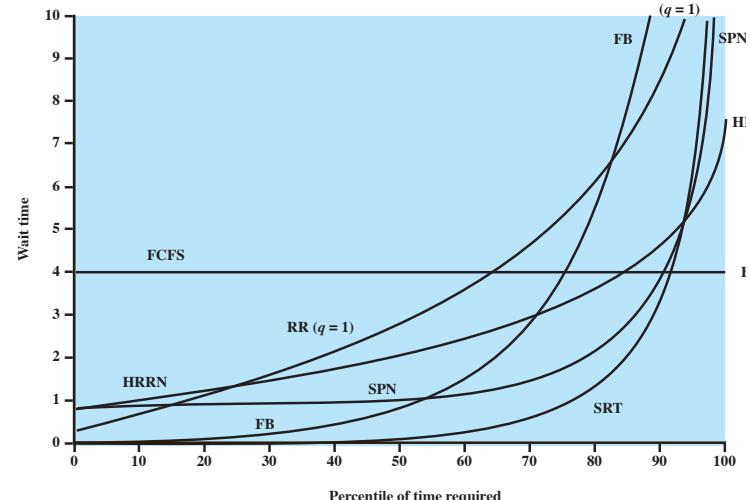


Figure 9.15 Simulation Results for Waiting Time

Linux Scheduling

- Scheduling Classes
 - **SCHED_FIFO** (limited preemption) & **SCHED_RR** (time-sliced) – (0-99)
 - **SCHED_NORMAL** (100-139)
 - > lower value == higher priority
 - > non-RT thread only executes if no RT threads are ready to execute

- Linux 2.6+ → **O(1) Scheduler**
 - Named so because time to select & run a process takes constant time.
 - Complex and not good to run in the kernel....

- Linux 2.6.23+ → **Completely Fair Scheduler (CFS)**
 - Maintain *virtual runtime* value for each task (normalized amt. of time spent executing so far)
 - *Sleeper Fairness* (i.e., ensure that processes that wait on I/O get fair access to processor)
 - Use Red-Black Tree to order runnable tasks (efficient inserts/deletes, and searches)

