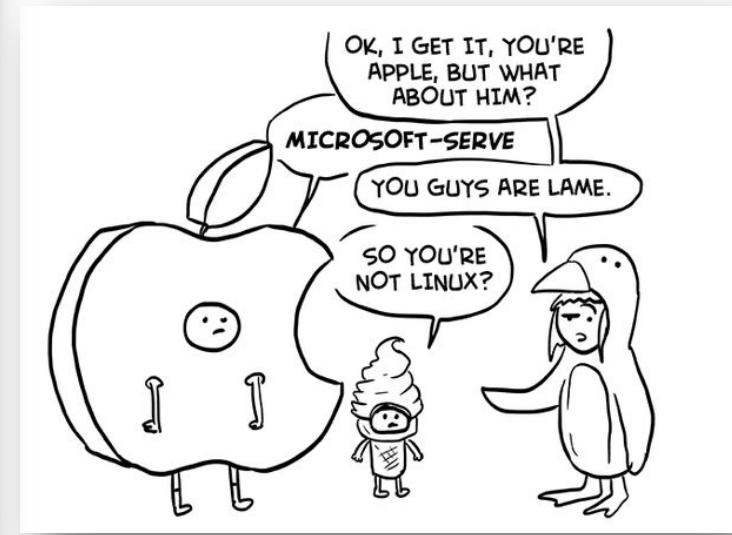
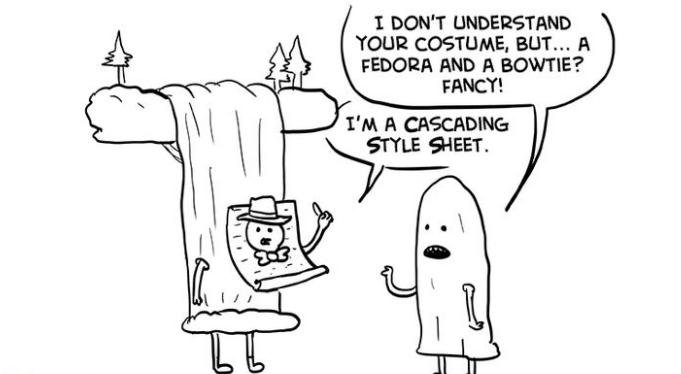
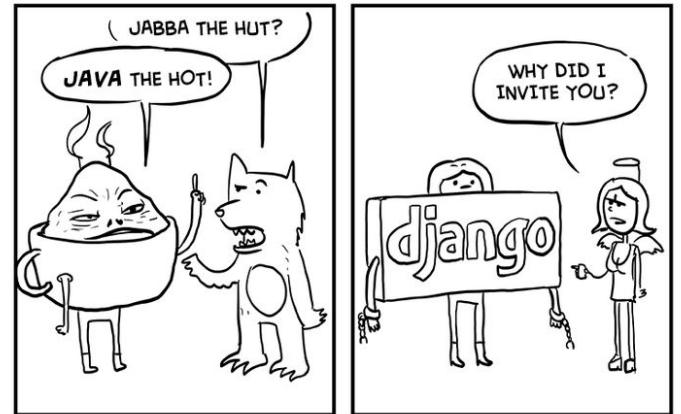
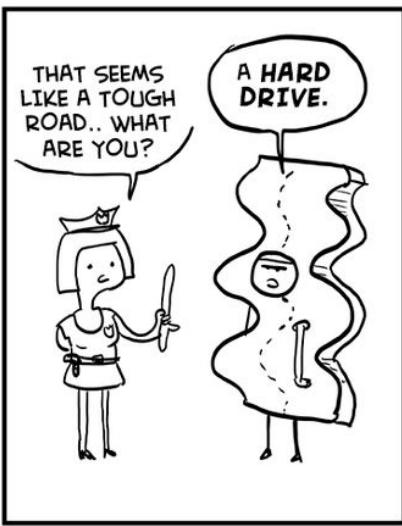
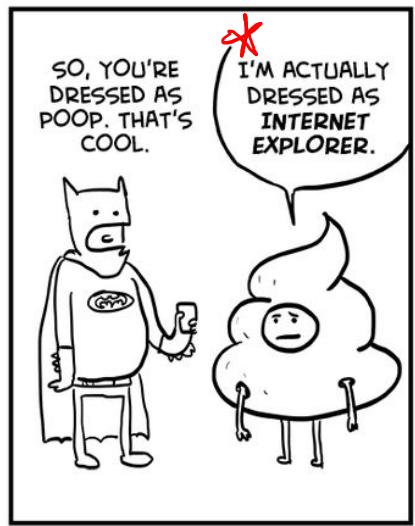
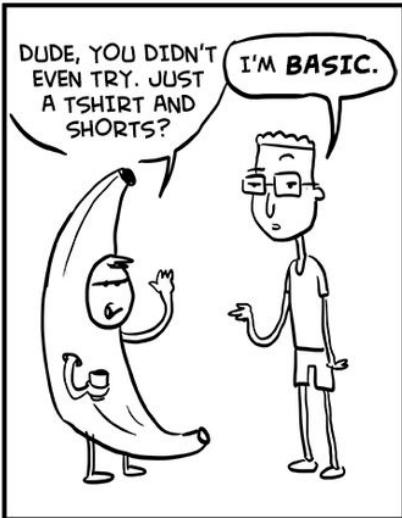


Halloween Costumes for Computer Scientists



Operating Systems!

Input/Output (I/O): **Files & File Systems**

Prof. Travis Peters

Montana State University

CS 460 - Operating Systems

Fall 2020

<https://www.cs.montana.edu/cs460>

Some diagrams and notes used in this slide deck have been adapted from Sean Smith's OS courses @ Dartmouth. Thanks, Sean!

Today

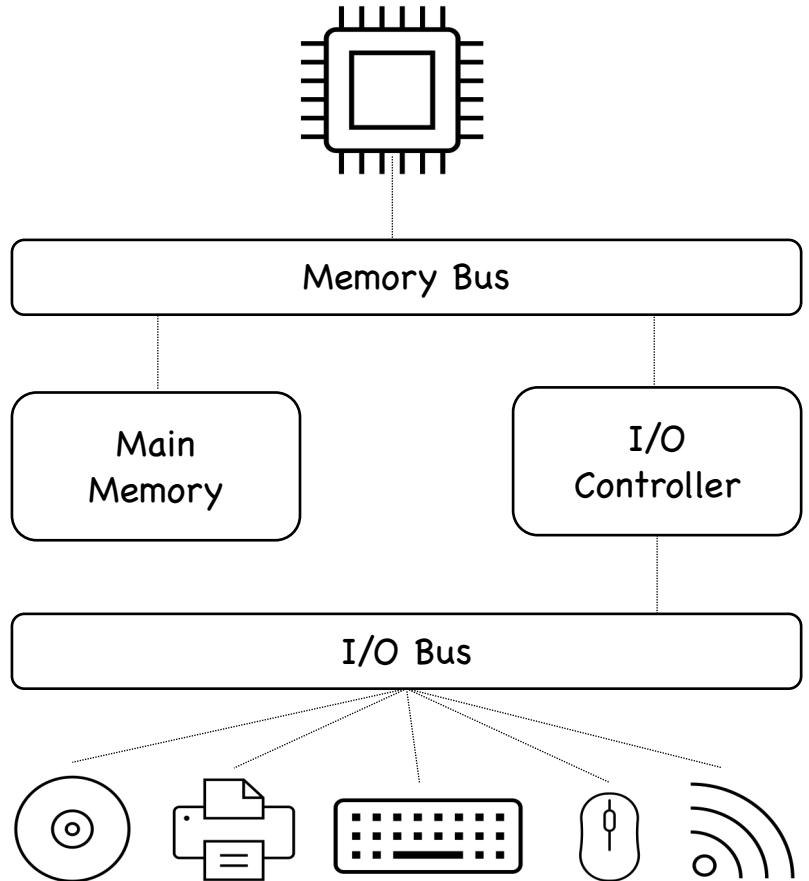
- Announcements
 - Almost done grading PA3! ☺
 - Exam 2 details coming soon (11/20? 11/23?)
- Upcoming Deadlines
 - **Project Submission (HARD DEADLINE)**
Sunday [11/15/2020] @ 11:59 PM (MST)
 - **Project Evaluations (HARD DEADLINE)**
Wednesday [11/18/2020] @ 11:59 PM (MST)
 - **Exam #2**
TBD





Today

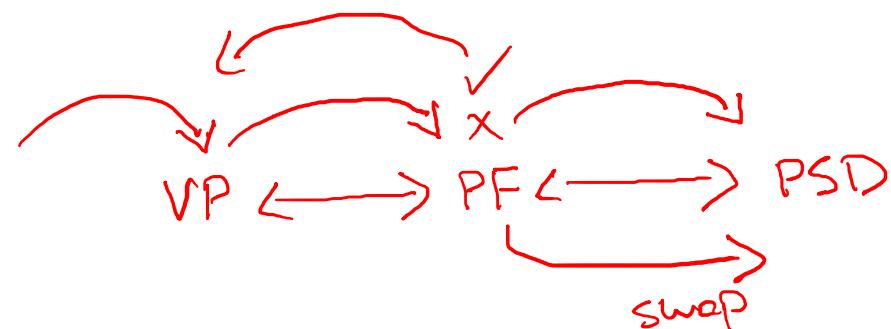
- Agenda
 - Key concepts behind I/O
 - ~~What is I/O? I/O Categories. Memory access.~~
 - ~~What is a file? What is a file system? etc.~~
 - Secondary Storage
 - ~~Emphasis on disks and disk scheduling~~
 - Files
 - Operations
 - Implementations
 - File System
 - ~~Virtual File Systems~~
 - Implementations
 - Security



Reflection

- What is a file? persistent data (stored on disk)
a place to put stuff

- What is a file system?
 $/tmp$ $tmpfs$
enables retrieval (operations)
an organization structure
- How do these concepts differ from the other storage concepts we've been discussing?
E.g.,
 - physical memory?
 - virtual memory?
 - backing store for nonresident pages?



Files

Files & File Metadata

- A place to store stuff!
(e.g., bytes)
 - Try: **xxd**
- Typical metadata attributes
 - name, size, identifier, location, protection, bookkeeping, type, etc.
 - Try: **stat**, **man 2 stat**

JULIA EVANS
@b0rk

file descriptors

Unix systems use integers to track open files

process → Open foo.txt → Okay! that's file #7 for you.

these integers are called **file descriptors**

lsof (list open files) will show you a process's open files

\$ lsof -p 4242 ← PID we're interested in

FD	NAME
0	/dev/pts/tty1
1	/dev/pts/tty1
2	pipe:29174
3	/home/b0rk/awesome.txt
5	/tmp/

↑ FD is for file descriptor

file descriptors can refer to:

- files on disk
- pipes
- sockets (network connections)
- terminals (like xterm)
- devices (your speaker! /dev/null!)
- LOTS MORE (eventfd, inotify, signalfd, epoll, etc etc)

not EVERYTHING on Unix is a file, but lots of things are

When you read or write to a file/pipe/network connection you do that using a file descriptor

connect to google.com → write GET / HTTP/1.1 to fd #5 → OS → ok! fd is 5!

Behind the scenes:

open file.txt → Python program → read from file #4 → OS → here are the contents!

ok! fd is 4

Let's see how some simple Python code works under the hood:

Python:
f = open("file.txt")
f.readlines()

"read from std in" means "read from the file descriptor 0"
could be a pipe or file or terminal

Files & File Types

- How do we know what is IN a file?

- How can we know the file type?

- Any issues?



Issues with Binding Applications/File Extensions

- File Extension Attacks:

- Spoofed file types

<https://cve.mitre.org/cgi-bin/cvename.cgi?name=CVE-2017-14604>

- Games w/ parsing file extensions

<https://cve.mitre.org/cgi-bin/cvename.cgi?name=cve-2009-4444>

- Visual truncation

<https://cve.mitre.org/cgi-bin/cvename.cgi?name=CAN-2004-2530>

File Systems

File System Operations

- *What do we do with files?*
 - create
 - read
 - write
 - delete
- To do these things, we often need information...
 - sequential access *current pos + len / # bytes*
vs.
 - direct/random access *Adelr*

File System Operations

- *What do we do with files?*
 - create
 - read
 - write
 - delete
- To do these things, we often need information...
 - sequential access → CURRENT LOCATION
vs.
• direct/random access → need ADDRESS of specific thing to access

File Systems & Bookkeeping

- We need to store this information somewhere...
 - Where?
 - When does it get set/reset?

File Systems & Bookkeeping

- We need to store this information somewhere...
 - Where?
 - When does it get set/reset?
- On open/close
→ **open file table!**
- I/O calls refer to this table
 - efficient references
 - easier sharing semantics
 - easier access control

JULIA EVANS
@b0rk

file descriptors

Unix systems use integers to track open files

process → Open foo.txt → PID we're interested in

Okay! that's file #7 for you.

these integers are called **file descriptors**

lsof (list open files) will show you a process's open files

FD	NAME
0	/dev/pts/tty1
1	/dev/pts/tty1
2	pipe:29174
3	/home/b0rk/awesome.txt
5	/tmp/
6	---
7	file.txt

FD is for file descriptor

file descriptors can refer to:

- files on disk
- pipes
- sockets (network connections)
- terminals (like xterm)
- devices (your speaker! /dev/null!)
- LOTS MORE (eventfd, inotify, signalfd, epoll, etc etc)

not EVERYTHING on Unix is a file, but lots of things are

When you read or write to a file/pipe/network connection you do that using a file descriptor

connect to google.com → write GET / HTTP/1.1 to fd #5 → done!

ok! fd is 5! → OS → Python program → read from file #4 → here are the contents!

Let's see how some simple Python code works under the hood:

Python:
`f = open("file.txt")
f.readlines()`

Behind the scenes:

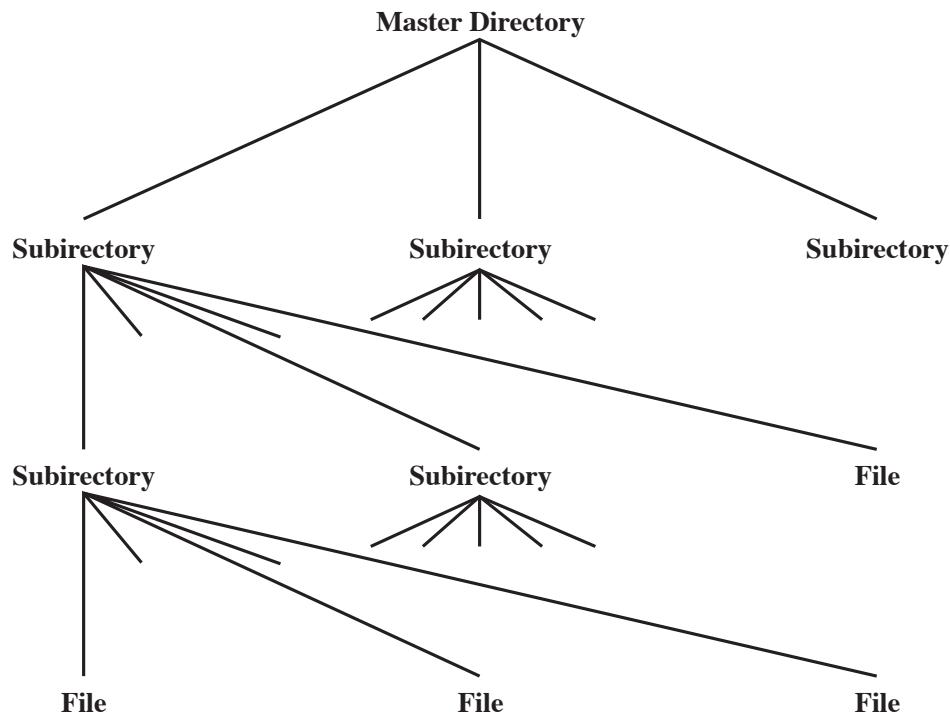
open file.txt → ok! fd is 4 → OS → here are the contents!

(almost) every process has 3 standard FDs

stdin → 0
stdout → 1
stderr → 2

"read from std in" means "read from the file descriptor 0"
could be a pipe or file or terminal

Bookkeeping (detailed)

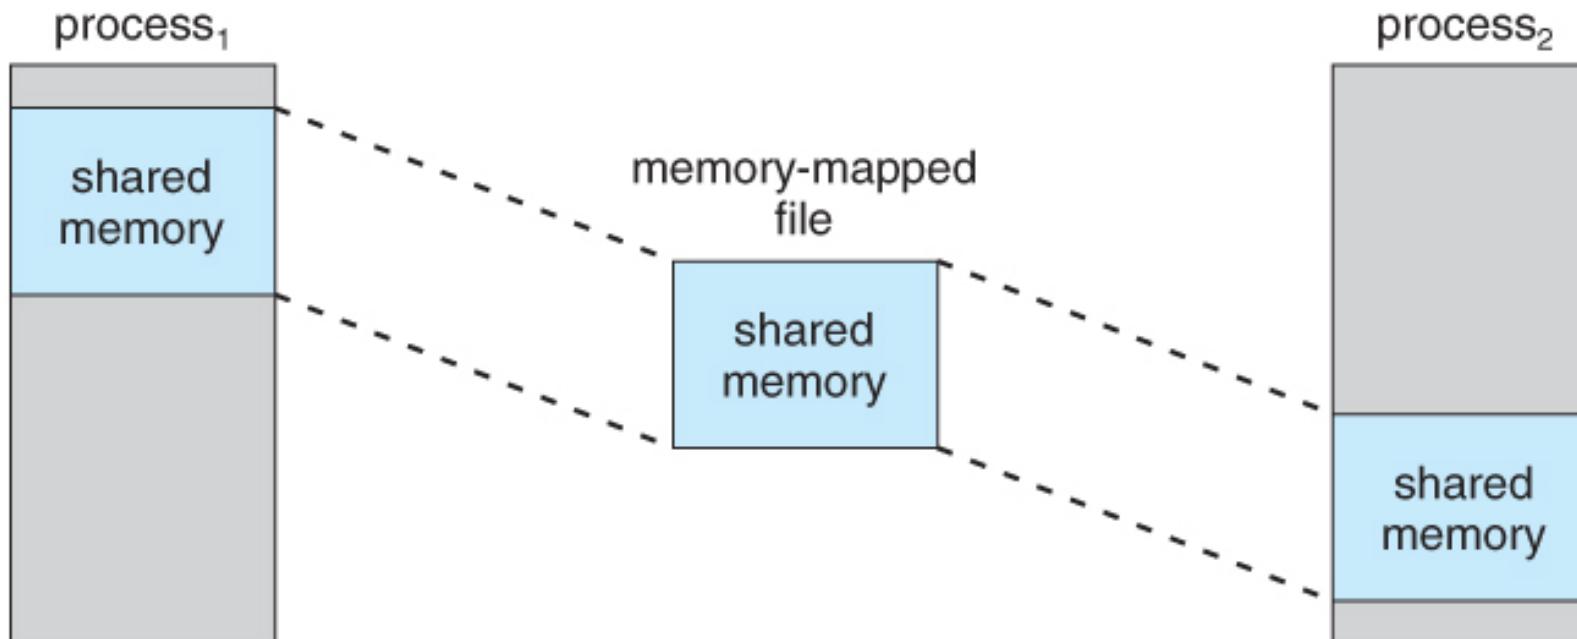


Basic Information	
File Name	Name as chosen by creator (user or program). Must be unique within a specific directory.
File Type	For example: text, binary, load module, etc.
File Organization	For systems that support different organizations
Address Information	
Volume	Indicates device on which file is stored
Starting Address	Starting physical address on secondary storage (e.g., cylinder, track, and block number on disk)
Size Used	Current size of the file in bytes, words, or blocks
Size Allocated	The maximum size of the file
Access Control Information	
Owner	User who is assigned control of this file. The owner may be able to grant/deny access to other users and to change these privileges.
Access Information	A simple version of this element would include the user's name and password for each authorized user.
Permitted Actions	Controls reading, writing, executing, transmitting over a network
Usage Information	
Date Created	When file was first placed in directory
Identity of Creator	Usually but not necessarily the current owner
Date Last Read Access	Date of the last time a record was read
Identity of Last Reader	User who did the reading
Date Last Modified	Date of the last update, insertion, or deletion
Identity of Last Modifier	User who did the modifying
Date of Last Backup	Date of the last time the file was backed up on another storage medium
Current Usage	Information about current activity on the file, such as process or processes that have the file open, whether it is locked by a process, and whether the file has been updated in main memory but not yet on disk

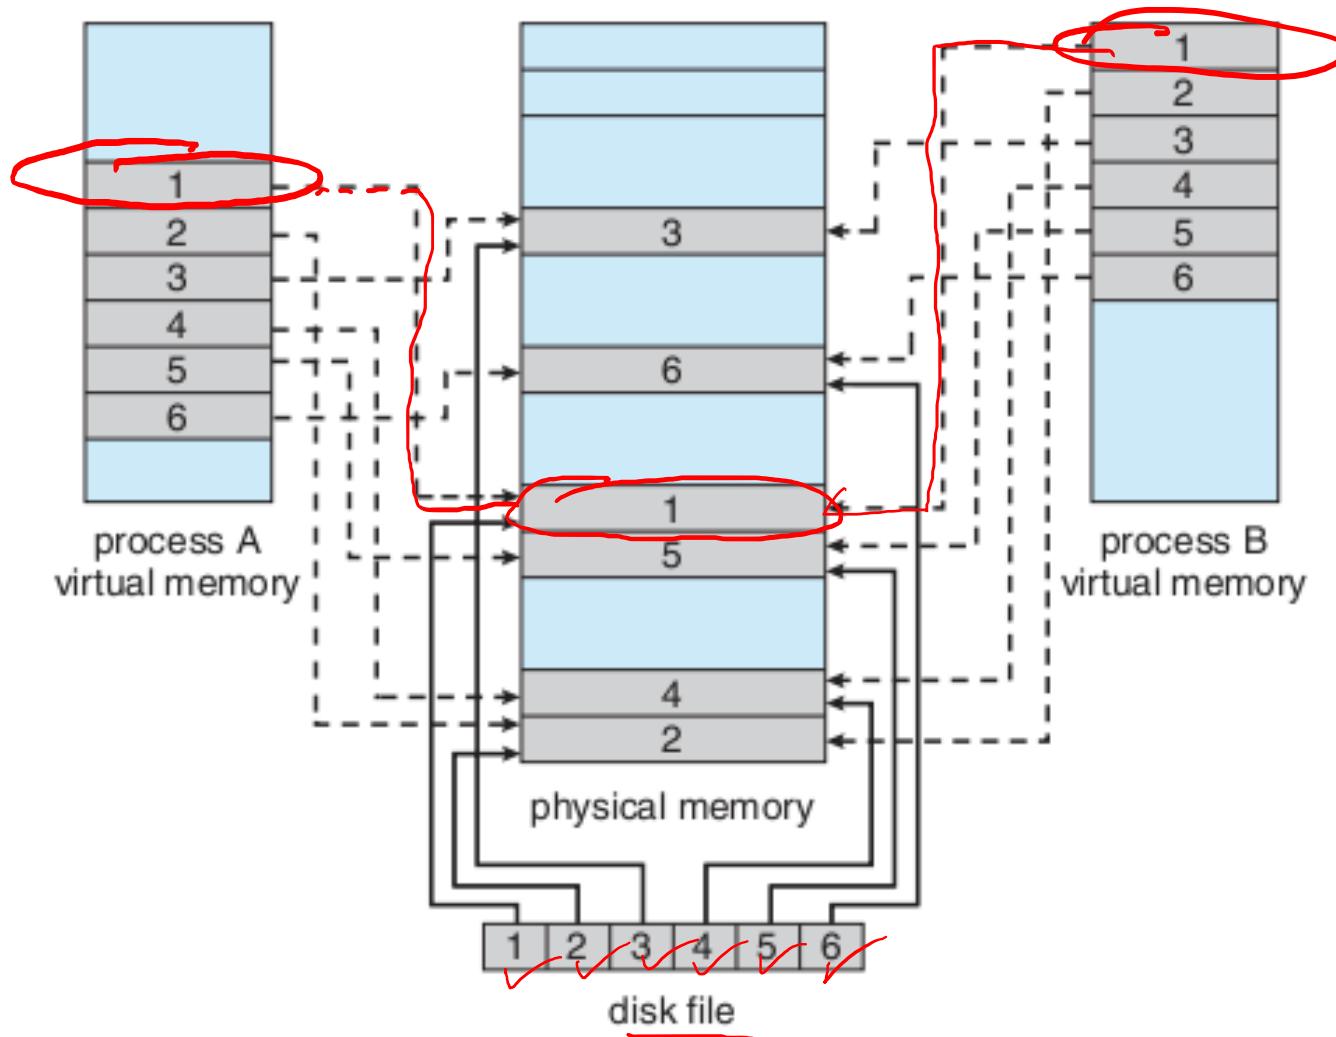
Memory Mapping

- How might we implement memory mapping?
- How might we implement memory mapped *files*?
- Advantages?
- Disadvantages?

Memory Mapping

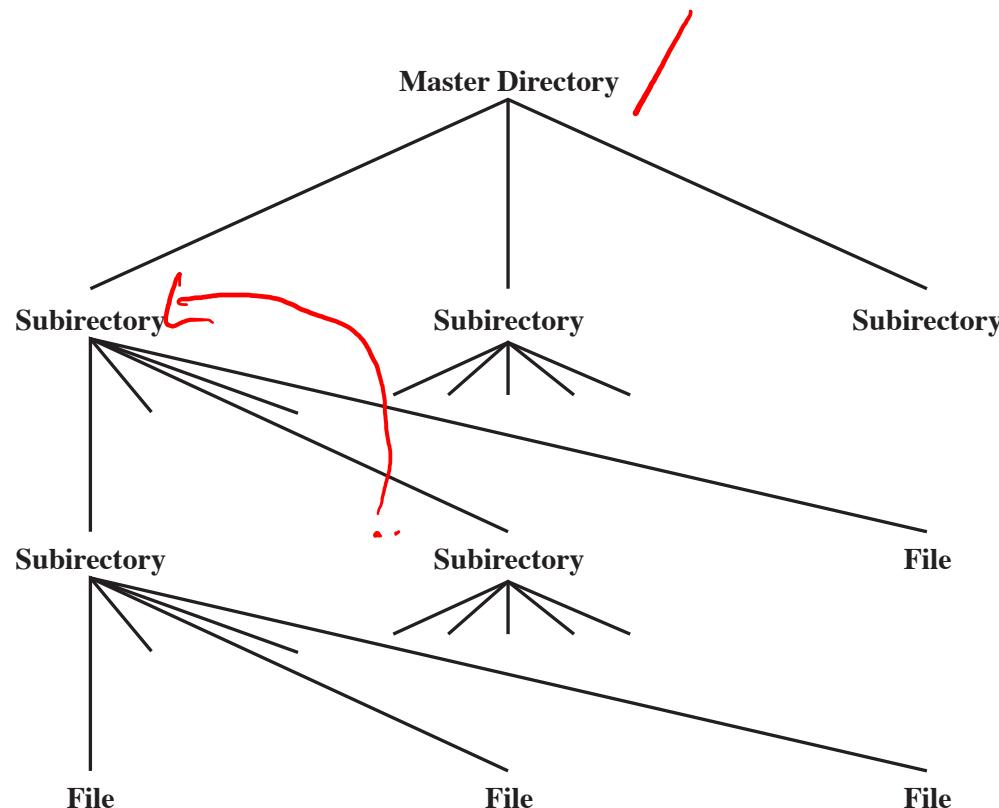


Memory Mapping



File System Organization

Basic File System Organization



- Put files in a **directory**!
- Put files and directories...
...inside directories!

(could there have ever been any other way?!)

- **Basic Operations:**
search, create, delete, list, traverse
- **Basic Concepts:**
current directory, absolute/relative paths, search path
PATH

Basic File System Organization

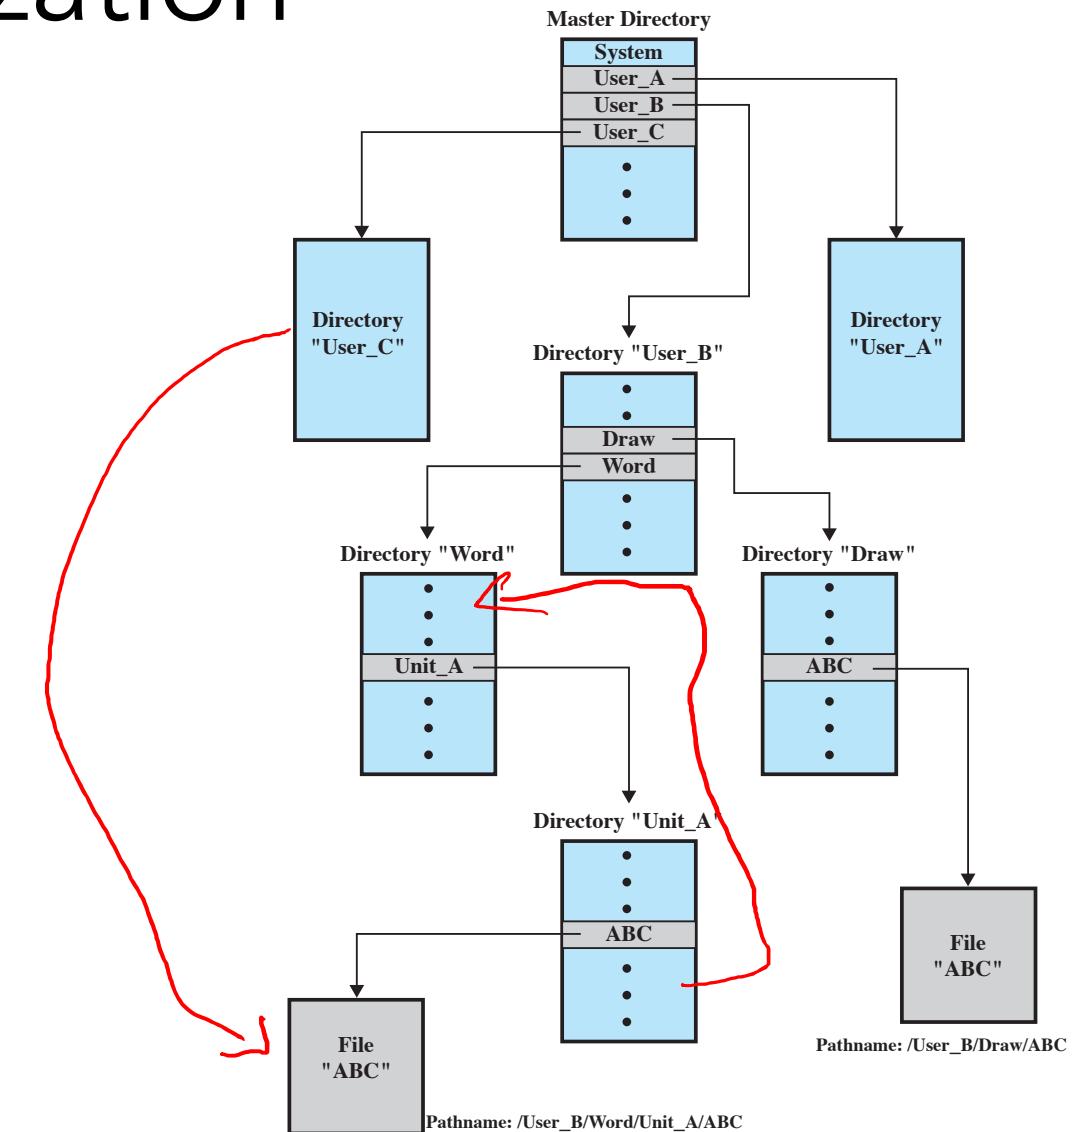
- The file system is a tree, right?

Maybe... maybe not...

- single-level ("flat file system")
- tree
- DAG
- graph

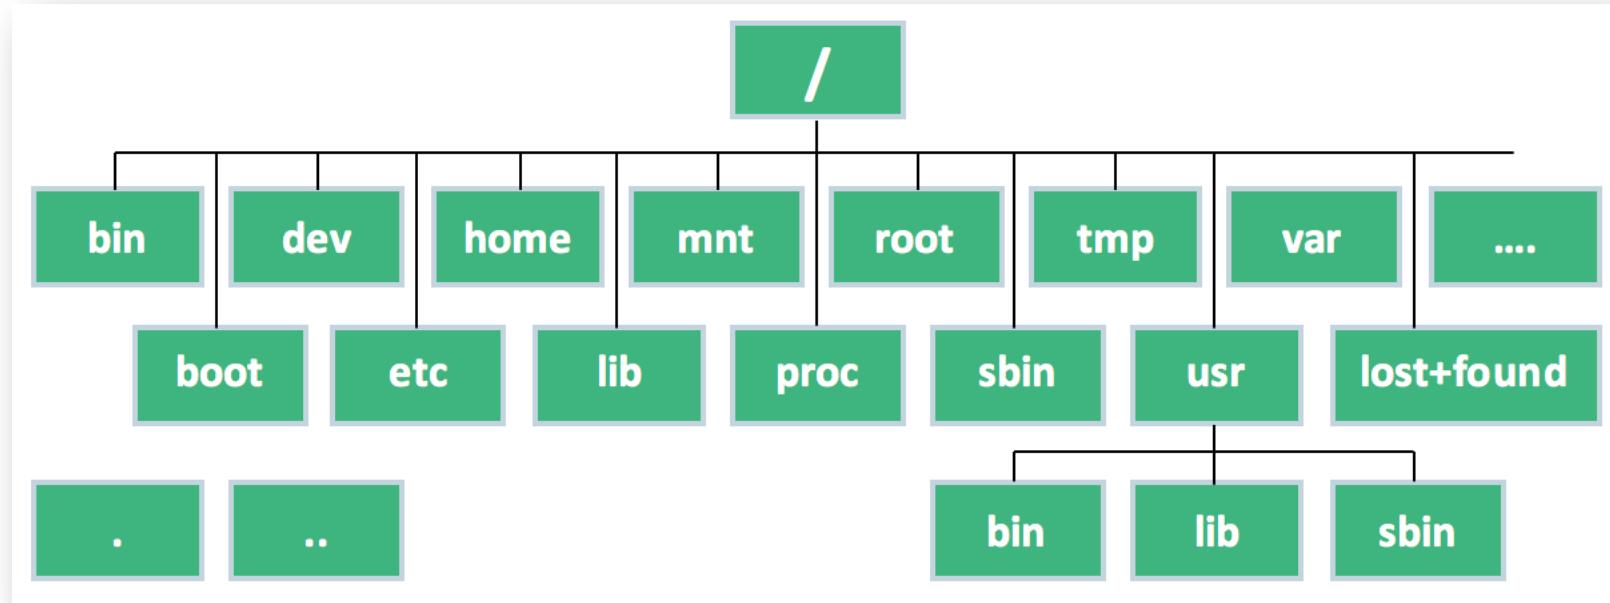
In UNIX/Linux:

- The FS is basically tree with **links**
 - hard links vs. soft links \n

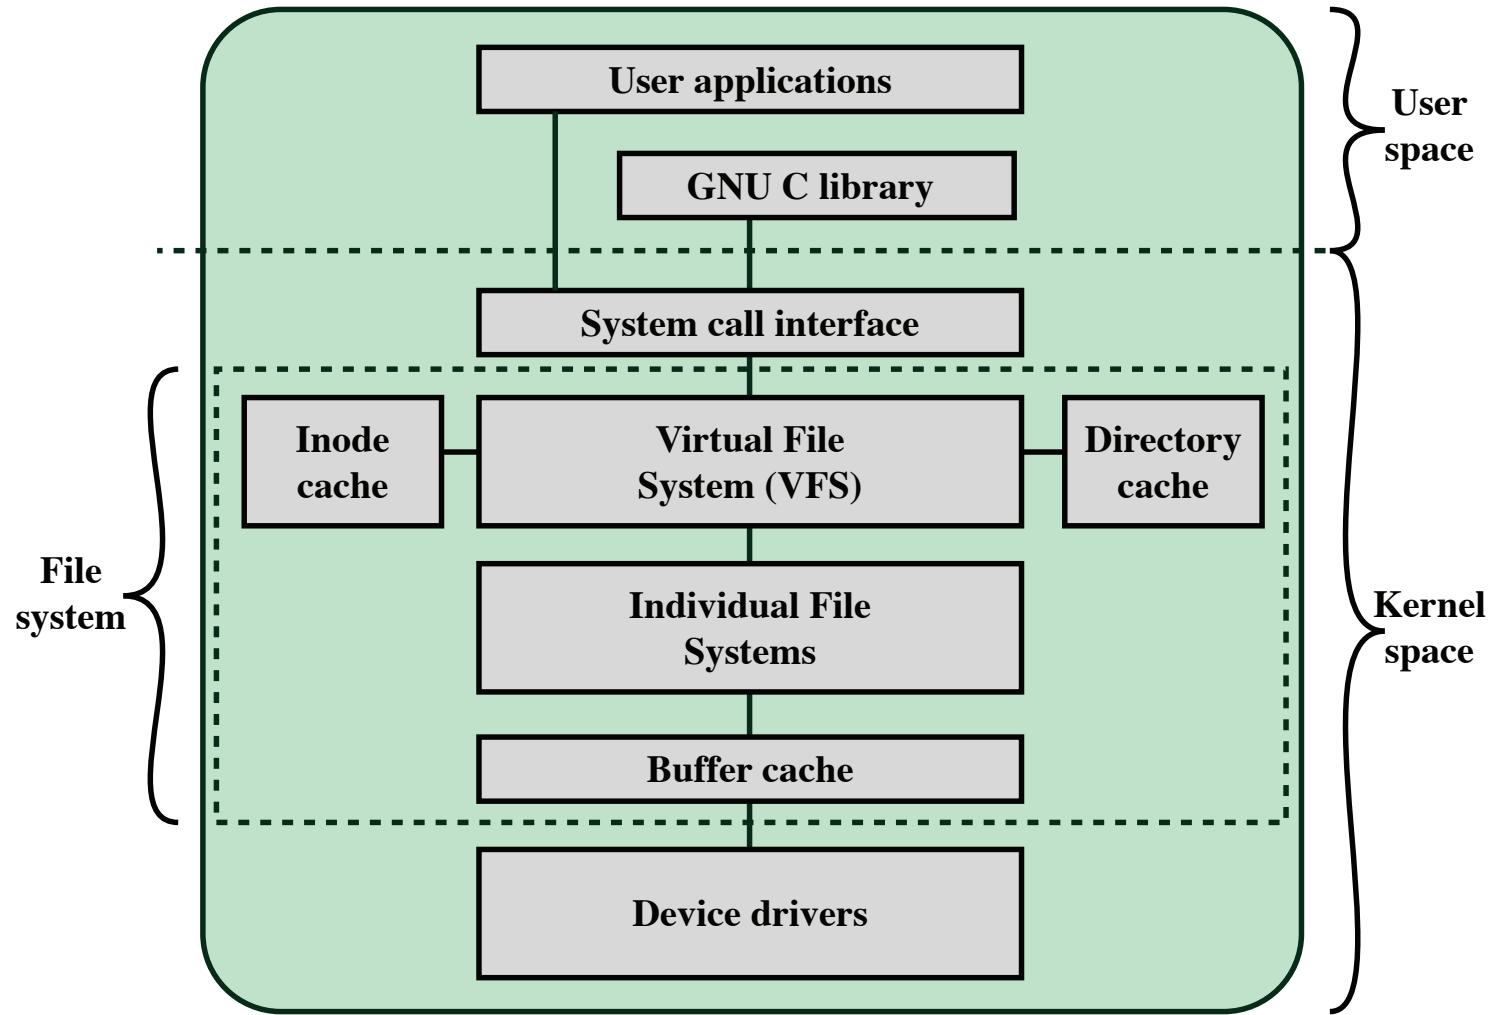


Kitchen Sink File Systems

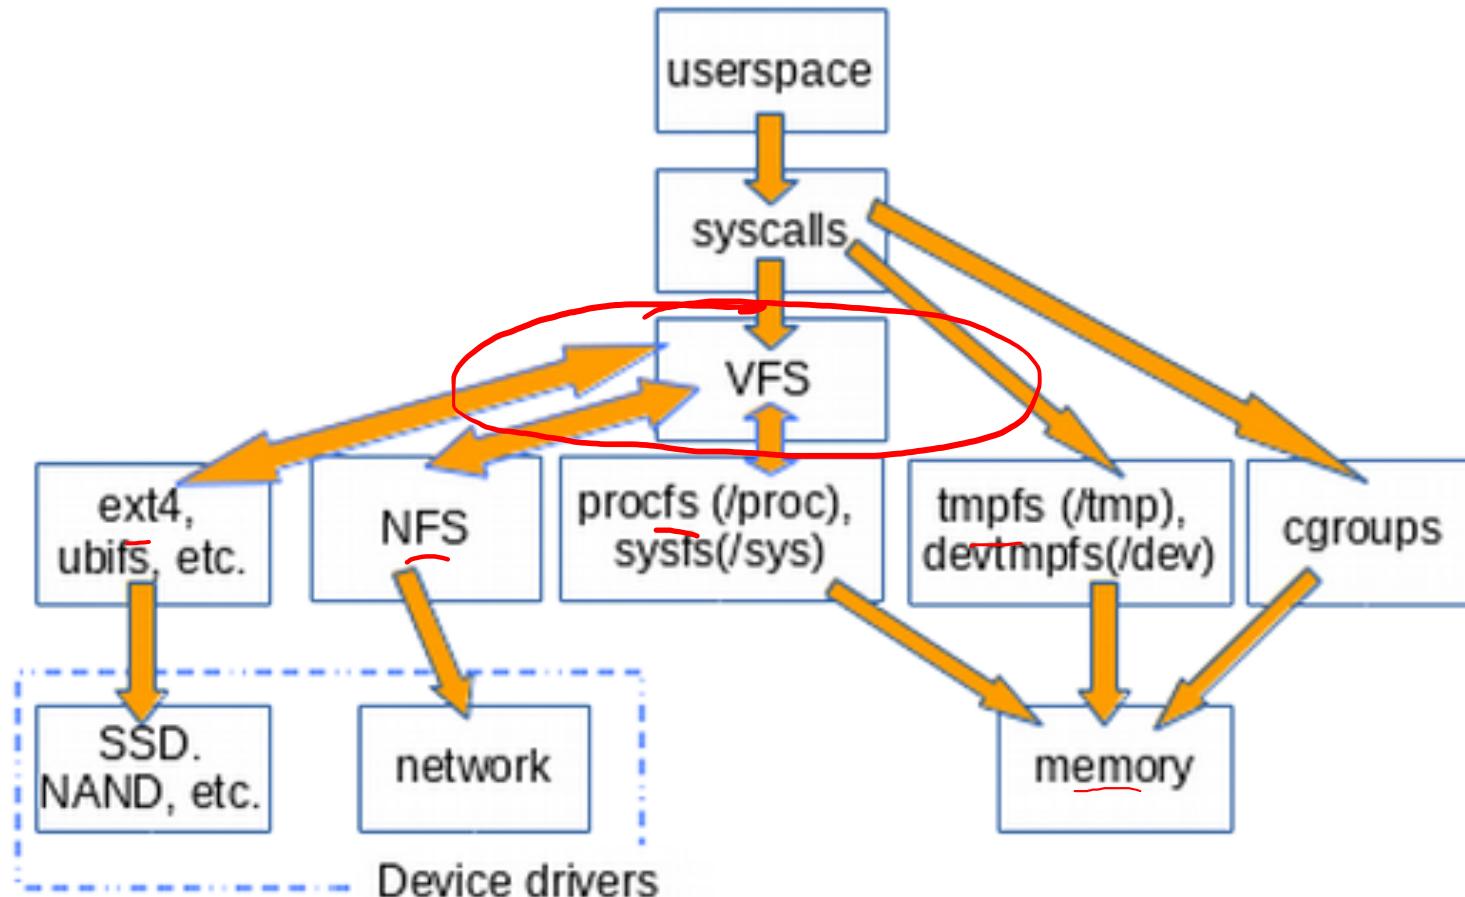
- Files can be more than just a bunch of stored bytes...



Linux VFS



Linux VFS



Summary

Big Ideas

- Files and file systems are where you keep (and organize) stuff
 - > they have various attributes: **name, id, type**, other **bookkeeping**
 - > files have typical operations: **create, delete, read, write** (+memory map, mount)
 - > file systems have typical operations: **create, delete, read, write** (+memory map, mount)
- Virtual File System (VFS)
 - > the abstraction that makes “everything is a file” possible!

What next?

- *Digging deeper into implementation details (e.g., inodes,)*
- *File system types (e.g., FAT, HFS+, EXT*, ZFS)*
- *Security (some next week!)*

Security

File Modes & Permissions

- Every Unix file has a set of **permissions** that determine whether someone* can read, write, or run the file.

- Try: `ls -l ~`
 - Try: `ls -l /dev`

- File mode (4 parts)

-> [file type][user][group][other]

- file type
 - ugo
 - rwx

```
$ls -l file  
-rwxr--r-- owner group ... file
```

- How do we change these settings?

- Ex: allow other members of my **group** to **write** "file"?

Getting Our Hands Dirty!

Some guided exploration of files & file systems - you try!

Fun with file/file system commands

- “.” vs “..”
- ls
- chmod
- umask
- chown / chgrp
- ln
- du
- dup / dup2
- open, read, write, lseek, creat
- tar

Fun with File Descriptors

If two processes each open the same file:

- Do we expect their reads to be interleaved or independent?
- What would this say about whether these two openings of this file have the same file entry, or different file entries?
-> Try filefun2.c with the -2 -s options to see.

If a fork gives the child a copy of the parent's file descriptors:

- If the parent had a file open, is this file open in the child?
- If so, does the child have a different file entry or the same one?
- What would this say about whether their reads would be interleaved or independent?
-> Try filefun2.c with the -2 -i options to see.

What about dup?

- > Try filefun2.c with the -1 -i options to see.

Fun with Links

- On Linux, go to a fairly empty directory.
- Get a medium file in there (say, at least 20K).
- (*I couldn't figure out how to get cat or such to do N bytes, so I wrote bytes.c instead. Recall... stdin, stdout, stderr....) (head -c might also work)*
- Make a hard link to it.
 - Use `ls -lh` to look at the size of things in the directory.
 - Use `du -h` to look at the sizes of individual things.
 - Use `du -h .` to look at the size of the directory.

-> Do things add up?

-> Try making a soft link, and using `du` to look at sizes. *What's going on?*

Fun with Permissions

- You might wonder, with group permissions, what if the requester is in several groups?
- (Here, we'll step over to Linux...)
- I created a dir **permtest/** with various files. I, as user **vagrant**, can read, say, file **file_test**.
- (I'm in many groups - grep in /etc/group)
- But what happens if I try to read **file_test_2**?
- Suppose I also create **subdir/file**, and two hardlinks and softlinks to it.
- What happens if I chmod **hardlink1**?
- What happens if I chmod **softlink1**?
- What happens if I chmod **subdir** to remove userx permissions? Do the softlinks still work? Do the hardlinks?
- What might you conclude from all of this about potential implementation details?