



Design Issues in Multiprocessor, Multicore, and Real-time Scheduling on Linux Machines

Hannah Cebulla, Bridget Wermers, John Hultman, Logan Shy



Multiprocessor Scheduling Design Issues

- Goal of Multiprocessor System
- Granularity
- Types of Multiprocessor Systems
 - Loosely Coupled
 - I/O processor
 - Tightly Coupled



Multicore Scheduling Design Issues

- Definition of Multicore system
- Scheduling
 - Locality
 - Caches
- Cooperative resource sharing
- Resource Contention



Real-time Scheduling Algorithms

- Hard vs Soft Deadlines
- Periodic vs Aperiodic
- Characteristics
 - Determinism
 - Responsiveness
 - User Control
 - Reliability
 - Fail Soft Operations
- Approaches
 - Static table-driven
 - Static priority-driven
 - Dynamic Planning-based
 - Dynamic Best Effort



Linux Kernel Scheduling Design

- Similar to UNIX
- Important scheduling changes
 - 1.2: Round Robin
 - 2.4: Priorities, Queues, and Multicore
 - 2.6: CFS
- CFS implementation
 - Dynamic time quantum
 - Target Latency
 - RB Tree
- Synchronization primitives
 - Atomic
 - Spinlock (r/w)
 - Mutex



Algorithm Comparison

- Stochastic Search Scheduling Method based on the Genetic Algorithm
- Dynamic Critical-Path Scheduling
- Earliest Deadline First (EDF)
- Rate Monotonic (RM)
- Declustering
- Priority Inversion
- Contention Aware Scheduling



Conclusions

- Scheduling is hard (NP-hard!)
- Key differences between multicore and multiprocessor systems
- Schedule design must vary by nature of system (laptop vs space shuttle)
- Many attempts at optimizing under varying circumstances
- Linux CFS and its many applications